

IFT6390

Fondements de l'apprentissage machine

Linear regression
and
Regularized linear regression

Professor: Ioannis Mitliagkas

Slides: Pascal Vincent

Supervised task

predict y from x

input

$$\mathbf{x} \in \mathbb{R}^d$$

target

(label)

y

n examples

\mathbf{x}_1	\mathbf{x}_2	\mathbf{x}_3	\mathbf{x}_4	\mathbf{x}_5	t
0.32	-0.27	+1	0	0.82	113
-0.12	0.42	-1	1	0.22	34
0.06	0.35	-1	1	-0.37	56
0.91	-0.72	+1	0	-0.63	77
...

Training set D_n

$$D_n = \{(x^{(1)}, t^{(1)}), \dots, (x^{(n)}, t^{(n)})\}$$

Learn a function

$f = f_\theta$ which
minimizes a cost (loss).

loss function

$$L(f_\theta(\mathbf{x}), y)$$

output/
prediction

$$f_\theta(\mathbf{x})$$

f_θ : parameters

-0.12 0.42 -1 1 0.22

input x

34

target y

Empirical risk minimization

We must specify:

- A parametric form for our functions, f_θ
- A specific cost (loss) function $L(y, t)$

So we define the empirical risk as:

$$\hat{R}(f_\theta, D_n) = \sum_{i=1}^n L(f_\theta(\mathbf{x}^{(i)}), t^{(i)})$$

i.e. total loss on the training set

Learning amounts to finding the optimal values for the parameters:

$$\theta^* = \arg \min_{\theta} \hat{R}(f_\theta, D_n)$$

It is the principle of empirical risk minimization.

Eg: Linear regression

A very simple learning algorithm

We select

A linear (affine) form for the function:

$$f_{\theta}(\mathbf{x}) = \underbrace{\langle \mathbf{w}, \mathbf{x} \rangle}_{\text{scalar product (inner product)}} + b$$

$$\theta = \{\mathbf{w}, b\}, \underbrace{\mathbf{w} \in \mathbb{R}^d}_{\text{"weight vector"}}, \underbrace{b \in \mathbb{R}}_{\text{bias}}$$

Cost: quadratic error:

$$L(y, t) = (y - t)^2$$

Principle of empirical risk minimization (ERM)

We look for the parameters that minimize the empirical risk

$$\theta^* = \arg \min_{\theta} \hat{R}(f_{\theta}, D_n)$$

Linear regression



Neural inspiration

Intuitive understanding of the scalar product

each component of \mathbf{x} has a weighted influence on the output y

$$y = f_{\theta}(\mathbf{x}) = \mathbf{w}_1\mathbf{x}_1 + \mathbf{w}_2\mathbf{x}_2 + \dots + \mathbf{w}_d\mathbf{x}_d + b$$

Neural terminology



Regularized empirical risk

It is often necessary to induce a "preference" for some parameter values rather than others to avoid overfitting

We define **regularized empirical risk** as follows:

$$\hat{R}_\lambda(f_\theta, D_n) = \underbrace{\left(\sum_{i=1}^n L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) \right)}_{\text{empirical risk}} + \underbrace{\lambda \Omega(\theta)}_{\text{regularization term (penalty)}}$$

Ω penalizes more or less the different parameter values.

$\lambda \geq 0$ the importance of this regularization term

(in relation to the empirical risk)

Eg: Ridge Regression

= linear regression + quadratic (L2) regularization

We penalize the large weights

$$\Omega(\theta) = \Omega(\mathbf{w}, b) = \|\mathbf{w}\|^2 = \sum_{j=1}^d \mathbf{w}_j^2$$

Neural terminology:
“weight decay” penalty

$$\hat{R}_\lambda(f_\theta, D_n) = \underbrace{\left(\sum_{i=1}^n L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) \right)}_{\text{empirical risk}} + \underbrace{\lambda \Omega(\theta)}_{\text{regularization term (penalty)}}$$

Eg: Ridge regression

= linear regression + quadratic (L2) regularization

Regularized empirical risk

$$\hat{R}_\lambda(f_\theta, D_n) = \underbrace{\left(\sum_{i=1}^n L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) \right)}_{\text{Empirical risk}} + \underbrace{\lambda \Omega(\theta)}_{\text{regularization term (penalty)}}$$

We are looking for the parameter values that minimize this objective

$$\{\mathbf{w}^*, b^*\} = \theta^* = \underset{\theta}{\operatorname{argmin}} \hat{R}_\lambda(f_\theta, D_n)$$

Eg: Ridge Regression

= linear regression + quadratic (L2) regularization

- For linear regression or ridge regression a little linear algebra gives us an **analytical solution**

we solve for $\theta = \{b, \mathbf{w}\}$: $\frac{\partial \hat{R}_\lambda(f_\theta, D_n)}{\partial \theta} = 0$

we obtain:
$$\begin{pmatrix} b^* \\ \mathbf{w}^* \end{pmatrix} = (\check{X}^T \check{X} + \lambda \check{I})^{-1} \check{X}^T \mathbf{t}$$

$$\text{où } \check{X} = \begin{pmatrix} 1 & \mathbf{x}_1^{(1)} & \dots & \mathbf{x}_d^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \mathbf{x}_1^{(n)} & \dots & \mathbf{x}_d^{(n)} \end{pmatrix}, \mathbf{t} = \begin{pmatrix} t^{(1)} \\ \vdots \\ t^{(n)} \end{pmatrix} \quad \check{I} = \begin{pmatrix} 0 & 0 & & \\ 0 & 1 & 0 & \\ & 0 & 1 & 0 \\ & & \ddots & \ddots \\ & & 0 & 1 \end{pmatrix}$$

- Most of the time (other choices of f and L) we do not have an analytical solution.
- More generally, we can use a **gradient descent method**.

Other possibility: optimization by gradient descent

$$\hat{R}_\lambda(f_\theta, D_n) = \underbrace{\left(\sum_{i=1}^n L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) \right)}_{\text{empirical risk}} + \underbrace{\lambda \Omega(\theta)}_{\text{regularization term}}$$

- we initialize the parameters randomly
- we update them iteratively following the gradient

Either **batch gradient descent** (whole dataset):

Loop: $\theta \leftarrow \theta - \eta \frac{\partial \hat{R}_\lambda}{\partial \theta}$

$$= \left(\sum_{i=1}^n \frac{\partial}{\partial \theta} L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) \right) + \lambda \frac{\partial}{\partial \theta} \Omega(\theta)$$


Or **stochastic gradient descent**:

Loop:

For i in 1...n

$$\theta \leftarrow \theta - \eta \frac{\partial}{\partial \theta} \left(L(f_\theta(\mathbf{x}^{(i)}), t^{(i)}) + \frac{\lambda}{n} \Omega(\theta) \right)$$

Or **other variants of the gradient descent idea**
(conjugate gradient, Newton's method, natural gradient, ...)



Various regularizers

«Ridge»: regularization, L_2

In Bayesian terms: corresponds to a Gaussian prior on the weights

$$\Omega(\theta) = \Omega(\mathbf{w}, b) = \|\mathbf{w}\|_2^2 = \sum_{j=1}^d \mathbf{w}_j^2$$

«Lasso»: regularization, L_1

In Bayesian terms: corresponds to a Laplacian prior on the weights

$$\Omega(\theta) = \Omega(\mathbf{w}, b) = \|\mathbf{w}\|_1 = \sum_{j=1}^d |\mathbf{w}_j|$$


=> automatic selection of components (a number of weights will be zero)

«Elastic net»: combination of the two

$$\Omega(\theta) = \Omega(\mathbf{w}, b) = \lambda_1 \|\mathbf{w}\|_1 + \lambda_2 \|\mathbf{w}\|_2^2$$

Etc...

Visualizing L_1 regularization





- If λ is big enough, the circle is very likely to intersect the diamond at one of the corners
- This makes L_1 regularization much more likely to make some weights *exactly* 0

Pros and cons of L_1 regularization

- If there are irrelevant input features, Lasso is likely to make their weights 0, while L_2 is likely to just make all weights small
- Lasso is biased towards providing *sparse solutions* in general
- Lasso optimization is computationally more expensive than L_2
- More efficient solution methods have to be used for large numbers of inputs (e.g. least-angle regression, 2003).
- L_1 methods of various types are very popular
- One can combine L_1 and L_2 regularization (elastic-net)

Example of L1 vs L2 effect



- Note the sparsity in the coefficients induced by L_1
- Lasso is an efficient way of performing the L_1 optimization