

# Computer Vision

## IFT6758 - Data Science

### Sources:

<http://www.cs.cmu.edu/~16385/>

<http://cs231n.stanford.edu/2018/syllabus.html>

<http://www.cse.psu.edu/~rtc12/CSE486/>

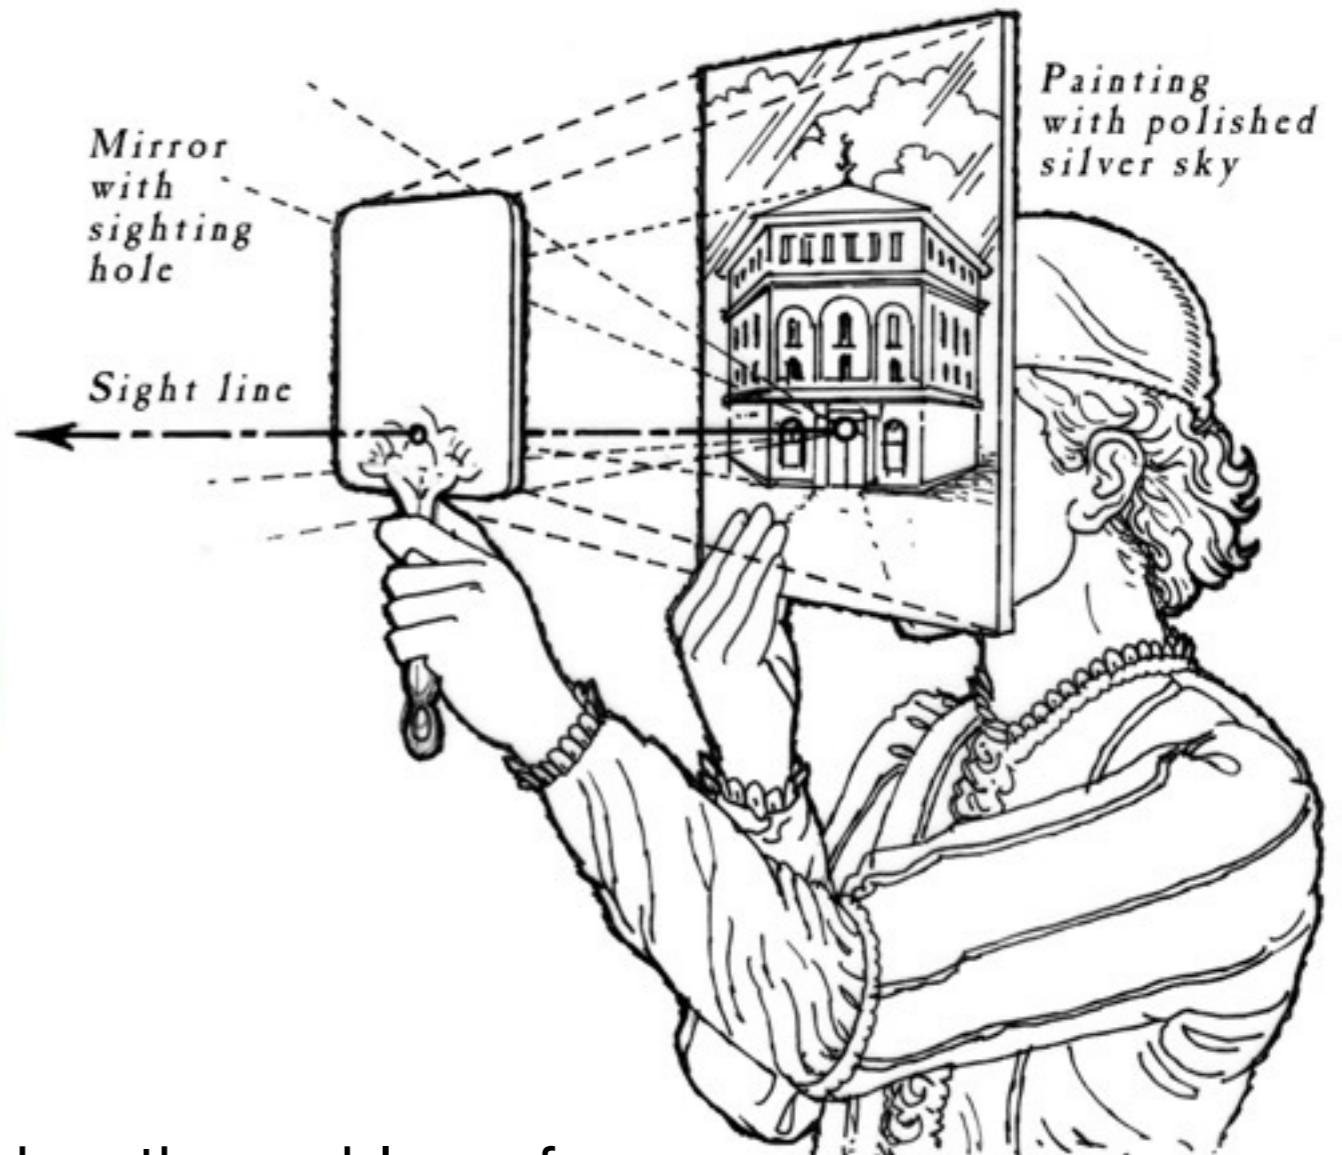
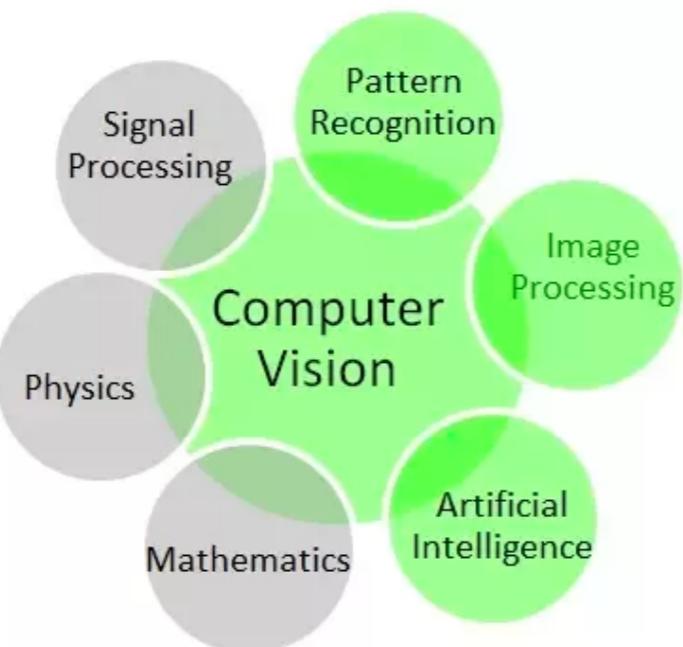
# What is Computer vision?

"

Vision is the act of knowing what is where by looking.

"

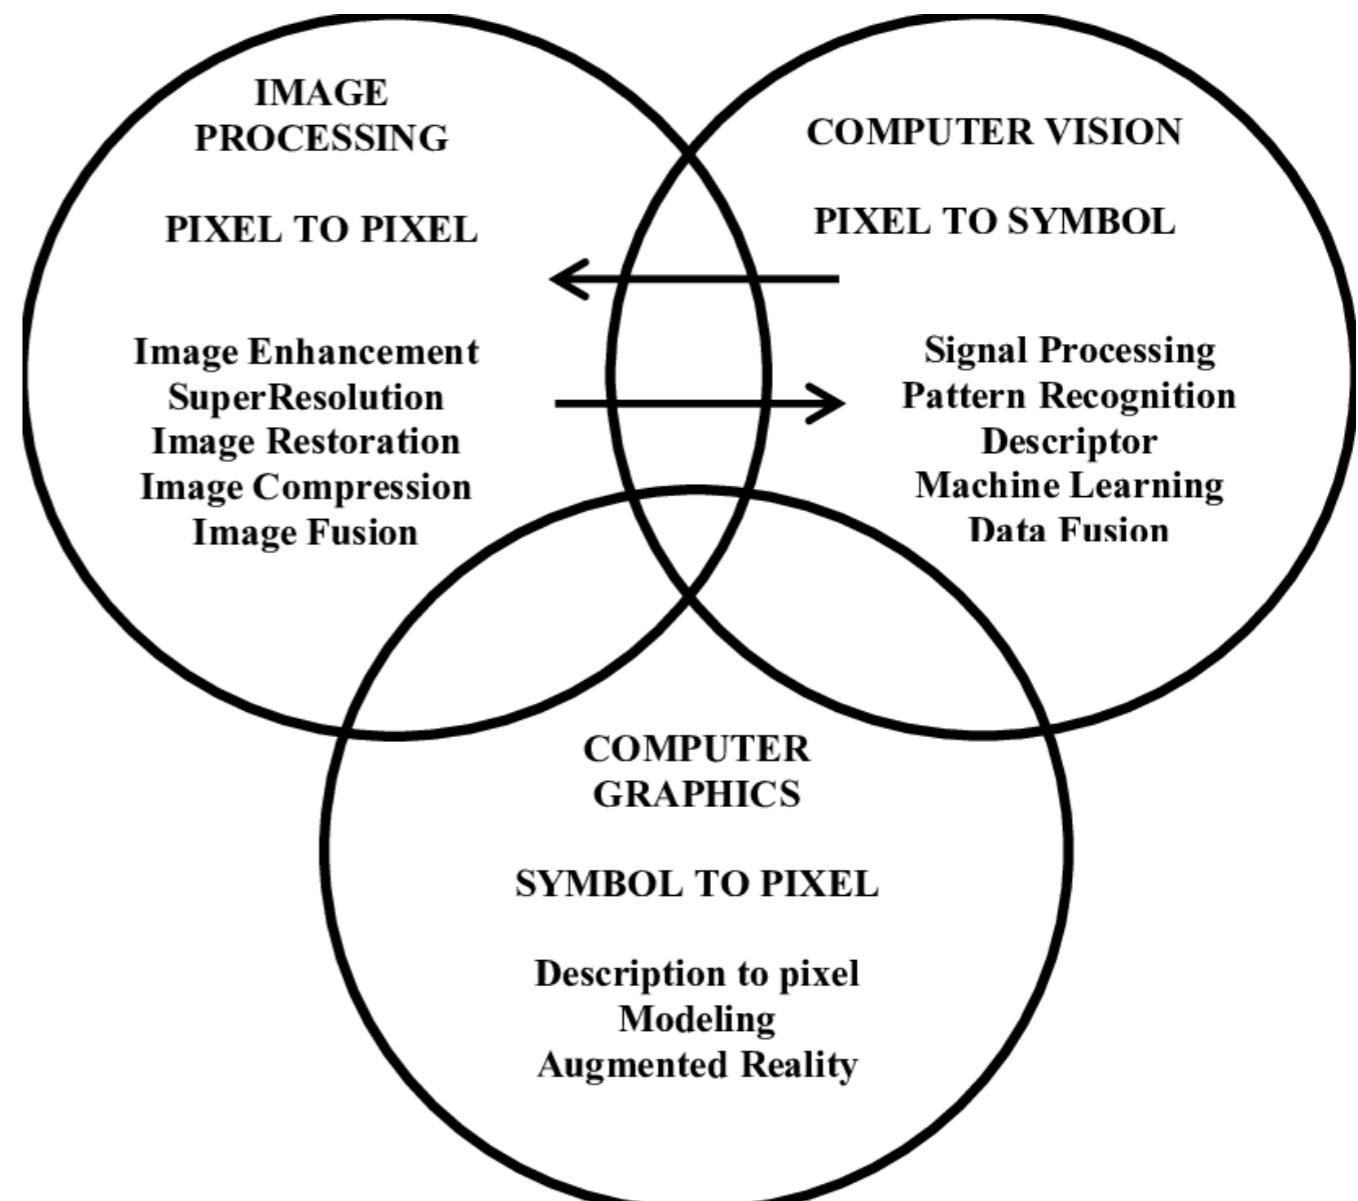
--Aristotle



- Computer vision is a field of study focused on the problem of helping computers to see.

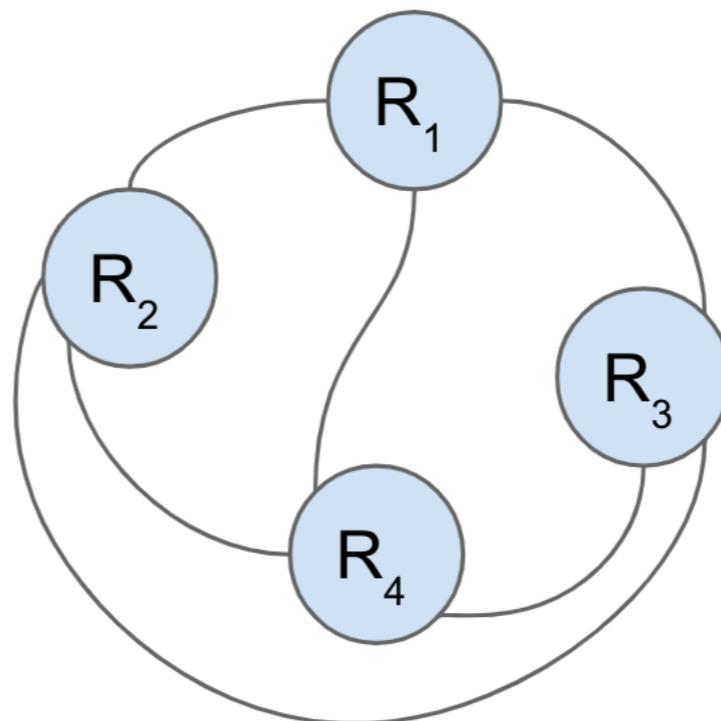
# Computer vision vs. Image processing

- Computer vision is distinct from image processing.
- Image processing is the process of **creating a new image** from an existing image, typically simplifying or enhancing the content in some way.
- Computer vision is concerned with understanding the content of an image.



# CV tasks (4 Rs)

1. Reconstruction
2. Registration
3. Reorganization
4. Recognition



# CV tasks (4 Rs)

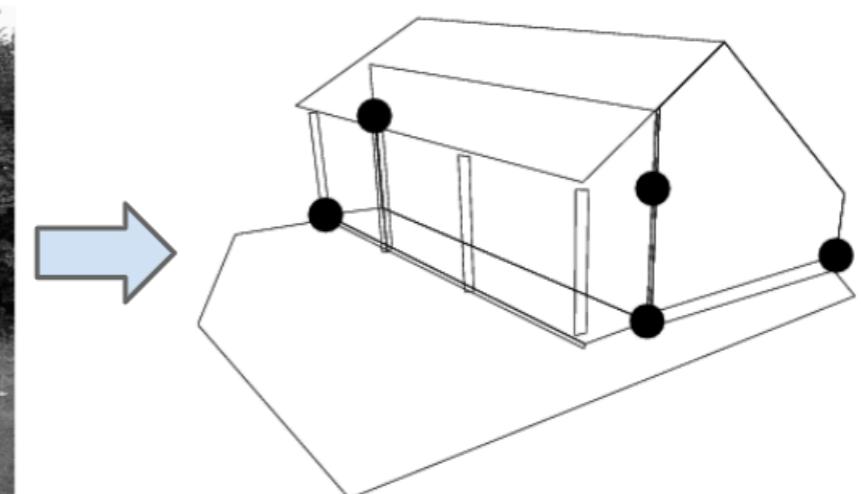
## 1. Reconstruction



Left View



Right View



Multiview Geometry, 3D Vision, Shape-from-X

# CV tasks (4 Rs)

## 2. Registration

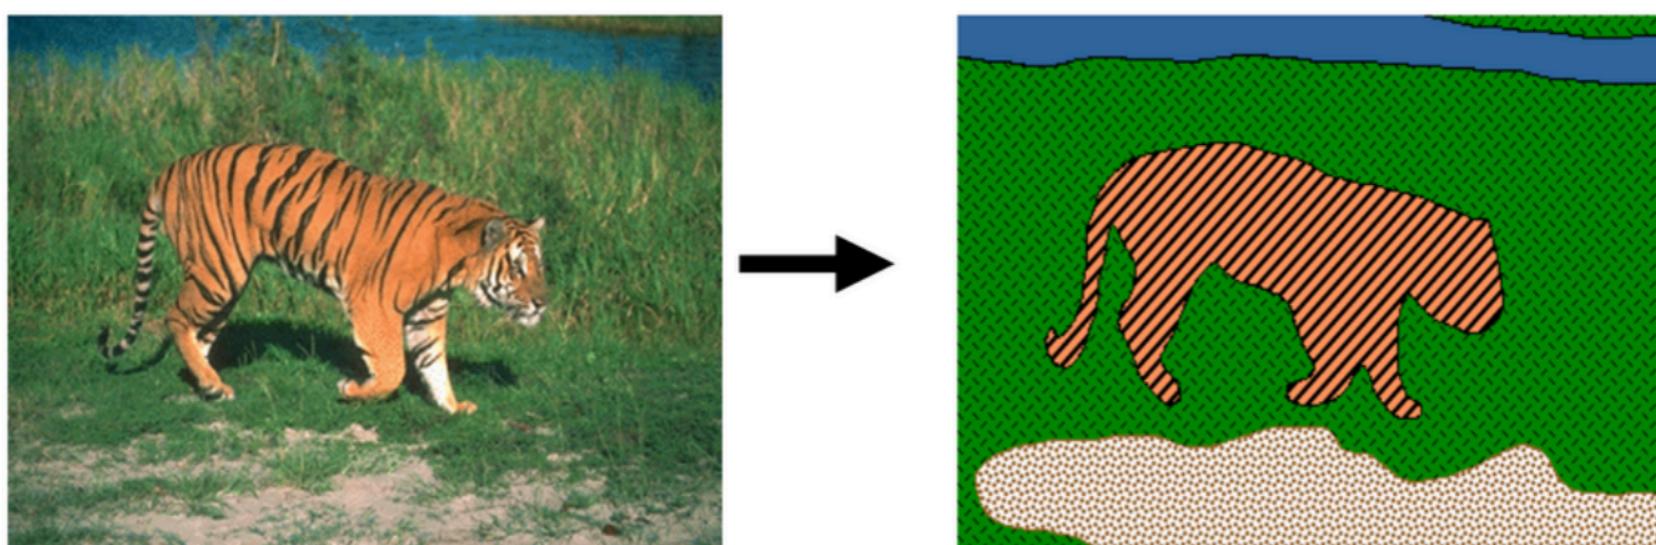


Tracking, Alignment, Optical Flow, Correspondence

# CV tasks (4 Rs)

Clustering, Unsupervised Learning, Segmentation, Perceptual Organization

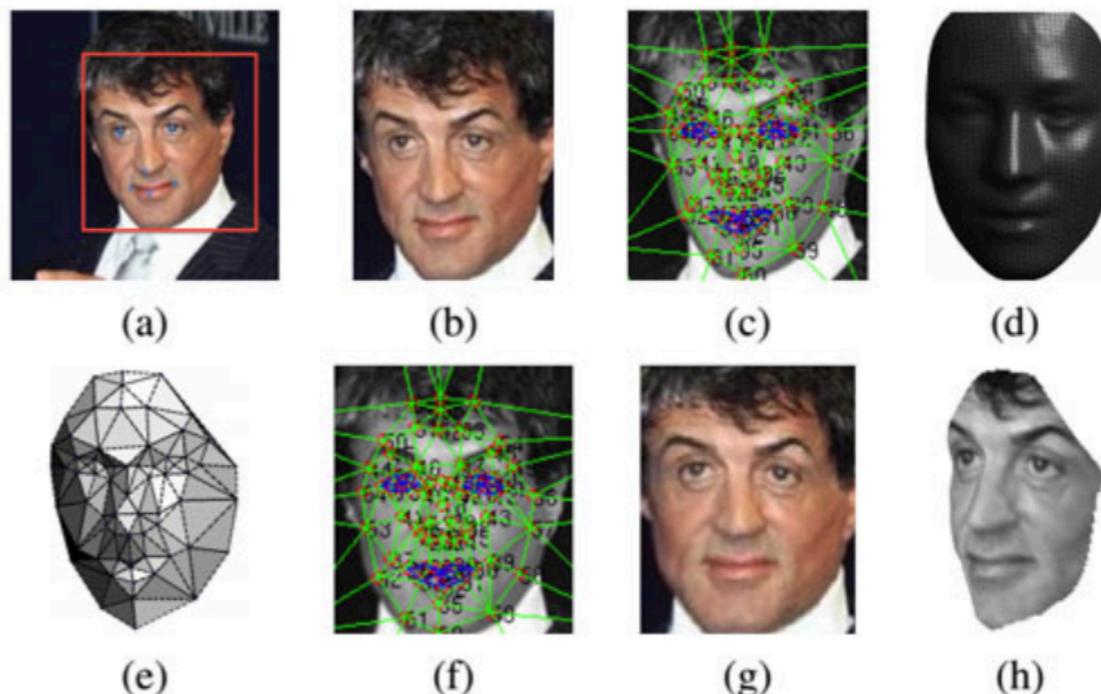
## 3. Reorganization



# CV tasks (4 Rs)



[Ricoh]



[DeepFace]

## 4. Recognition

Verification, Identification, Detection

# Why study Computer Vision?

- Images and movies are everywhere
- Fast-growing collection of useful applications
  - building representations of the 3D world from pictures
  - automated surveillance (who's doing what)
  - movie post-processing
  - face finding
- Greater understanding of human vision

# Earth view (3d modelling)

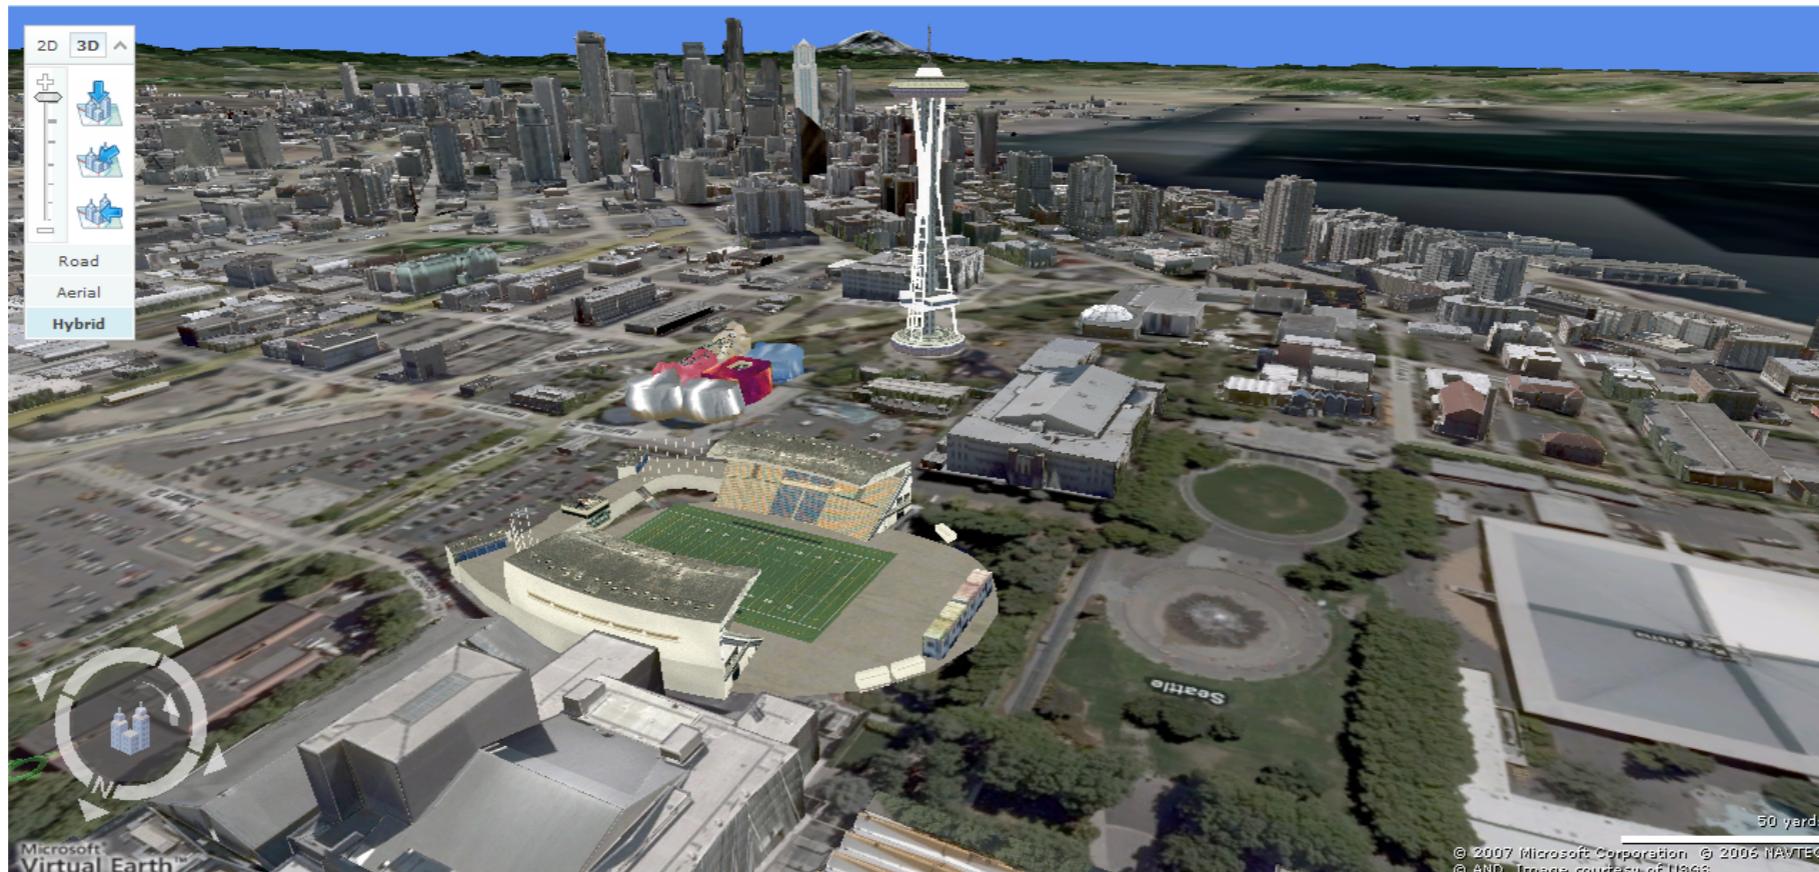
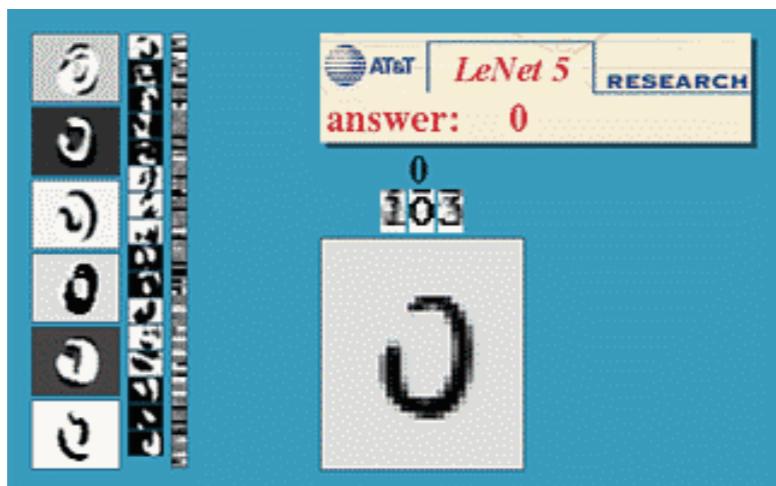


Image from Microsoft's [Virtual Earth](#)  
(see also: [Google Earth](#))

# Optical character recognition (OCR)

Technology to convert scanned docs to text

- If you have a scanner, it probably came with OCR software



Digit recognition, AT&T labs  
<http://www.research.att.com/~yann/>



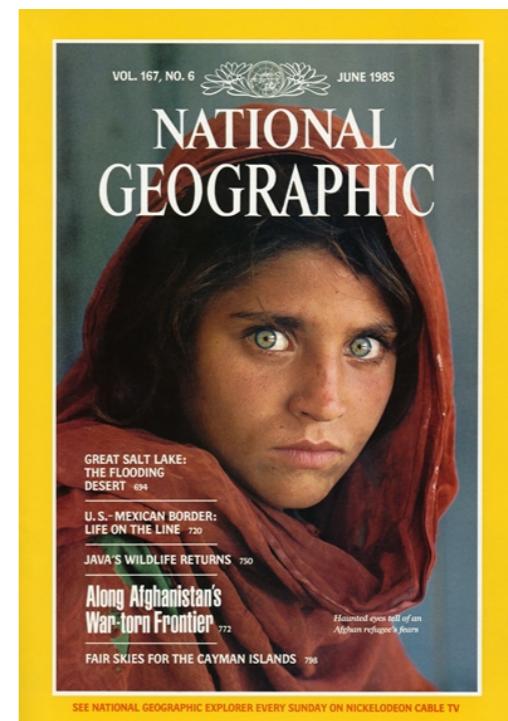
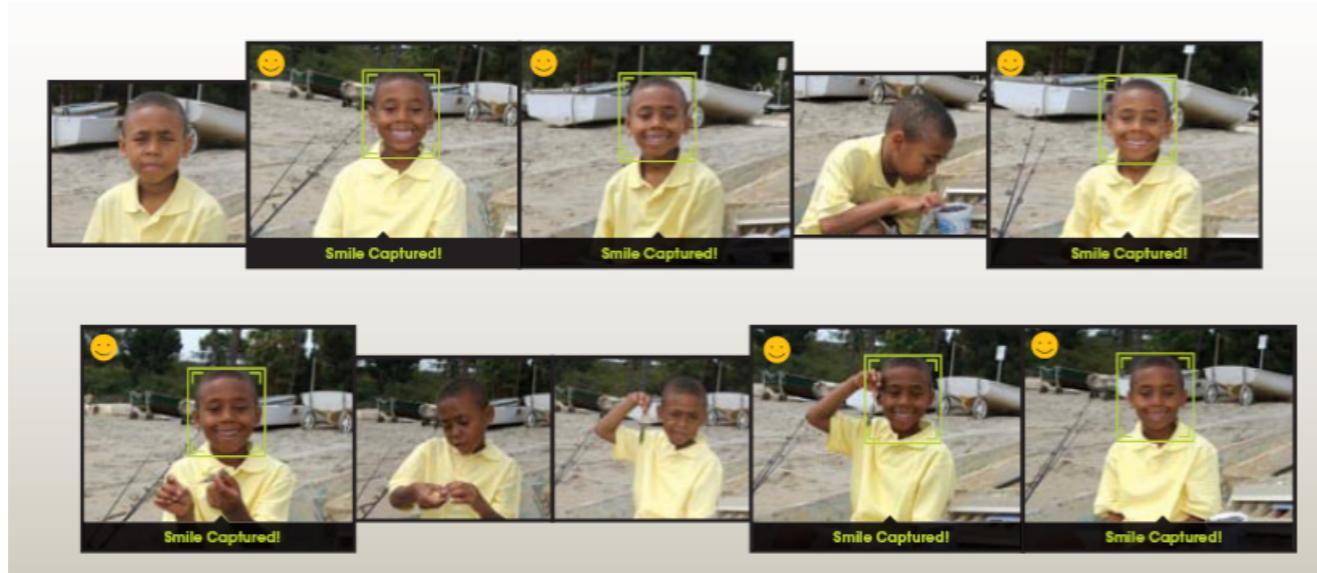
License plate readers  
[http://en.wikipedia.org/wiki/Automatic\\_number\\_plate\\_recognition](http://en.wikipedia.org/wiki/Automatic_number_plate_recognition)

# Face and smile detection



Many new digital cameras now detect faces

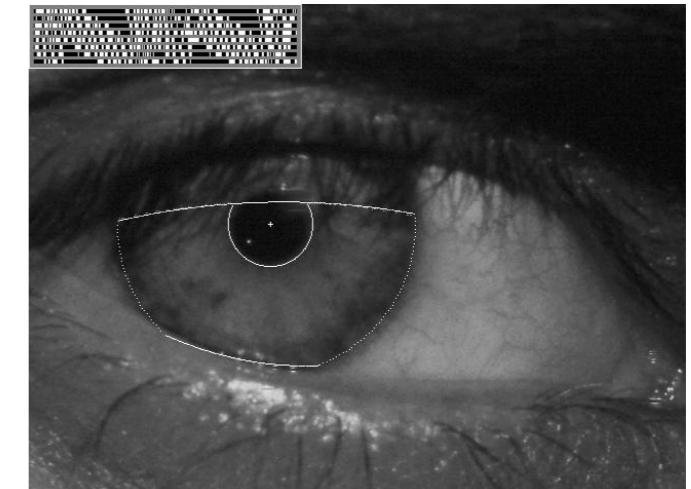
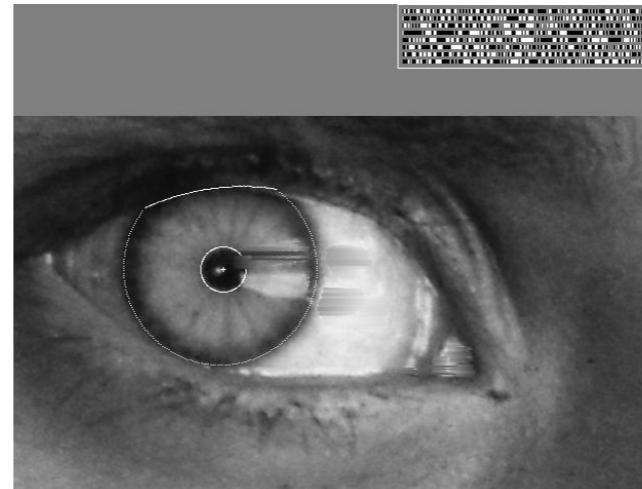
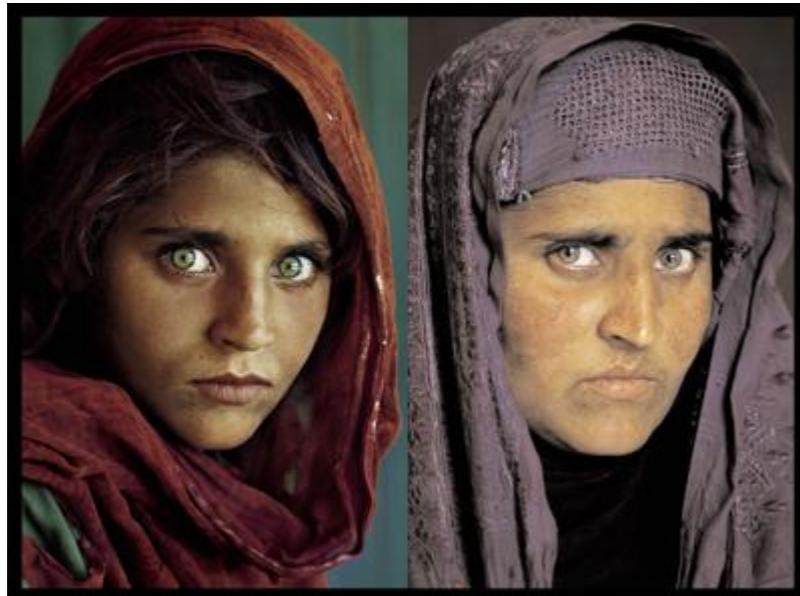
- Canon, Sony, Fuji, ...



Who is she?

[Sony Cyber-shot® T70 Digital Still Camera](#)

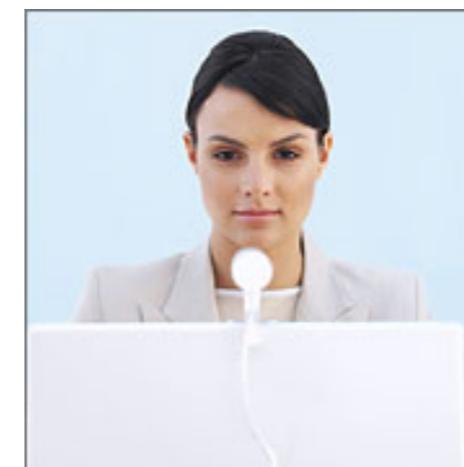
# Vision biometric



*"How the Afghan Girl was Identified by Her Iris Patterns"* Read the [story](#)



Fingerprint scanners on  
many new laptops,  
other devices



Face recognition systems now  
beginning to appear more widely  
<http://www.sensiblevision.com/>

# Object recognition



This is becoming real:

- Microsoft Research
- [Point & Find, Nokia](#)

## [LaneHawk by EvolutionRobotics](#)

“A smart camera is flush-mounted in the checkout lane, continuously watching for items...”

# Sports and games



*Sportvision* first down line  
Nice [explanation](#) on [www.howstuffworks.com](http://www.howstuffworks.com)

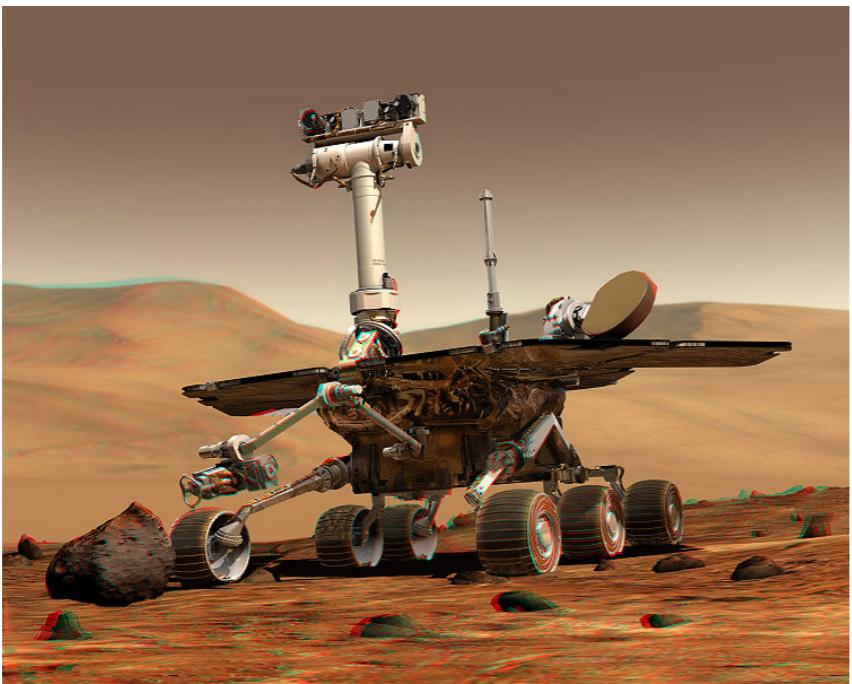


[Digimask](#): put your face on a 3D avatar.

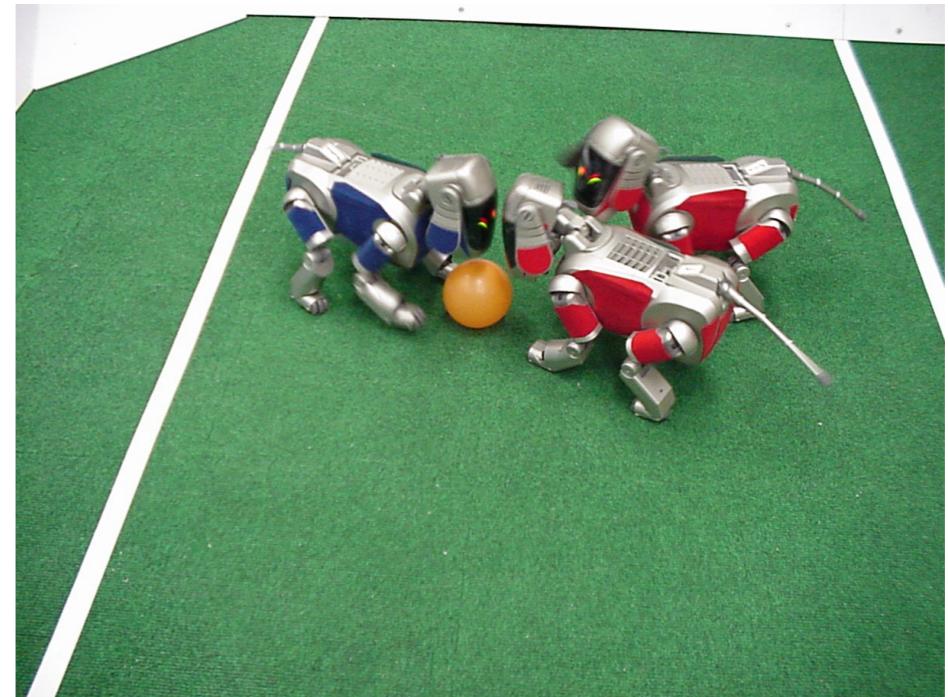


Nintendo Wii has camera-based IR tracking built in.

# Robotics

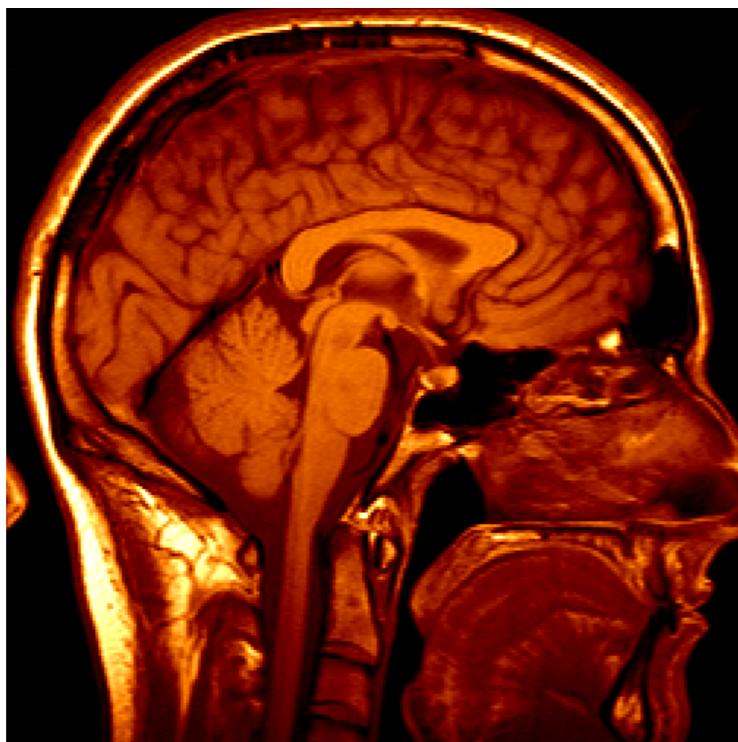


NASA's Mars Spirit Rover  
[http://en.wikipedia.org/wiki/Spirit\\_rover](http://en.wikipedia.org/wiki/Spirit_rover)



<http://www.robocup.org/>

# Medical imaging



3D imaging  
MRI, CT

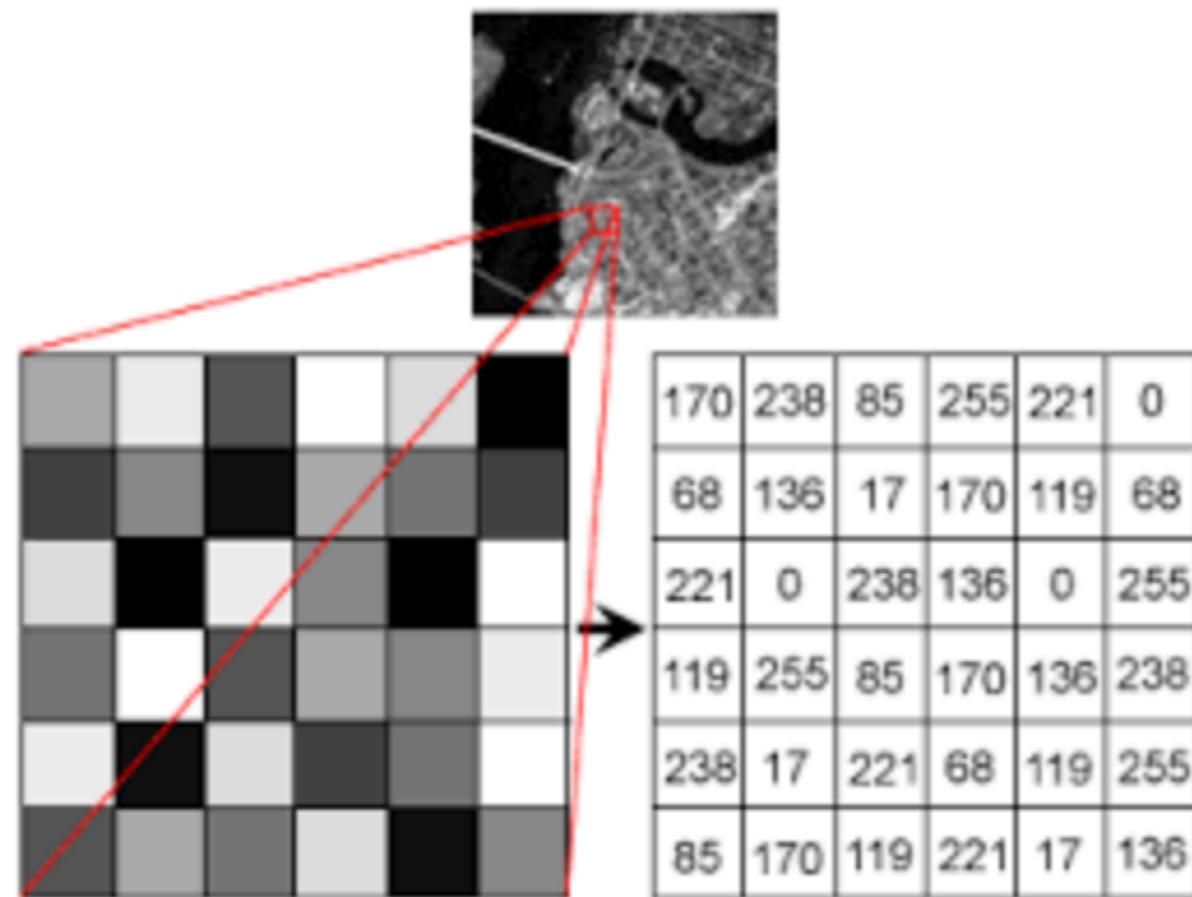


Image guided surgery  
[Grimson et al., MIT](#)

# CV Challenges

# How machines see an image?

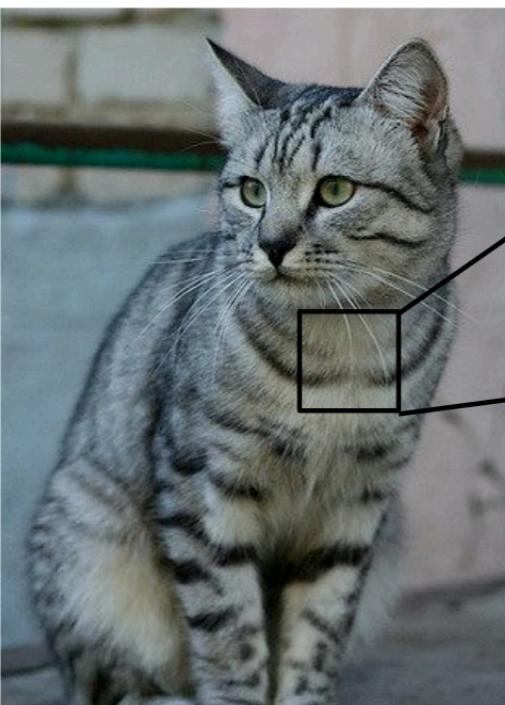
- Machines see and process everything using numbers, including images and text. How do you convert images to numbers?



# Image

- Every number represents the **pixel intensity** at that particular location. e.g., for a grayscale image where every pixel contains only one value i.e. the intensity of the black color at that location.

## The Problem: Semantic Gap



This image by Nikita is  
licensed under CC-BY 2.0

[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
[ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
[ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
[ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
[115 114 108 123 150 148 131 118 113 109 108 92 74 65 72 78]
[ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
[ 63 77 86 81 77 79 102 123 117 115 117 125 130 115 87]
[ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
[ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
[ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
[164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
[123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]

What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3  
(3 channels RGB)

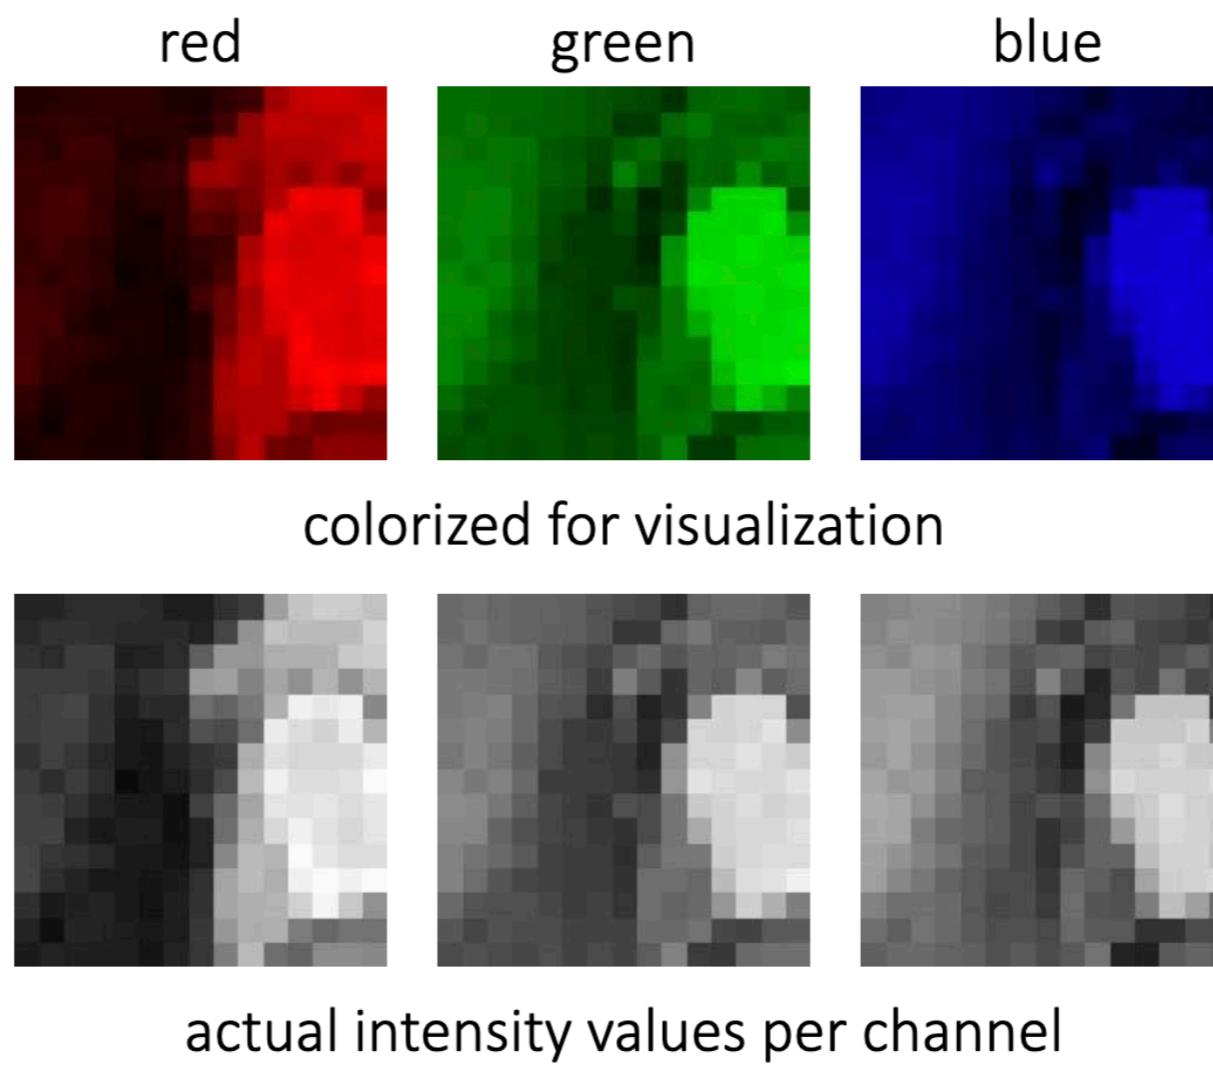
# What is an image?

- Color images will have multiple values for a single pixel. These values represent the **intensity of respective channels** – **Red, Green and Blue channels for RGB images**, for instance.



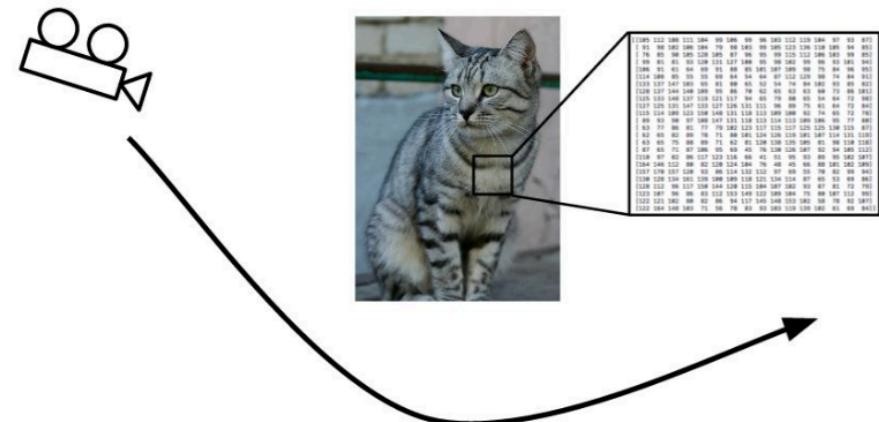
color image patch

How many bits are  
the intensity values?



# Challenges of recognition

Viewpoint

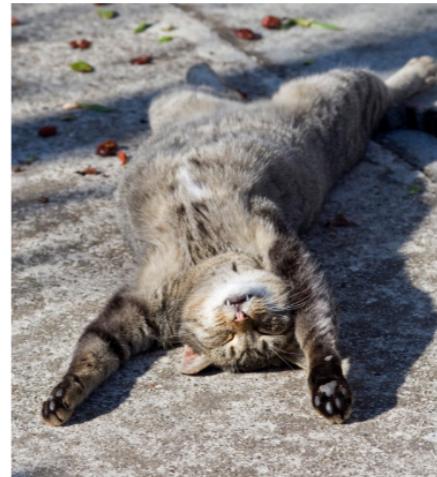


Illumination



[This image is CC0 1.0 public domain](#)

Deformation



[This image by Umberto Salvagnin is licensed under CC-BY 2.0](#)

Occlusion



[This image by jonsson is licensed under CC-BY 2.0](#)

Clutter



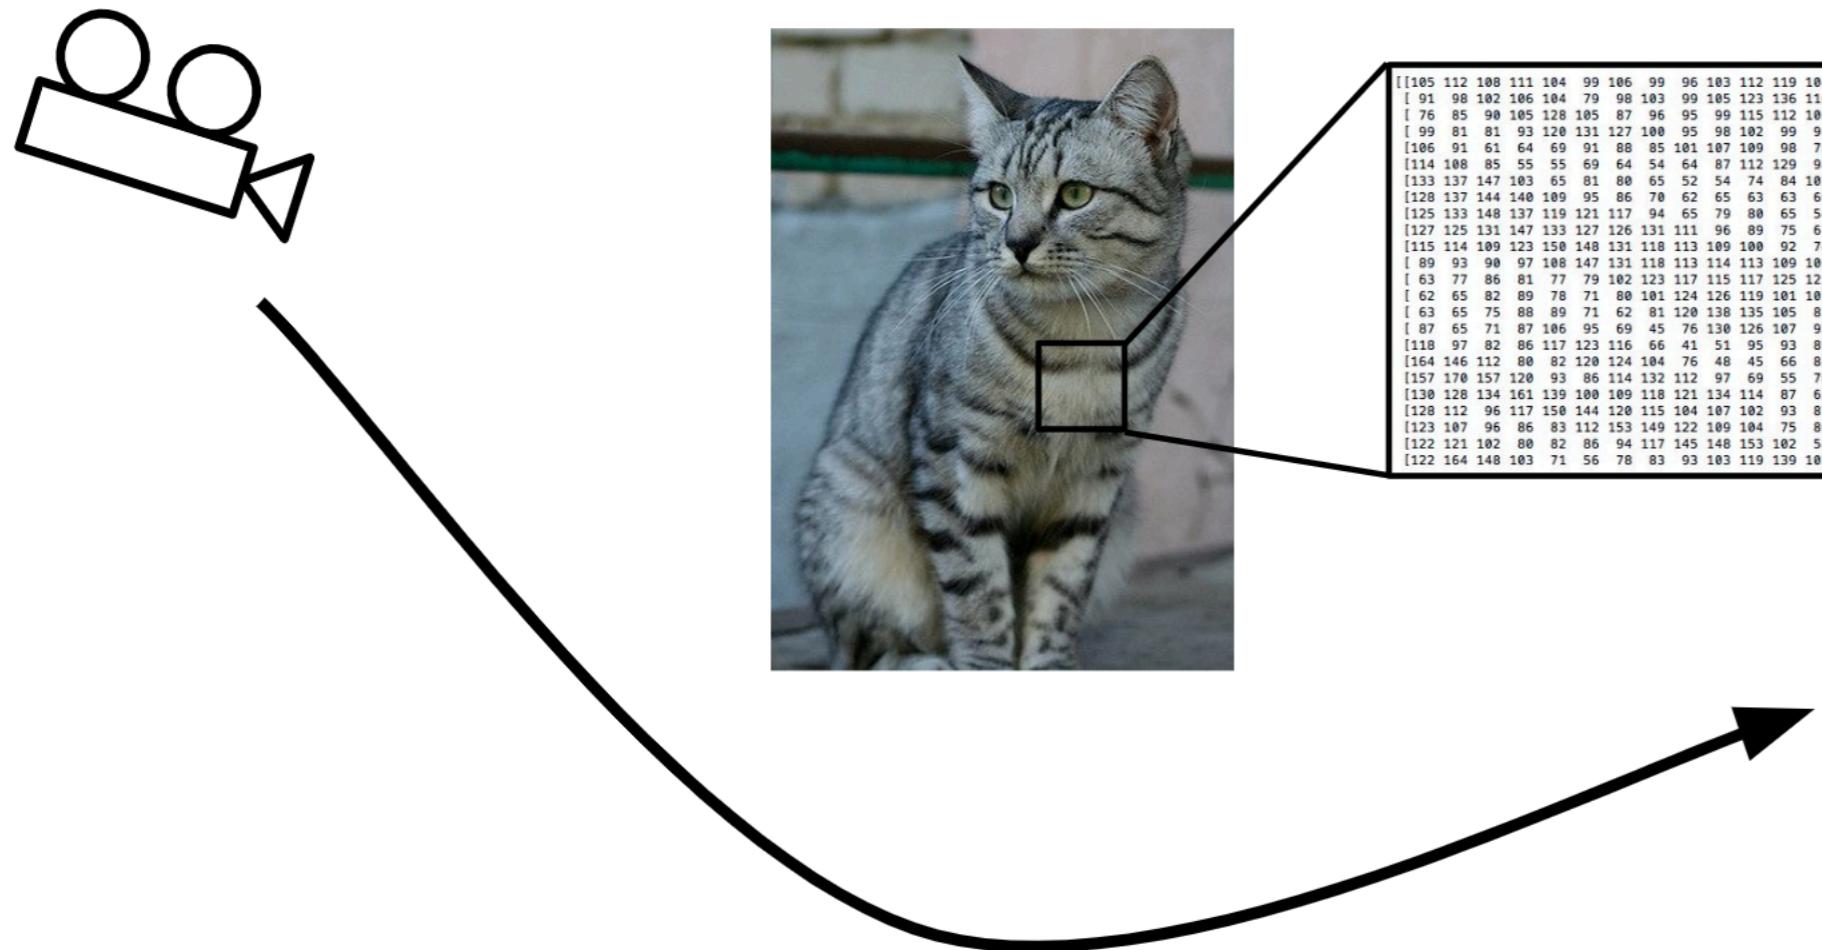
[This image is CC0 1.0 public domain](#)

Intraclass Variation



[This image is CC0 1.0 public domain](#)

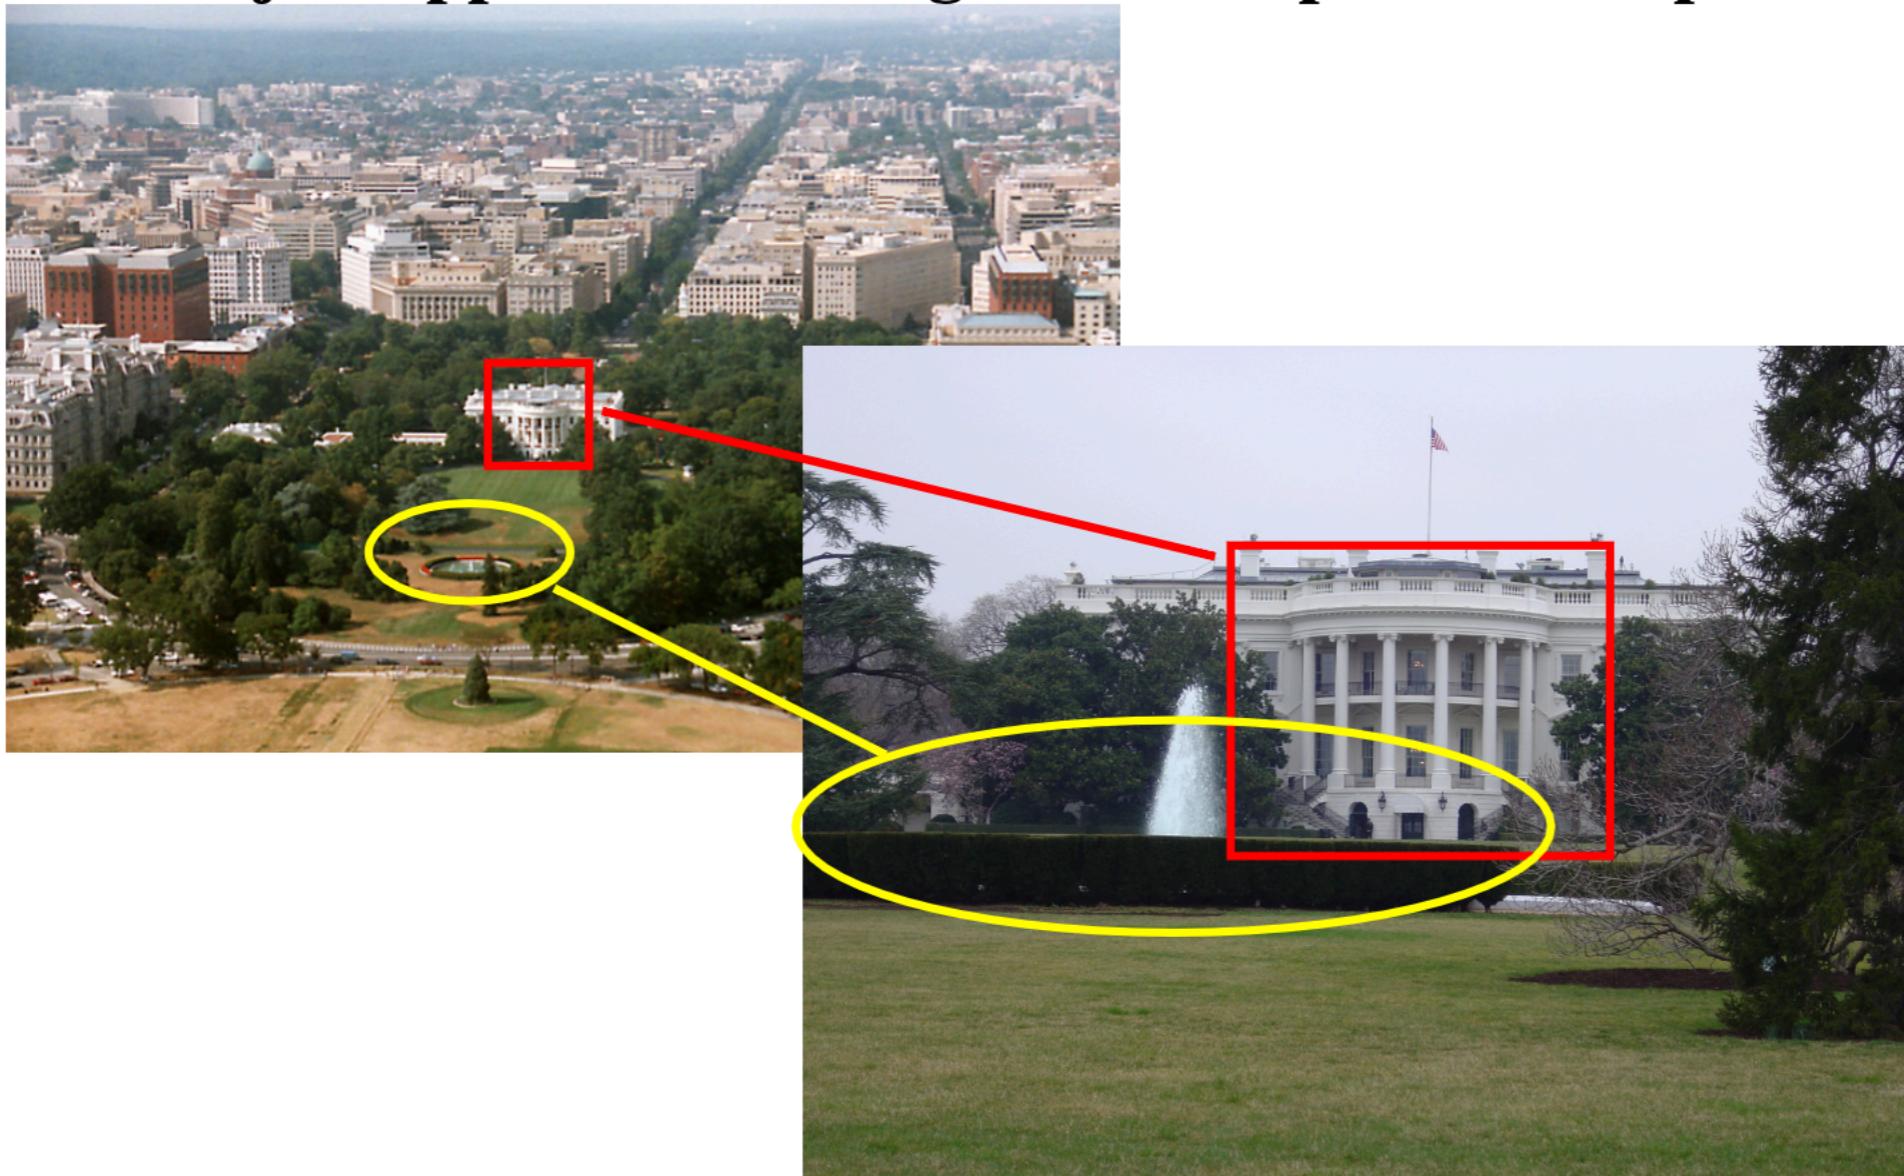
# Challenge: Viewpoint variation



All pixels change when  
the camera moves!

# Challenge: Viewpoint variation

Object appearance changes with respect to viewpoint



# Challenge: Illumination



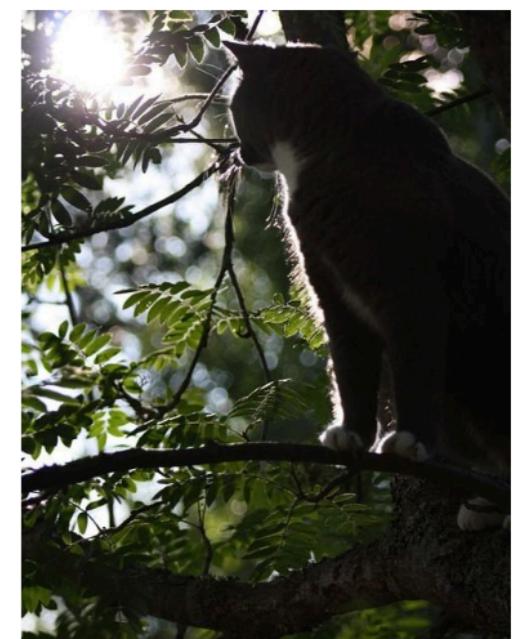
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



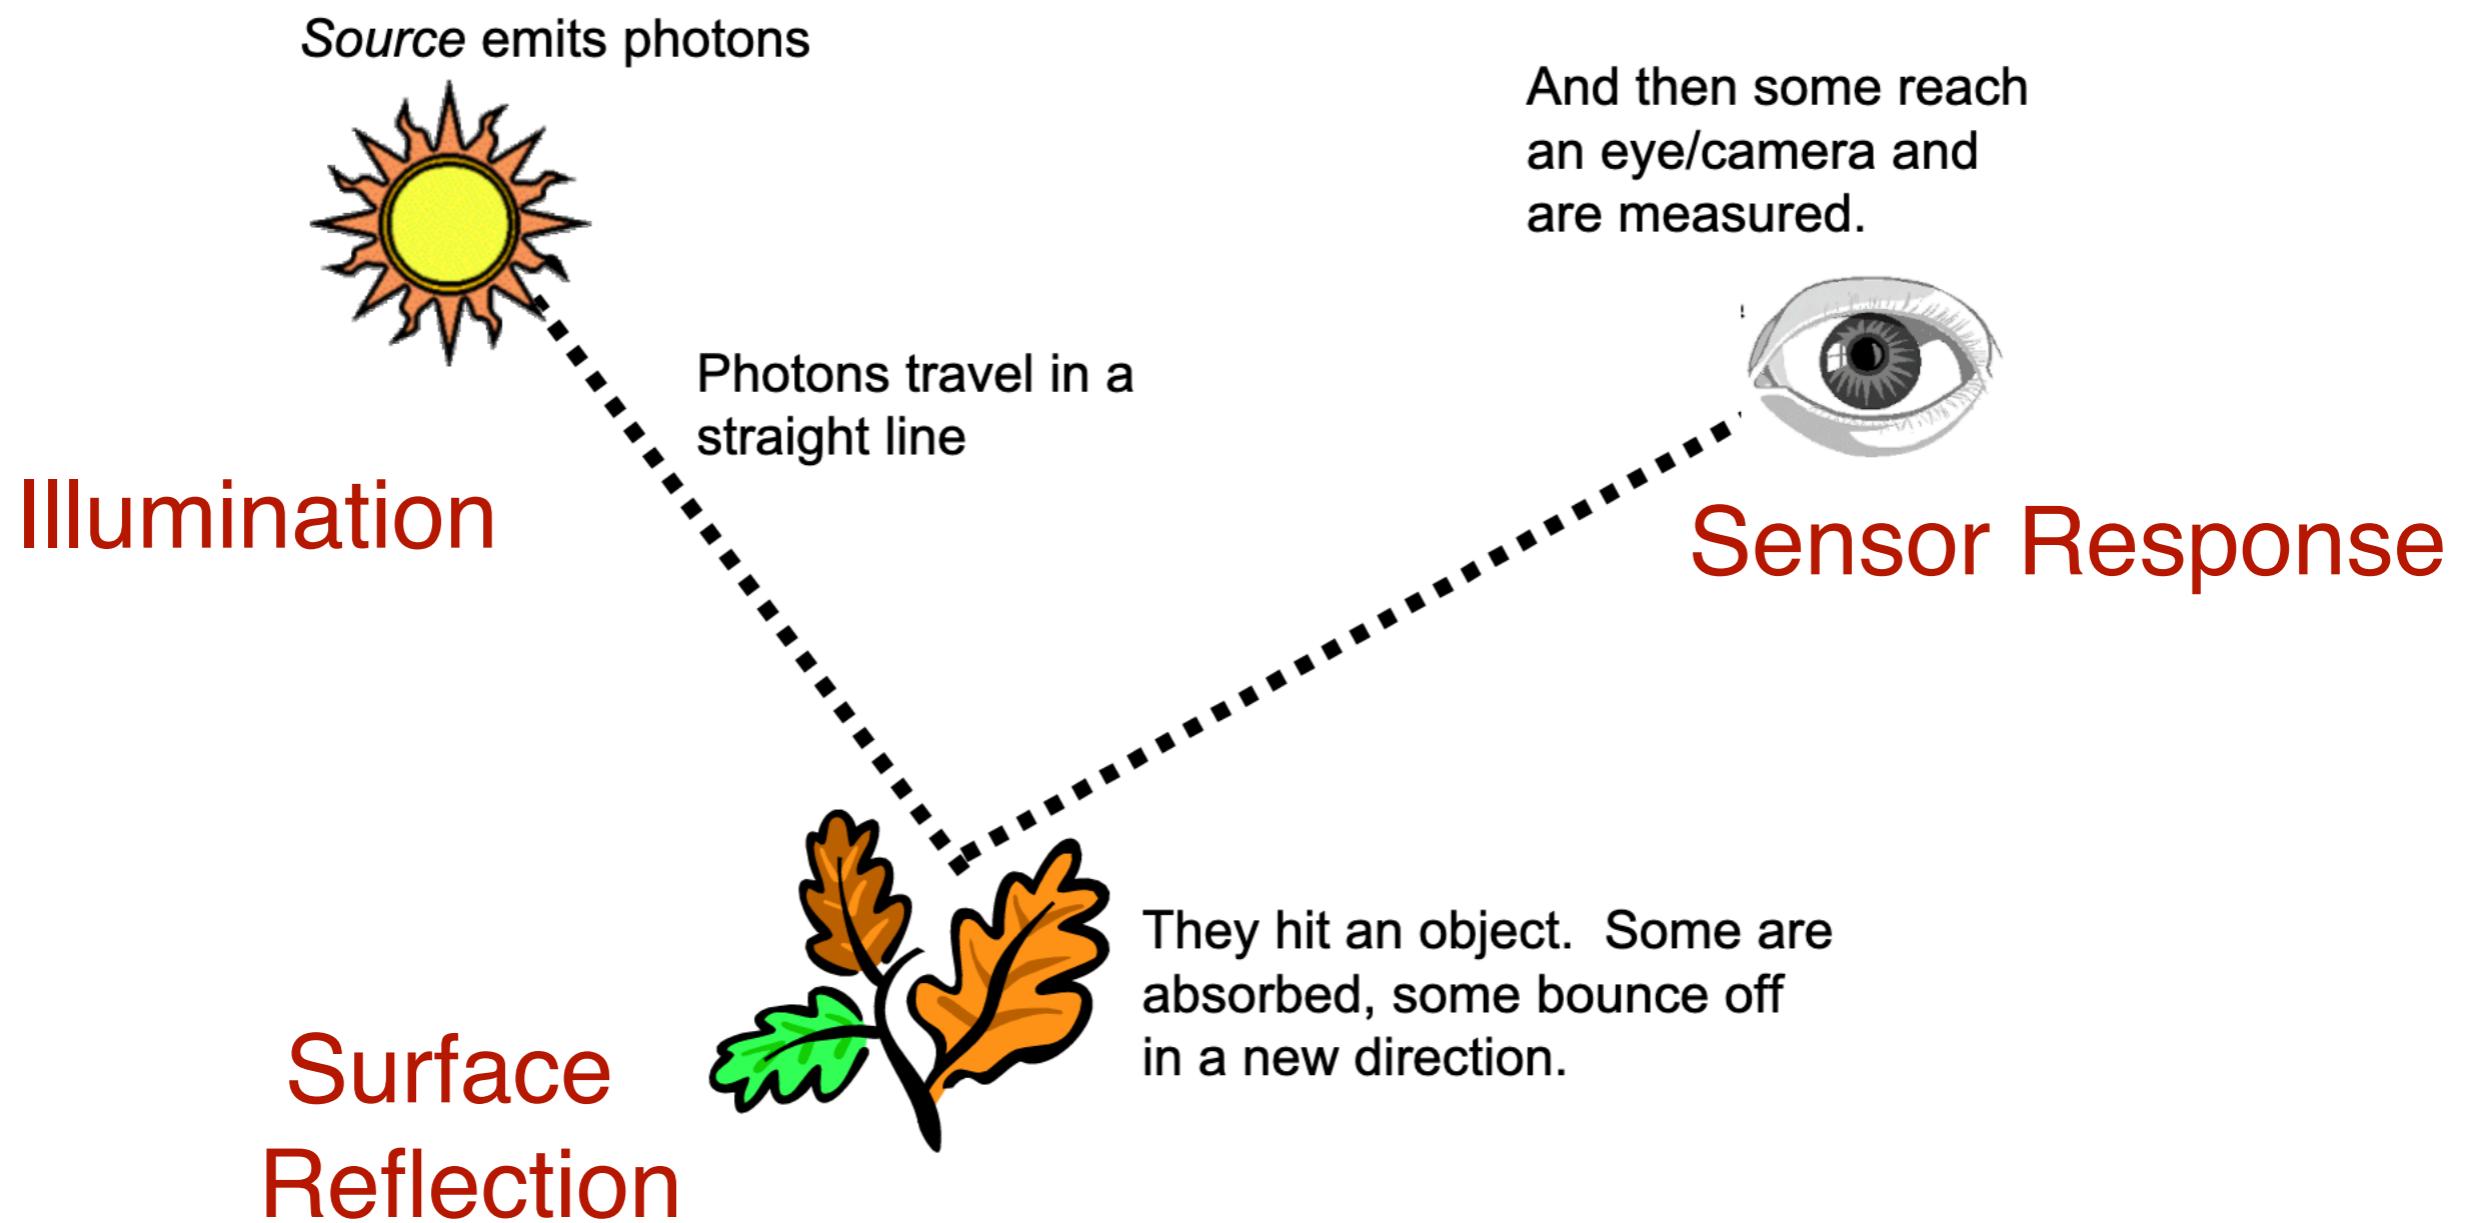
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

**Object appearance changes with respect to lighting magnitude and direction.**

# What is Color?



# Color

- Color percepts are a composition of three factors (**illumination, surface reflectance, sensor response**)
- We can't easily factor the color we see in the image to infer illumination and material (even if sensor properties are fixed and known).

## Some things to think about:

- “red” typically means “appears red to a human observer under white light”.
- white objects appear red under a red light.
- nothing looks red if you are red/green color blind.



“normal” color perception



red/green color blind

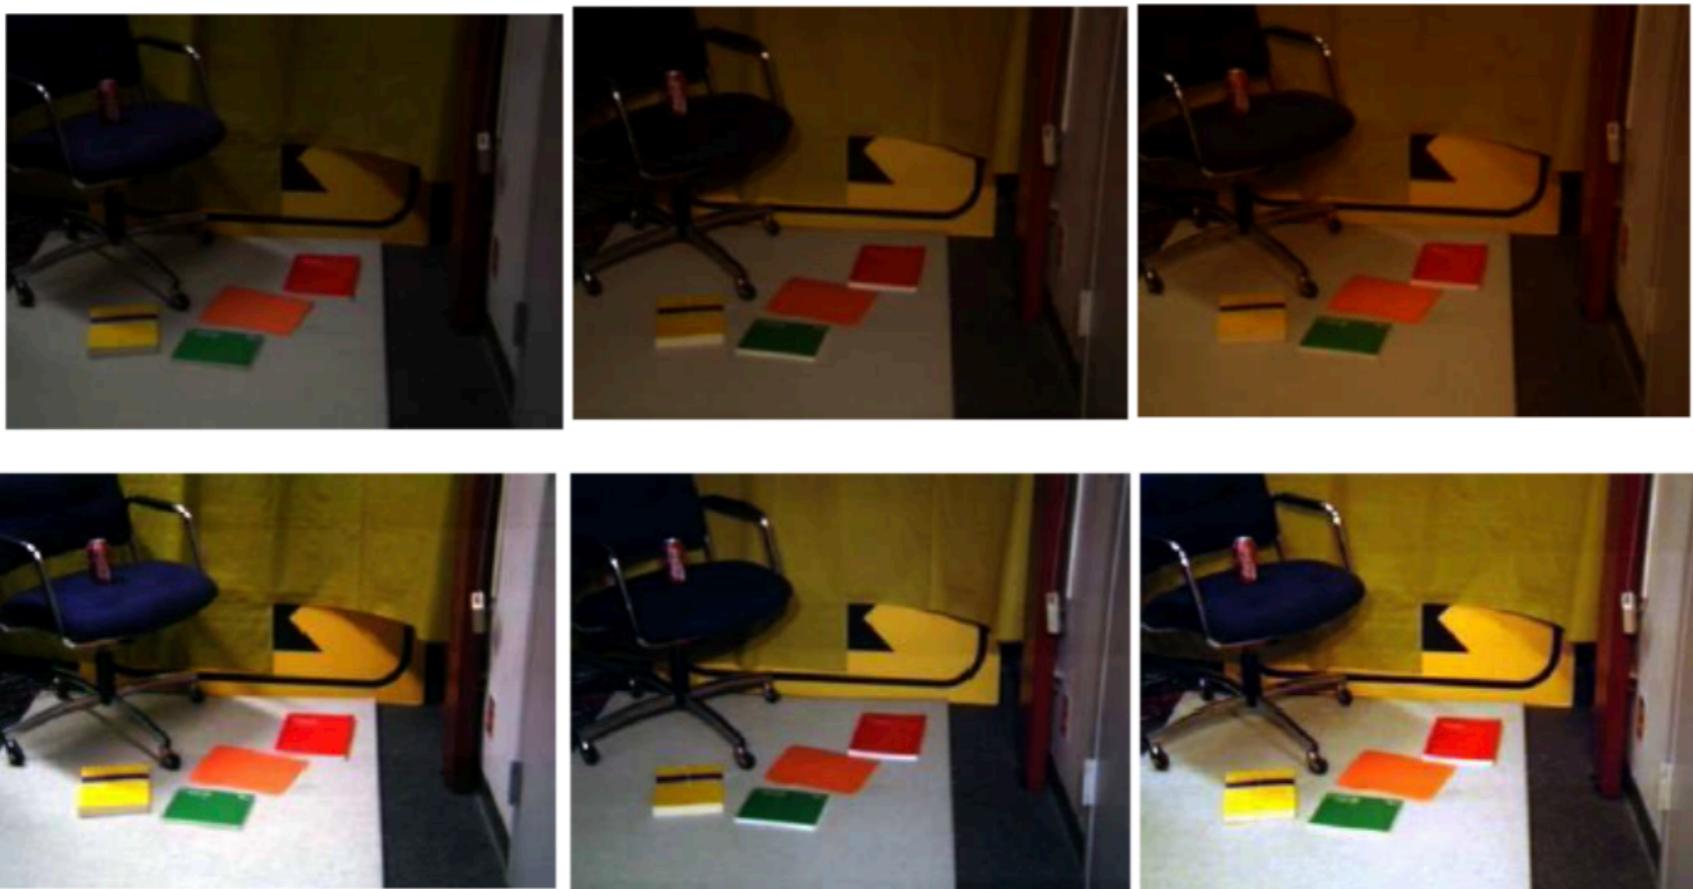
# Is The Dress Blue and Black or White and Gold?



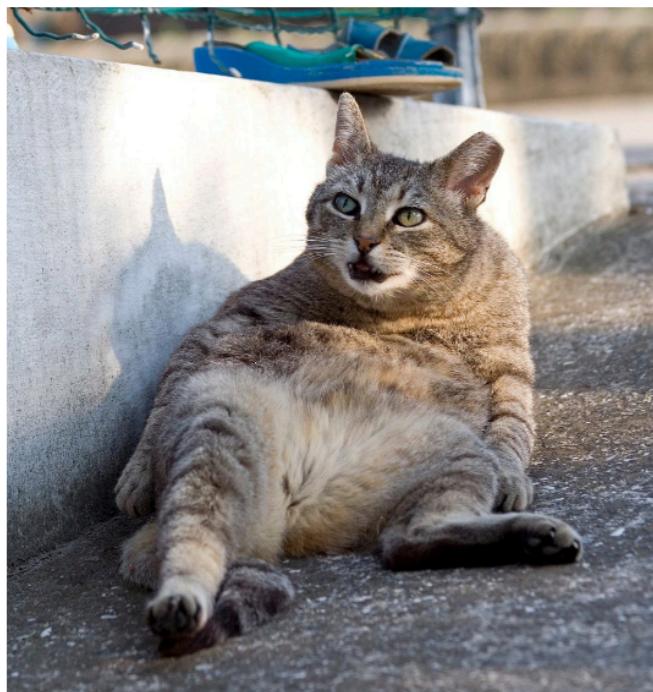
This dress manages to simultaneously gather more than 670,000 people on [Buzzfeed](#), and convince 900,000 visitors to take a poll.

# Challenge: Color Constancy

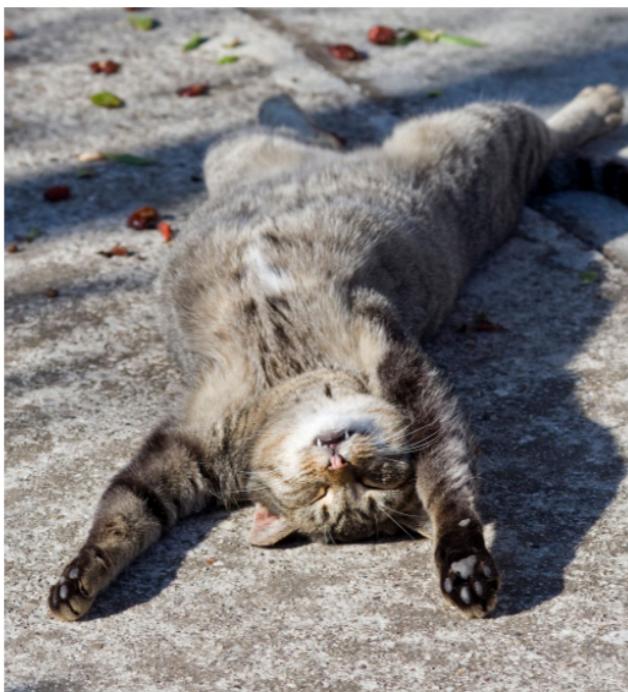
Humans are very good at recognizing the same material colors under different illumination. Not clear how this is achieved in the general case.



# Challenge: Deformation



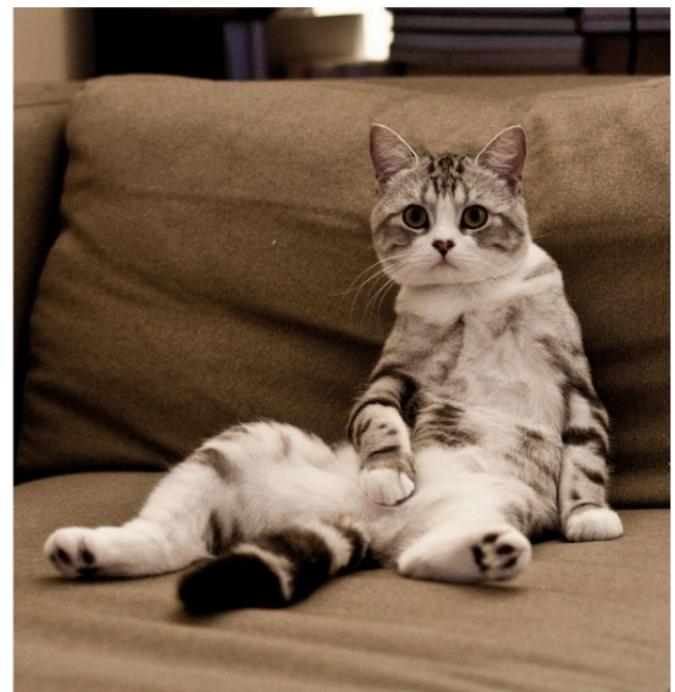
[This image by Umberto Salvagnin](#)  
is licensed under [CC-BY 2.0](#)



[This image by Umberto Salvagnin](#)  
is licensed under [CC-BY 2.0](#)



[This image by sare bear](#) is  
licensed under [CC-BY 2.0](#)



[This image by Tom Thai](#) is  
licensed under [CC-BY 2.0](#)

# Challenge: Occlusion



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image by jonsson is licensed under CC-BY 2.0](#)

# Challenge: Variation



This image is [CC0 1.0 public domain](#)

# Distance Metric to compare images

**L1 distance:**

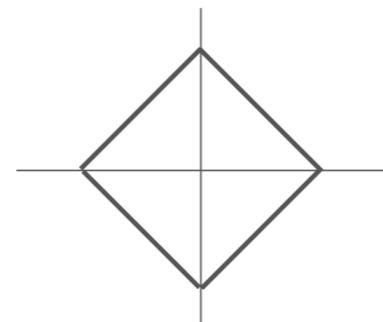
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

add → 456

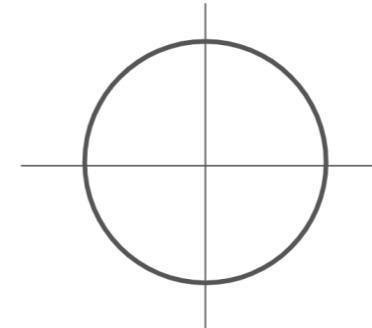
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



# Distance metrics on pixels



[Original image is  
CC0 public domain](#)

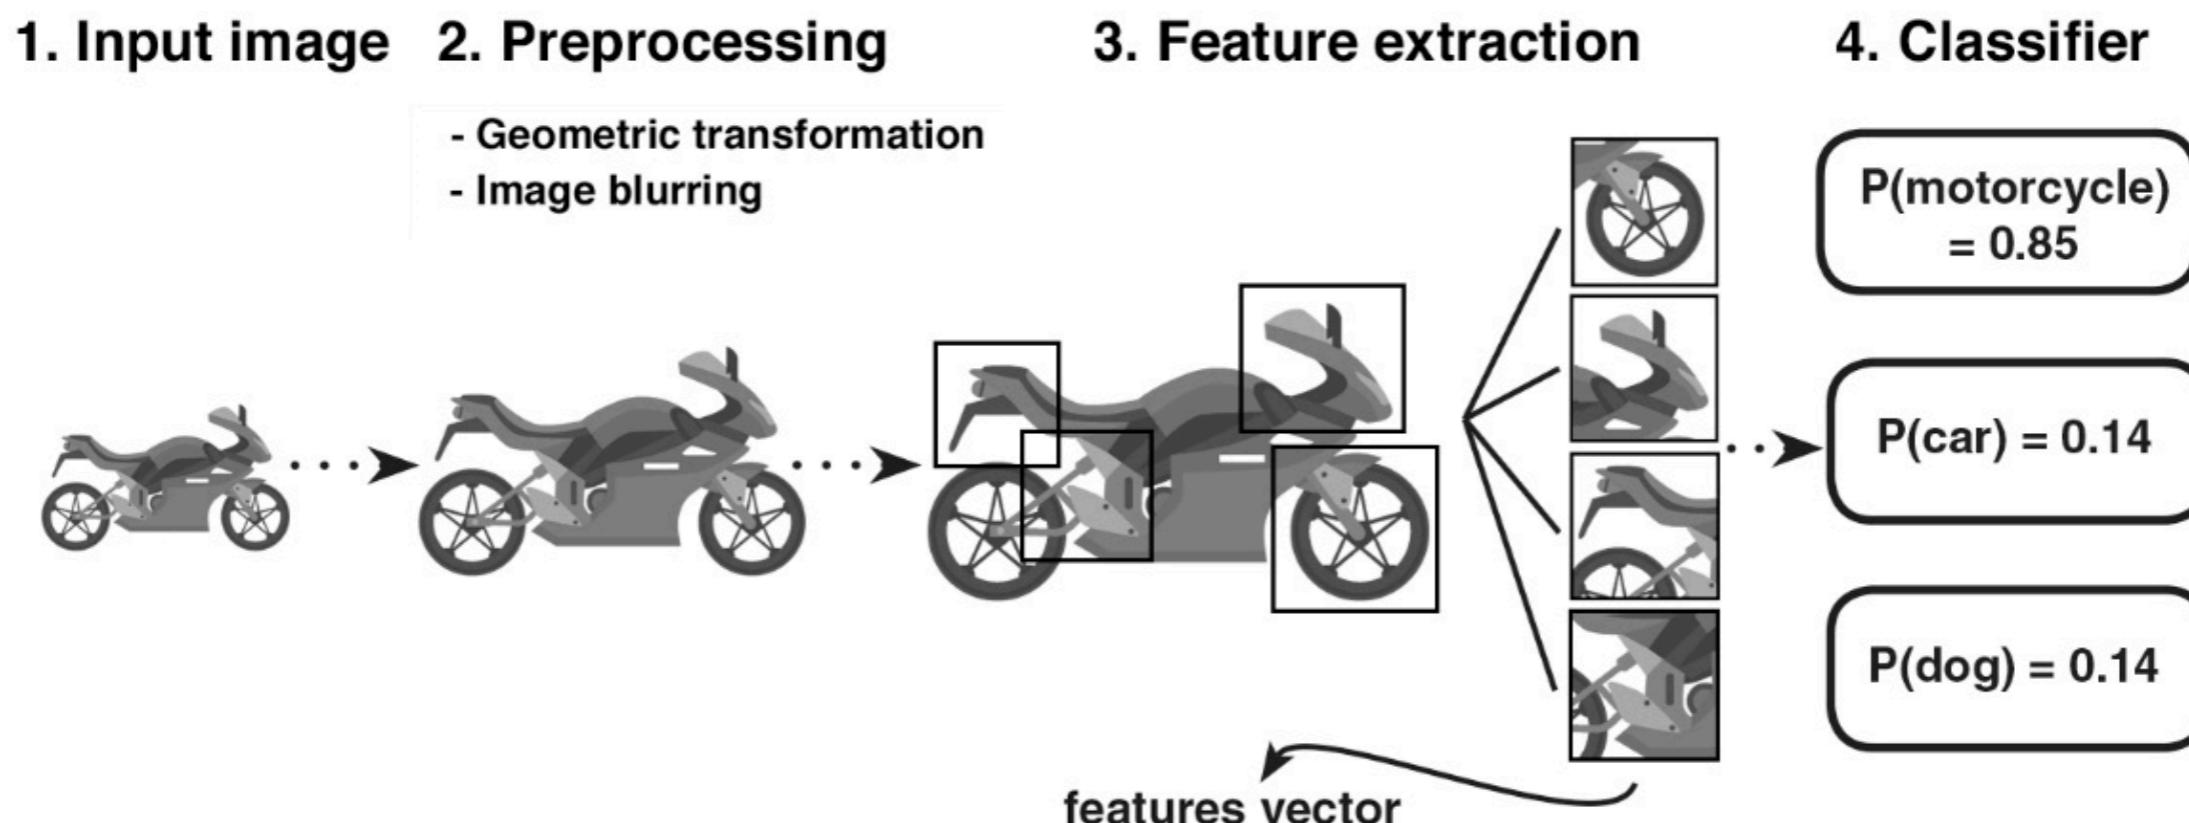
(all 3 images have same L2 distance to the one on the left)

**k-Nearest Neighbor on images **never used**.**

- Very slow at test time
- Distance metrics on pixels are not informative

# CV main Operations

# CV pipeline



# CV main Operations

## Filtering and Smoothing



Linear operators  
Convolution  
Smoothing

# CV main Operations

## Feature Extraction

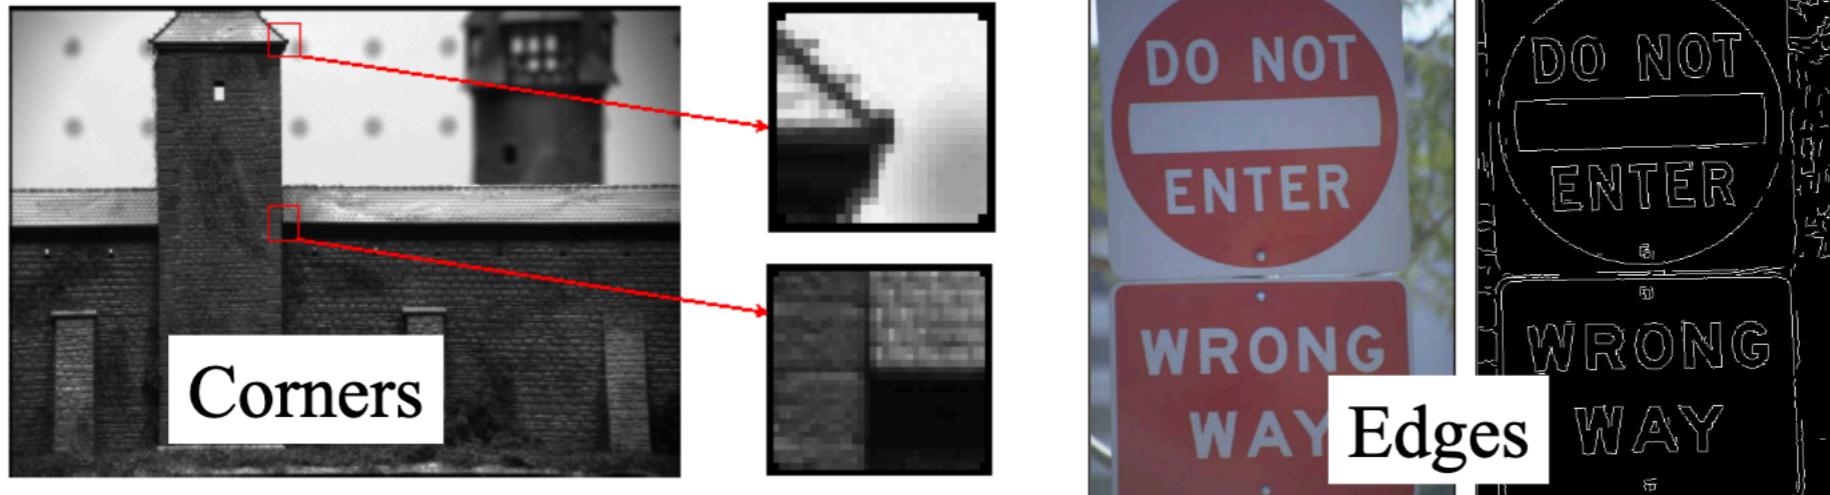


Image derivatives  
Gradient operators  
DoG/LoG operators  
Harris corner detector

**Why?**  
**Seek more unique descriptors  
(than pixels) for matching**

# CV main Operations

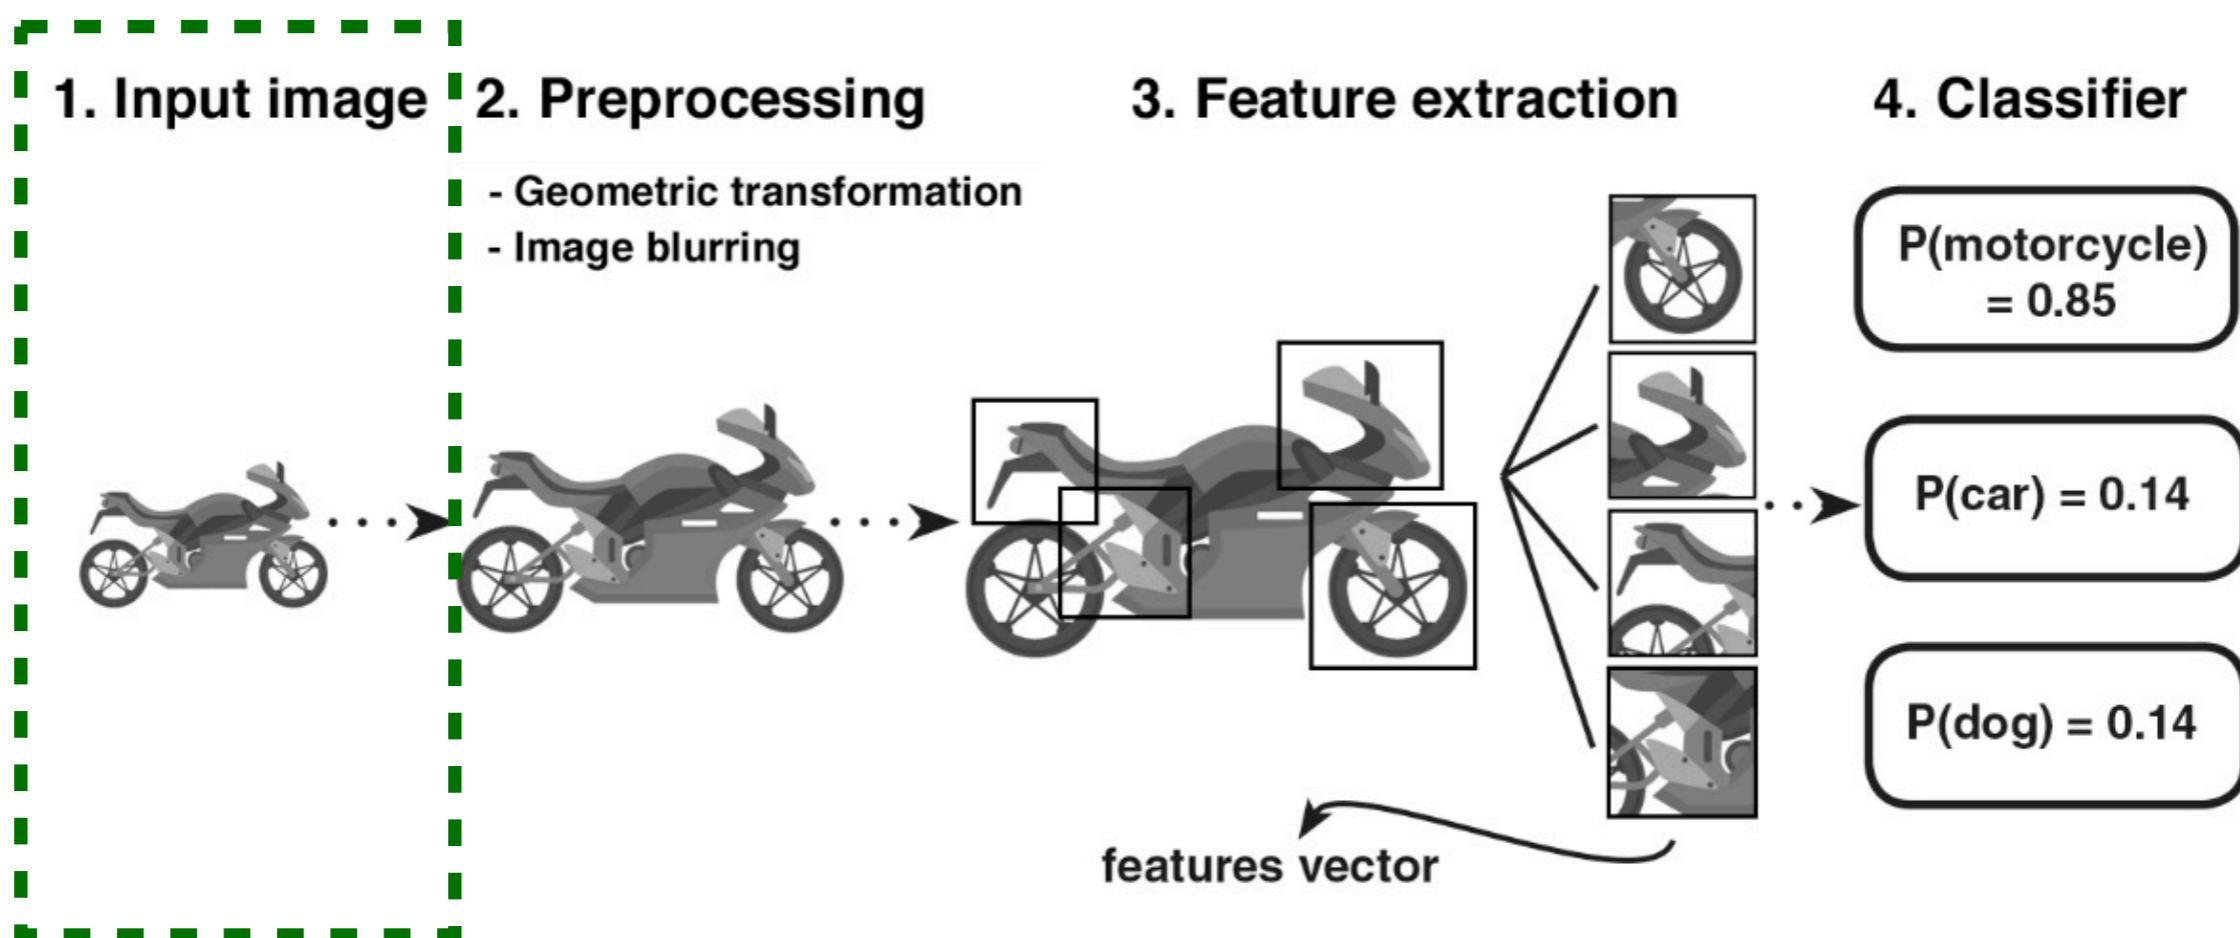
## Color and Light



Radiance / Reflection  
Illumination / Shading  
Chromaticity  
Color Constancy



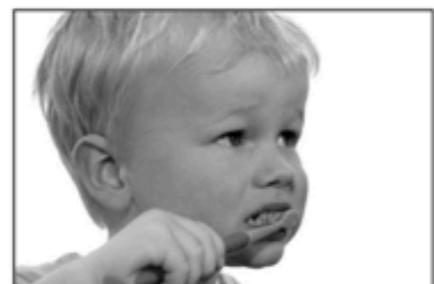
# CV Pipeline



# Images as functions

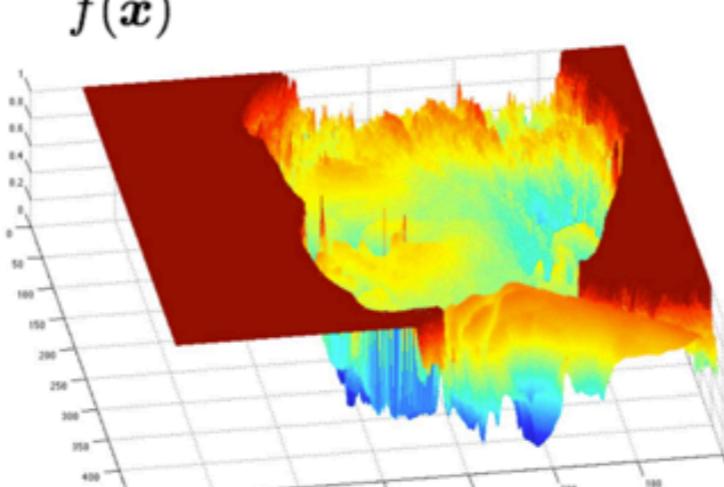
- An Image as a function  $f$  from  $\mathbb{R}^2$  to  $\mathbb{R}^M$ :
  - $f(x, y)$  gives the **intensity** at position  $(x, y)$
  - Defined over a rectangle, with a finite range:

$$f: \underbrace{[a,b] \times [c,d]}_{\text{Domain support}} \rightarrow \underbrace{[0,255]}_{\text{range}}$$



grayscale image

What is the range of  
the image function  $f$ ?



$$\text{domain } \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

A (grayscale)  
image is a 2D  
function.

# Images as functions

- An **Image** as a function  $f$  from  $\mathbb{R}^2$  to  $\mathbb{R}^M$ :
  - $f(x, y)$  gives the **intensity** at position  $(x, y)$
  - Defined over a rectangle, with a finite range:

$$f: \underbrace{[a,b] \times [c,d]}_{\text{Domain support}} \rightarrow \underbrace{[0,255]}_{\text{range}}$$

- A color image:  $f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$

# Input Image

```
1 #import the libraries
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import cv2
5 %matplotlib inline
6
7 #reading the image
8
9 image = cv2.imread('index.png')
10 image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
11 #plotting the image
12 plt.imshow(image)
13
14 #saving image
15 cv2.imwrite('test_write.jpg',image)
```

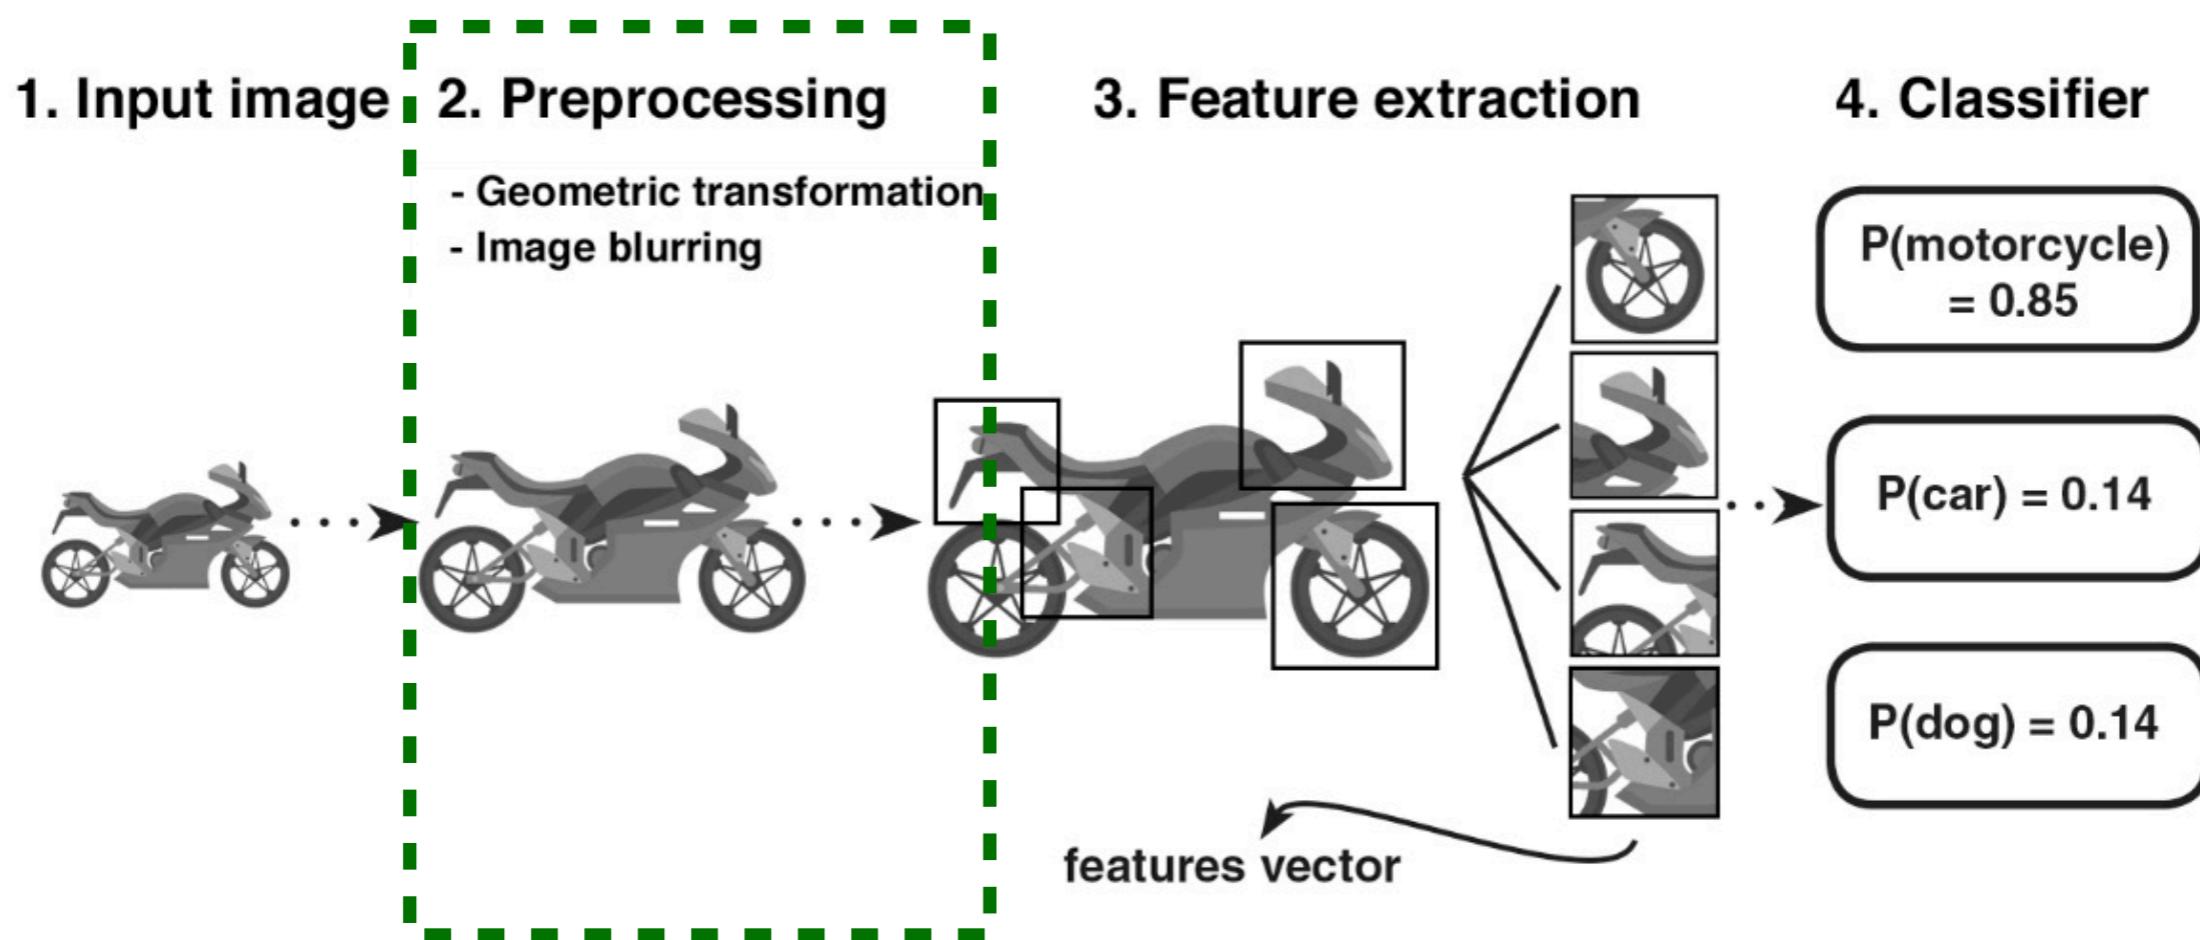
- By default, the *imread* function reads images in the BGR (Blue-Green-Red) format. We can read images in different formats using extra flags in the *imread* function:

**cv2.IMREAD\_COLOR:** Default flag for loading a color image

**cv2.IMREAD\_GRAYSCALE:** Loads images in grayscale format

**cv2.IMREAD\_UNCHANGED:** Loads images in their given format, including the alpha channel. Alpha channel stores the transparency information – the higher the value of alpha channel, the more opaque is the pixel

# CV Pipeline

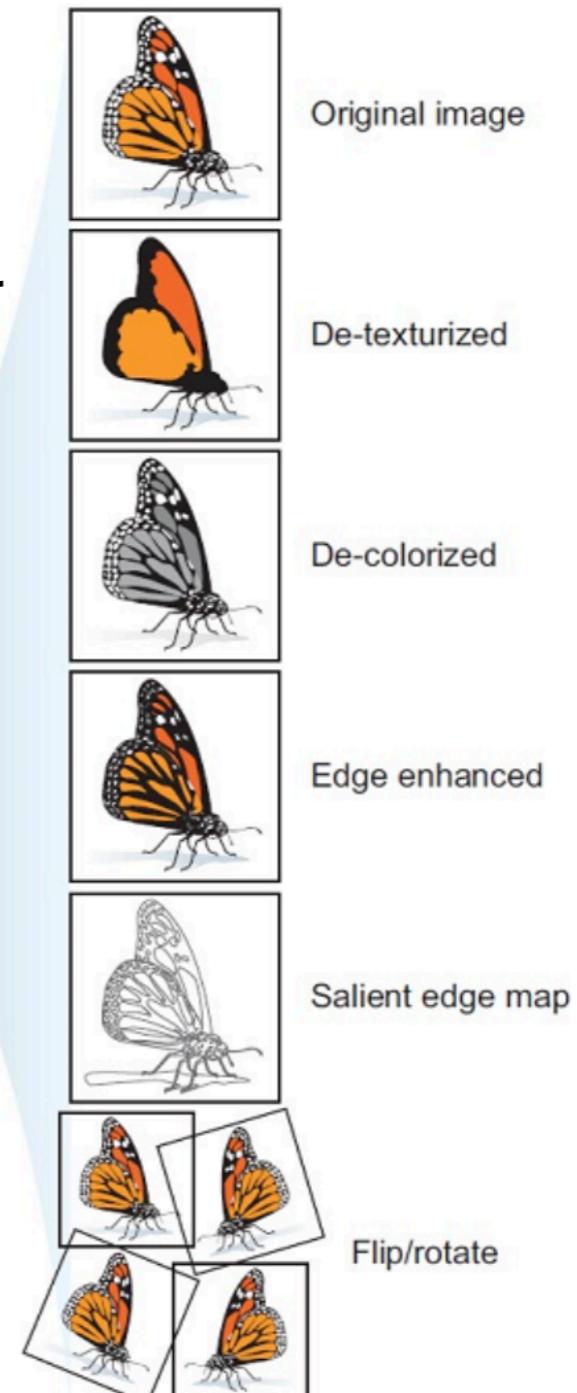


# Image Augmentation

- Data augmentation uses the available data samples to produce the new ones, by applying image operations like **rotation**, **scaling**, **translation**, etc. This makes our model robust to changes in input and leads to better generalization.



Data augmentation



# Image transformations



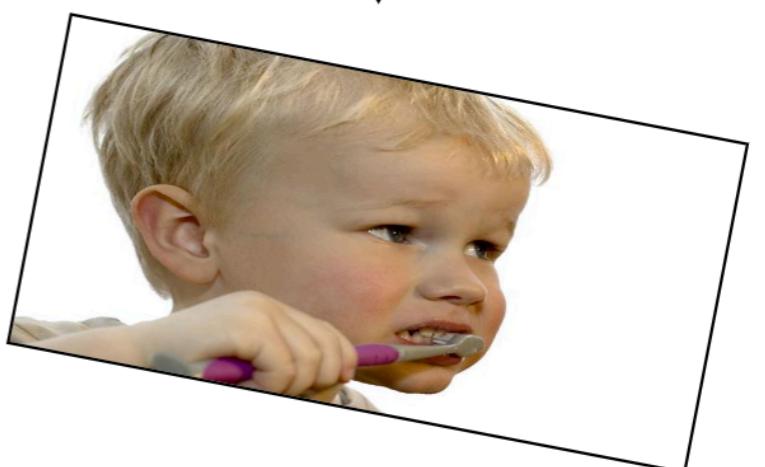
Filtering



changes pixel *values*

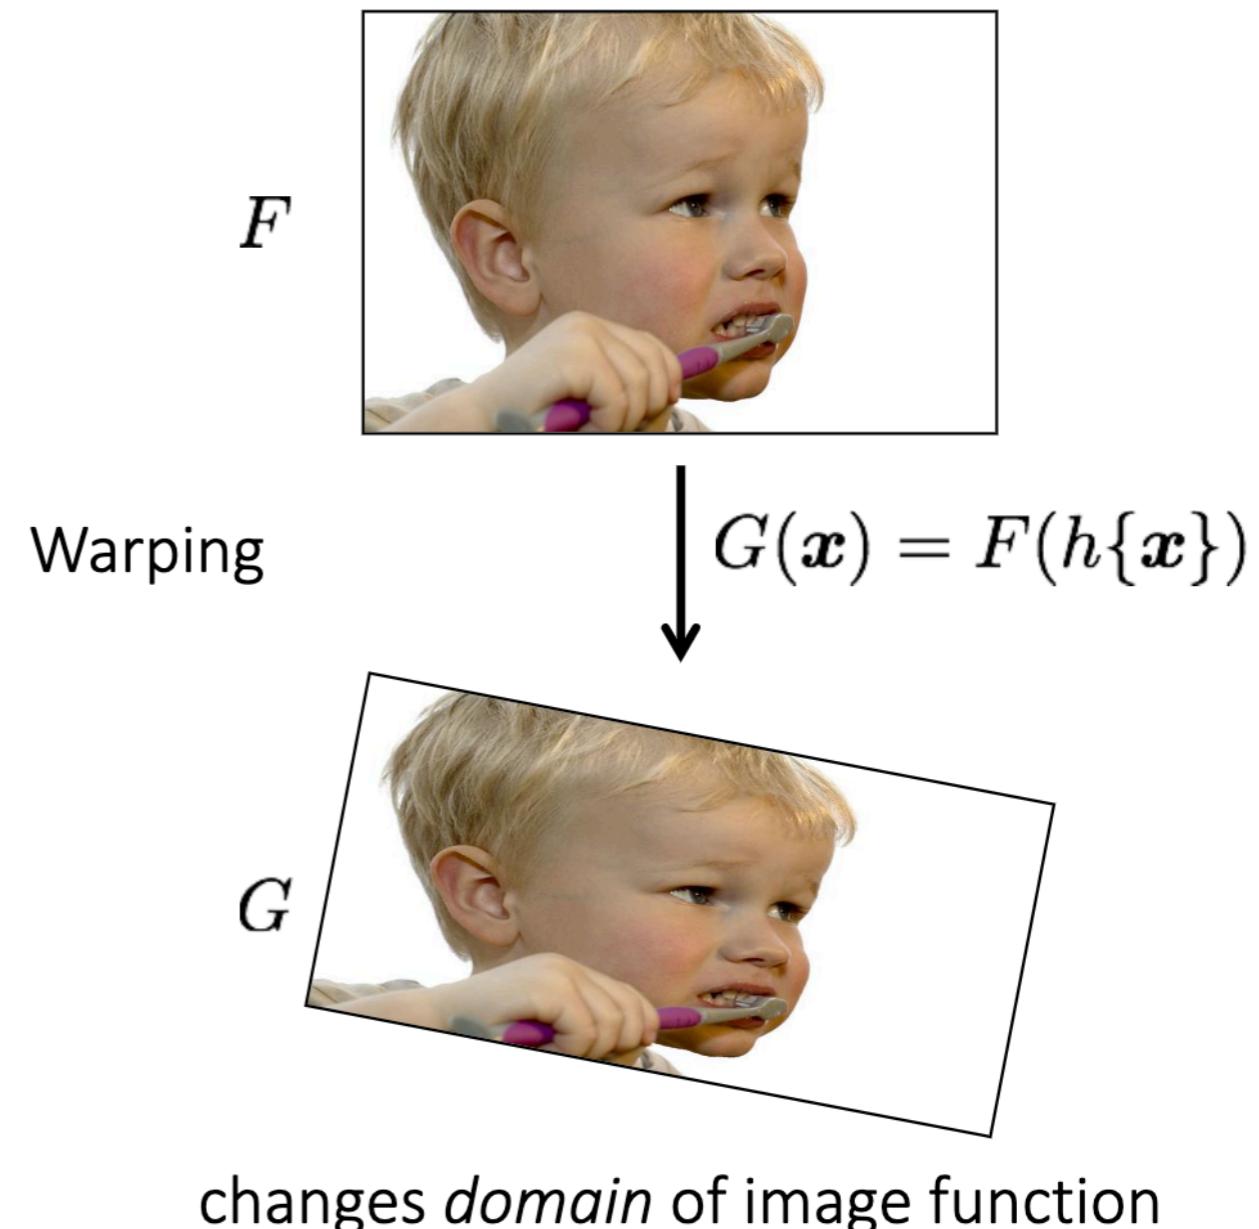
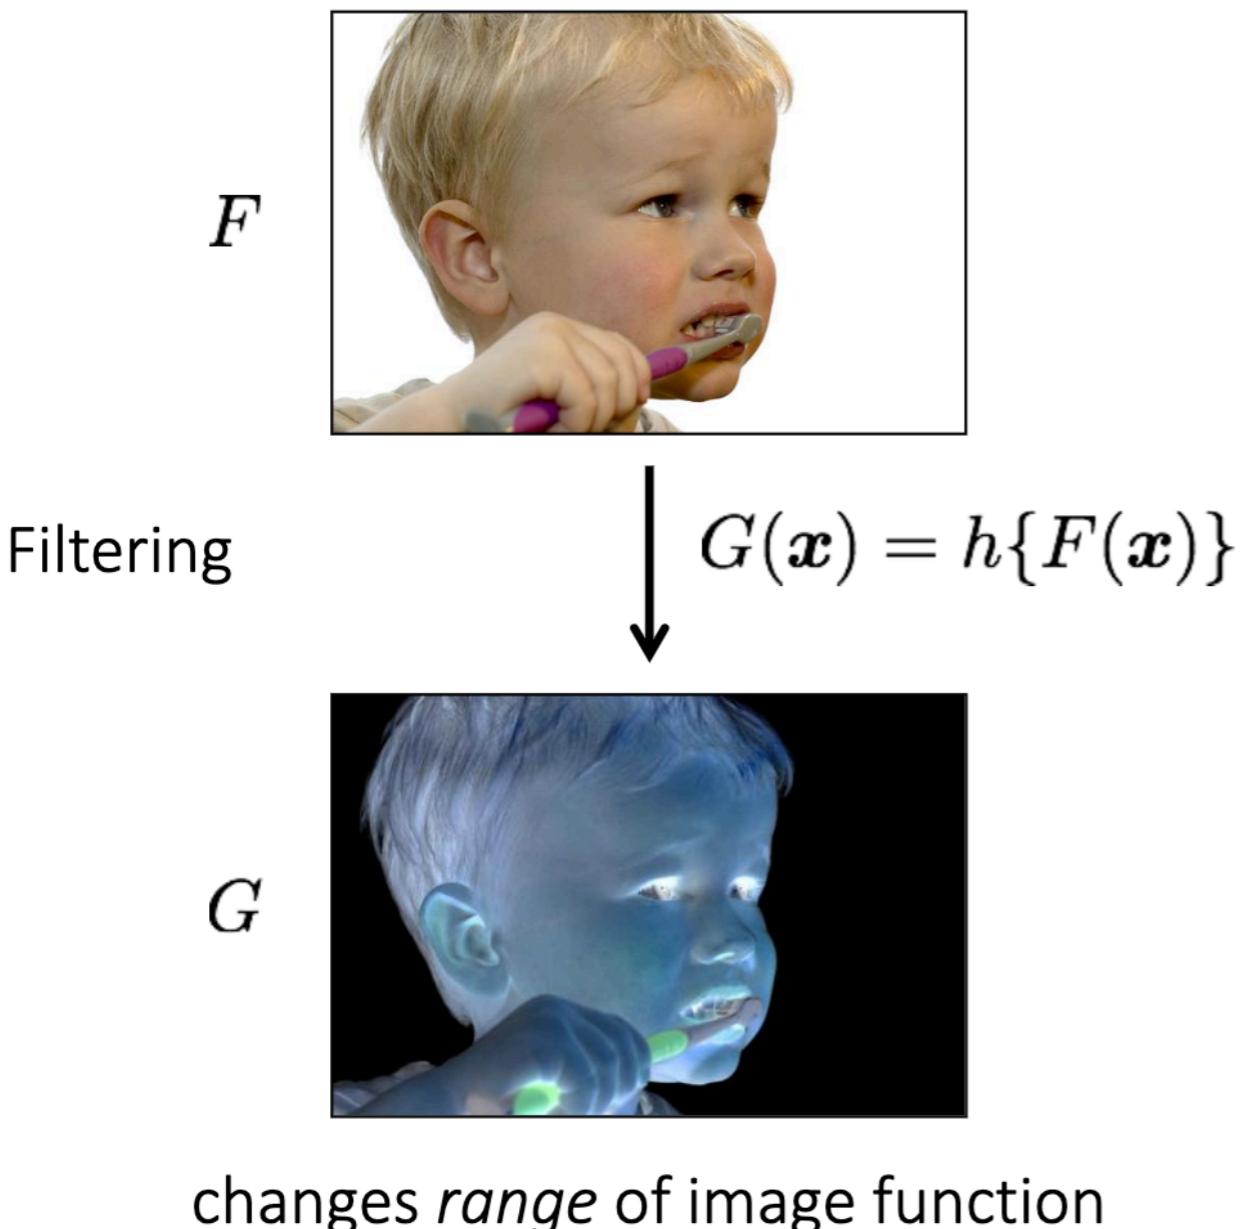


Warping

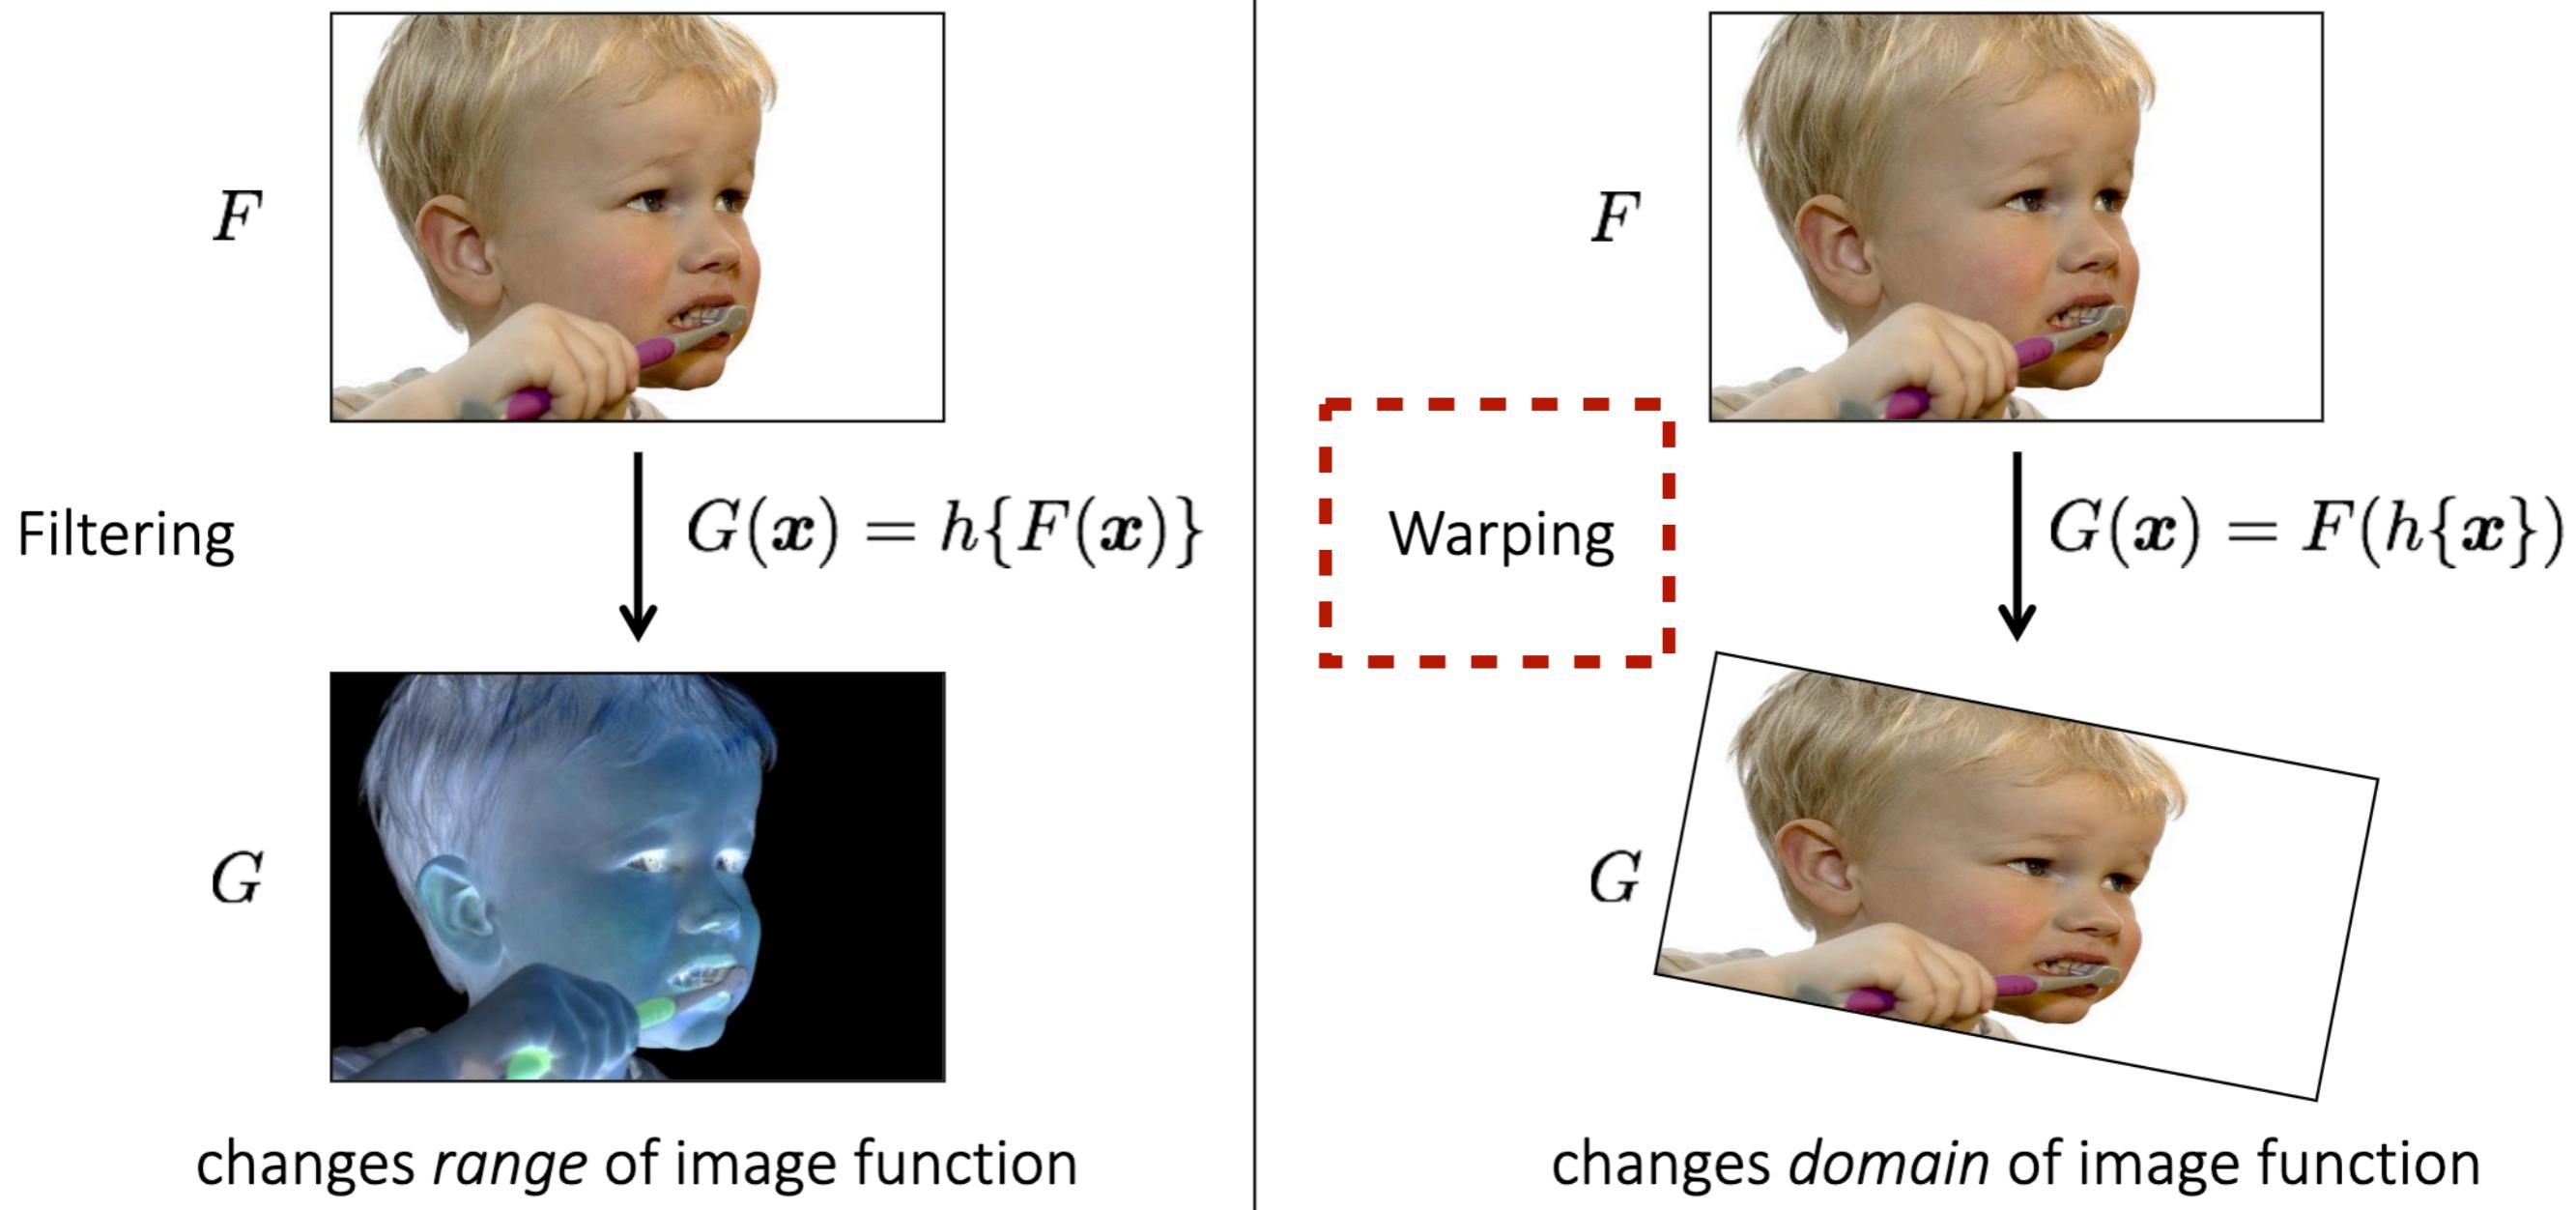


changes pixel *locations*

# Image transformations

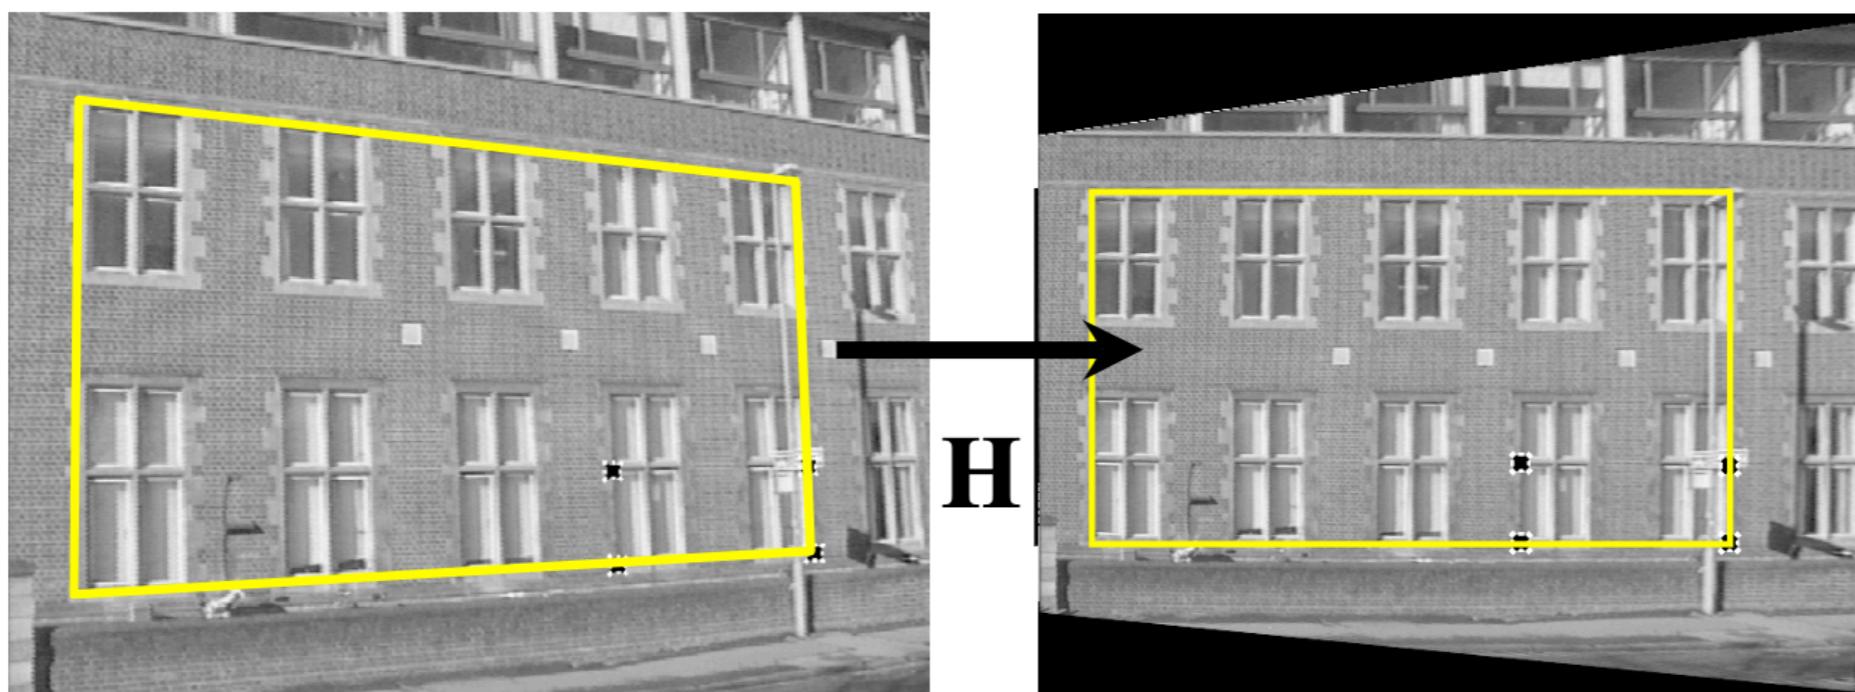


# Image transformations



# Transformation: Warping

Transformation in this case is a projective transformation  
(general  $3 \times 3$  matrix, operating on homogeneous coords)

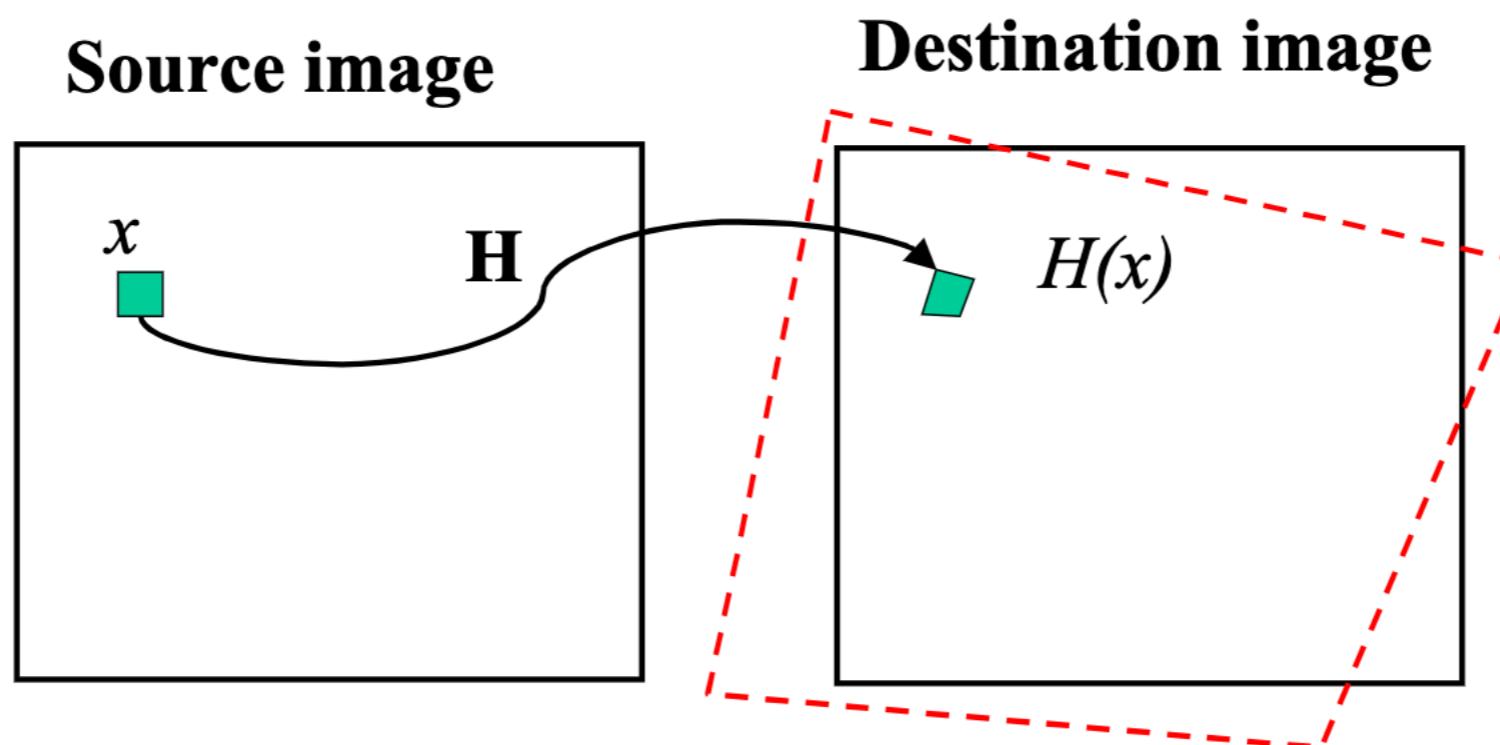


from Hartley & Zisserman

**Source Image**

**Destination image**

# Forward Warping

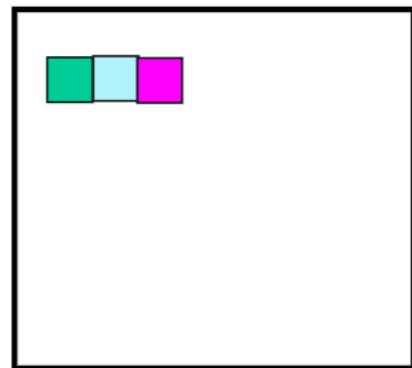


- For each pixel  $x$  in the source image
- Determine where it goes as  $H(x)$
- Color the destination pixel

**Problems?**

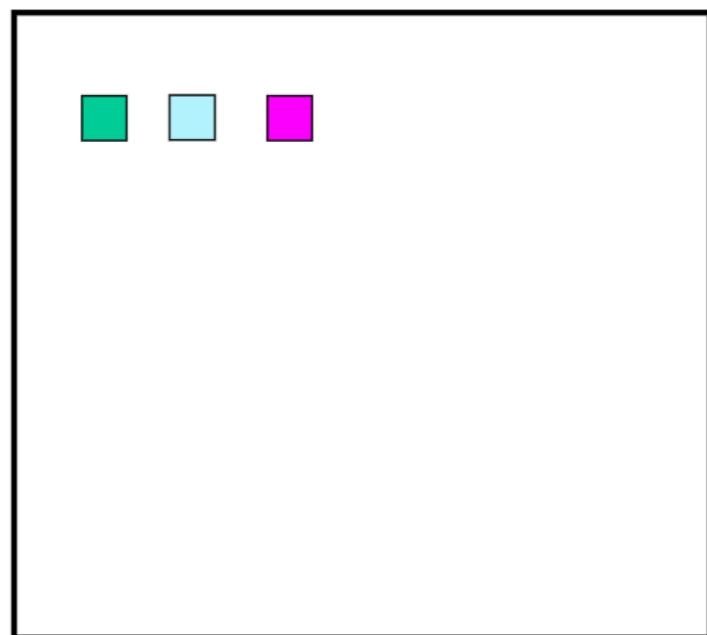
# Forward Warping (resizing)

**Source image**



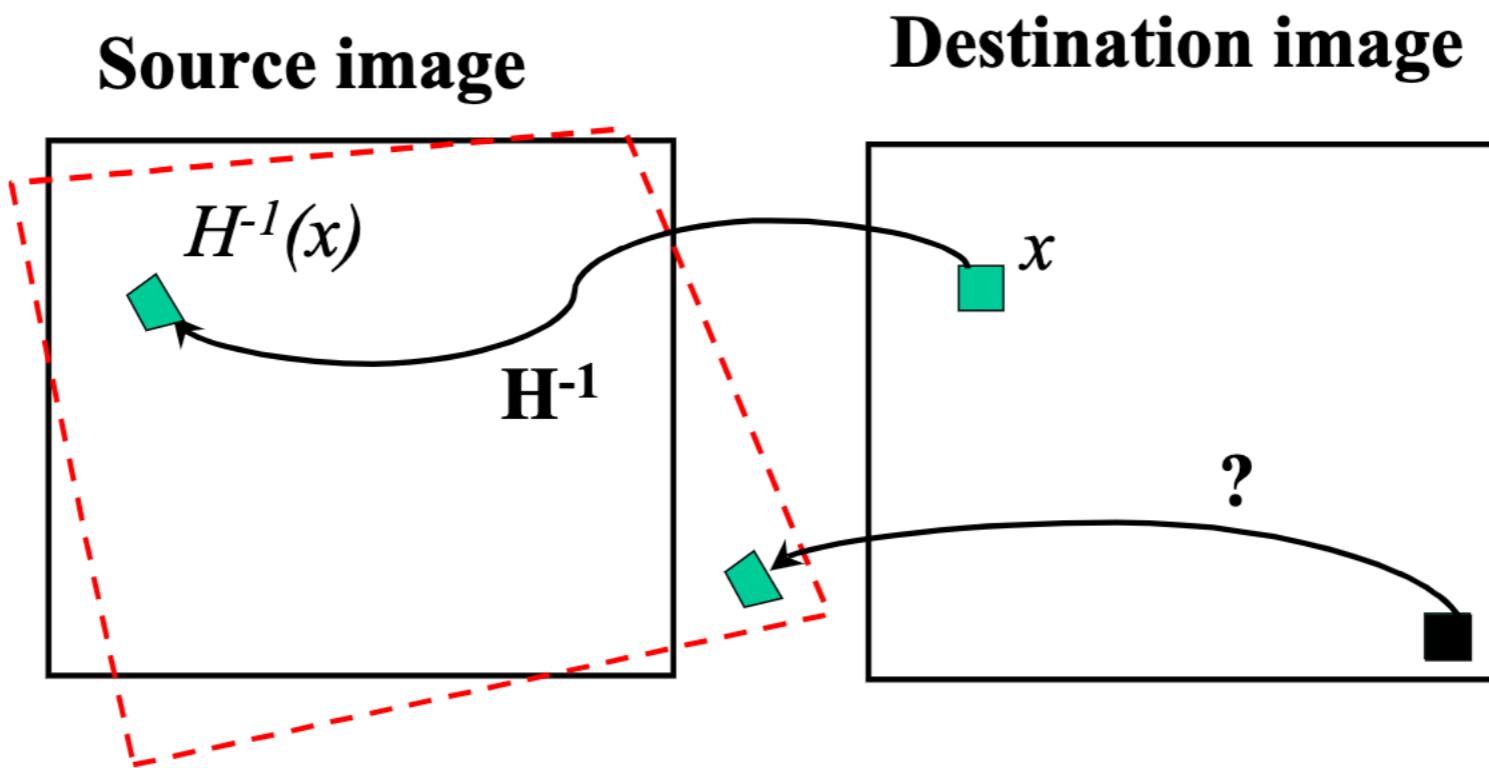
→  
**magnified**

**Destination image**



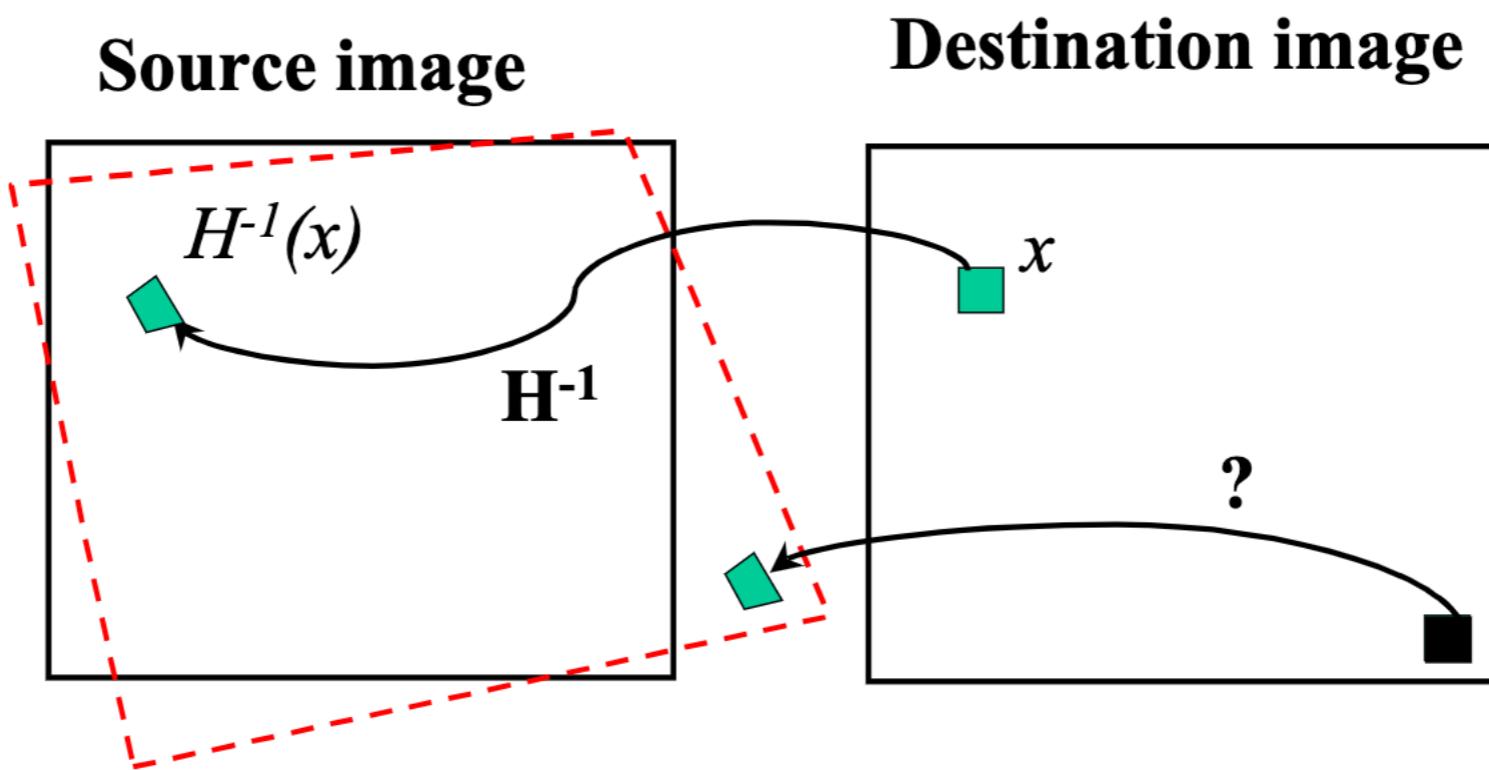
**Can leave gaps!**

# Backward Warping



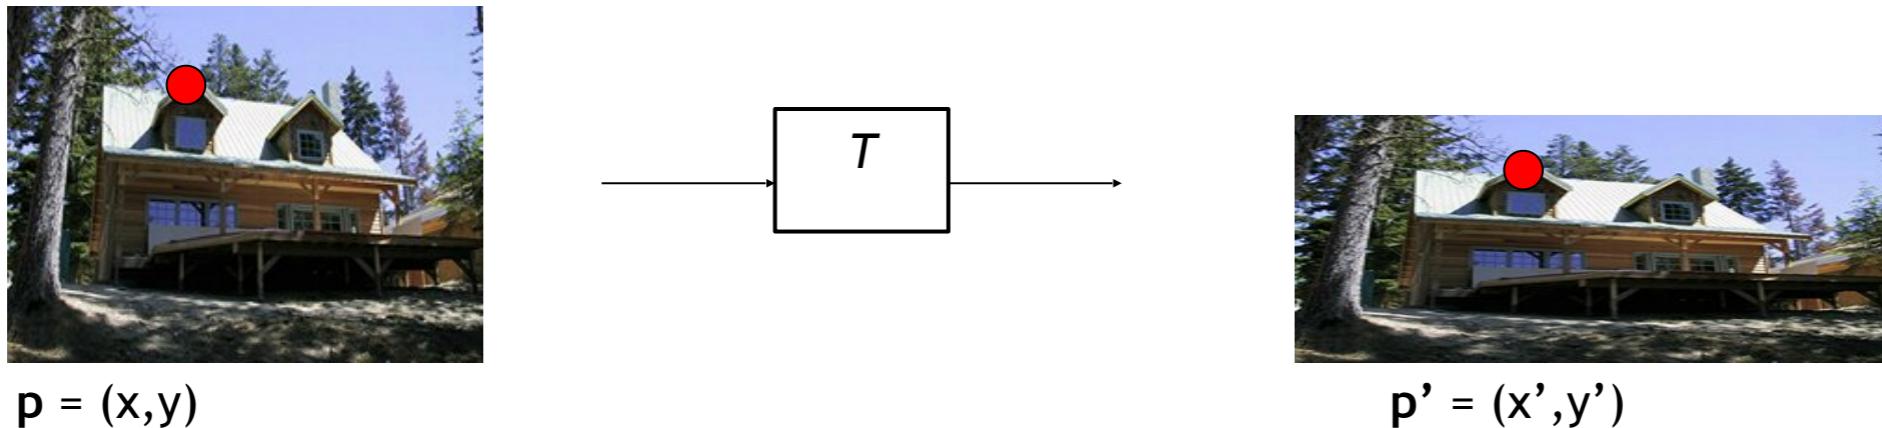
- For each pixel  $x$  in the destination image
- Determine where it comes from as  $H^{-1}(x)$
- Get color from that location

# Backward Warping



- For each pixel  $x$  in the destination image
- Determine where it comes from as  $H^{-1}(x)$
- Get color from that location

# Where the pixels go?



- Transformation  $T$  is a coordinate-changing machine:  
$$p' = T(p)$$
- What does it mean that  $T$  is global?
  - Is the same for any point  $p$
  - can be described by just a few numbers (parameters)
- Let's consider *linear* xforms (can be represented by a 2D matrix):

$$p' = Tp \quad \begin{bmatrix} x' \\ y' \end{bmatrix} = T \begin{bmatrix} x \\ y \end{bmatrix}$$

# Linear Transformation

- Uniform scaling by  $s$ :



(0,0)



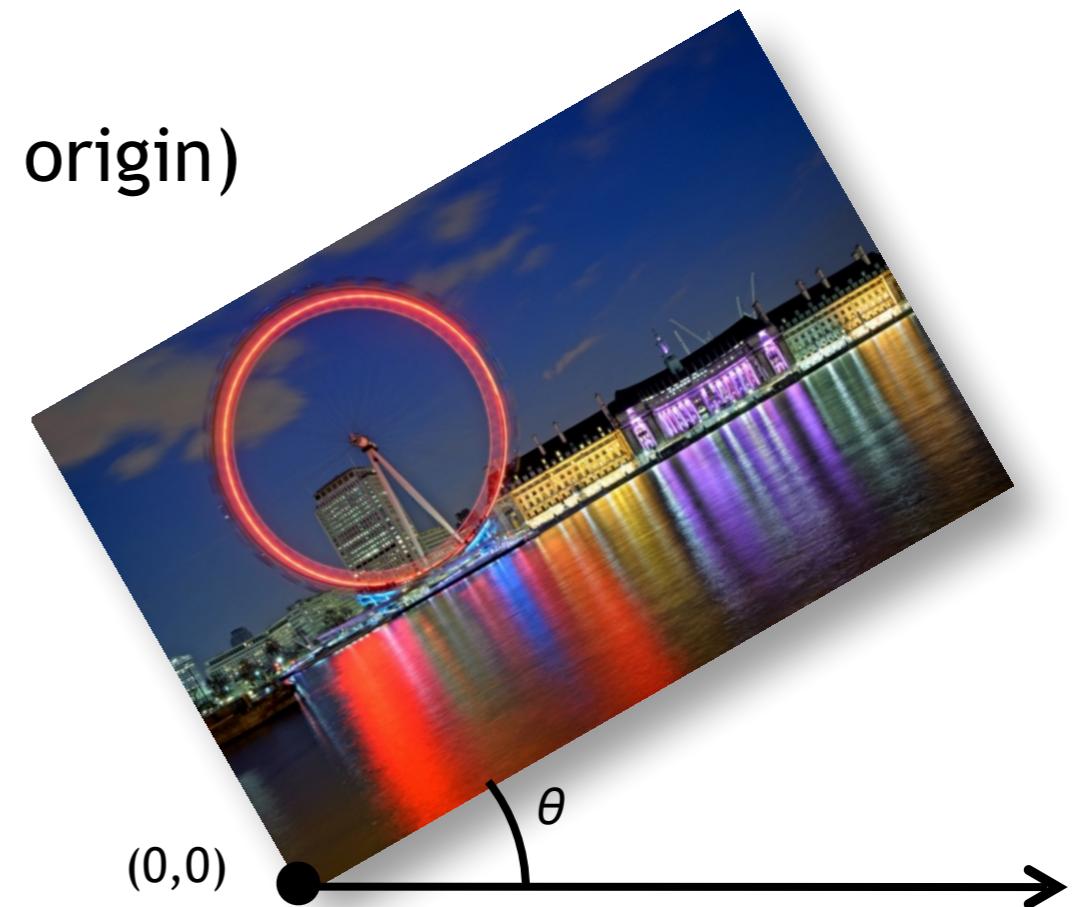
(0,0)

$$S = \begin{bmatrix} s & 0 \\ 0 & s \end{bmatrix}$$

What is the inverse?

# Linear Transformation

- Rotation by angle  $\theta$  (about the origin)



$$\mathbf{R} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

What is the inverse?

For rotations:

$$\mathbf{R}^{-1} = \mathbf{R}^T$$

# Transformation with 2x2 Matrices

- What types of transformations can be represented with a 2x2 matrix?

2D mirror about Y axis?

$$\begin{aligned}x' &= -x \\y' &= y\end{aligned}$$

$$\mathbf{T} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

2D mirror across line  $y = x$ ?

$$\begin{aligned}x' &= y \\y' &= x\end{aligned}$$

$$\mathbf{T} = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

# Affine Transformation

- Affine transformations are combinations of ...
  - Linear transformations, and
  - Translations

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Properties of affine transformations:
  - Origin does not necessarily map to origin
  - Lines map to lines
  - Parallel lines remain parallel
  - Ratios are preserved
  - Closed under composition

# Affine Transformation

$$\mathbf{T} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \quad \leftarrow \text{any transformation with last row } [0 \ 0 \ 1] \text{ we call an } \textit{affine} \text{ transformation}$$

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

# Basic transformation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

2D *in-plane* rotation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

# Projective transformation

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \quad \text{affine transformation}$$


what happens when we mess with this row?

# Projective transformation

$$\mathbf{H} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix}$$

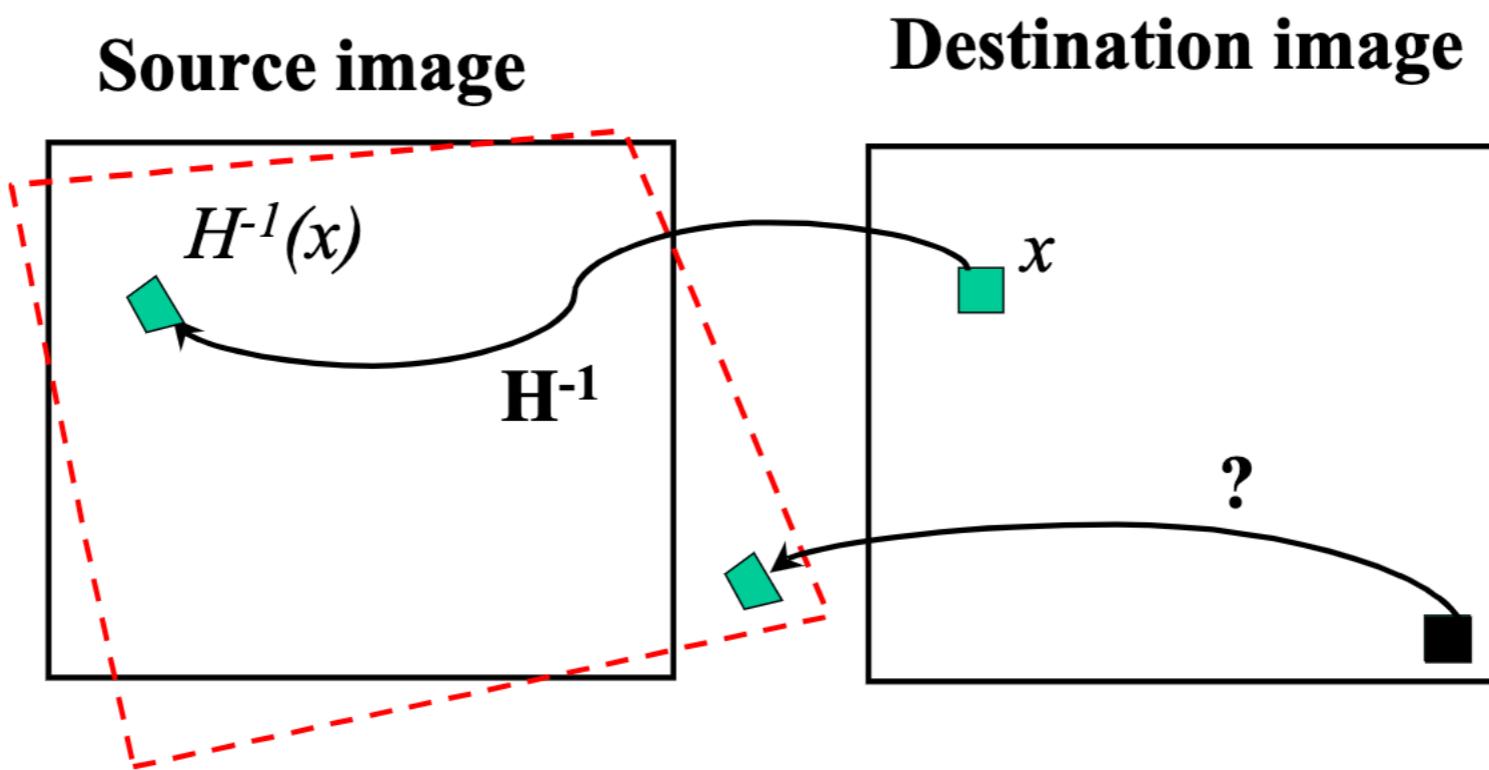
Called a *homography*  
(or *planar perspective map*)



# Projective transformation

- Projective transformations ...
    - Affine transformations, and
    - Projective warps
  - Properties of projective transformations:
    - Origin does not necessarily map to origin
    - Lines map to lines
    - Parallel lines do not necessarily remain parallel
    - Ratios are not preserved
    - Closed under composition
- $$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

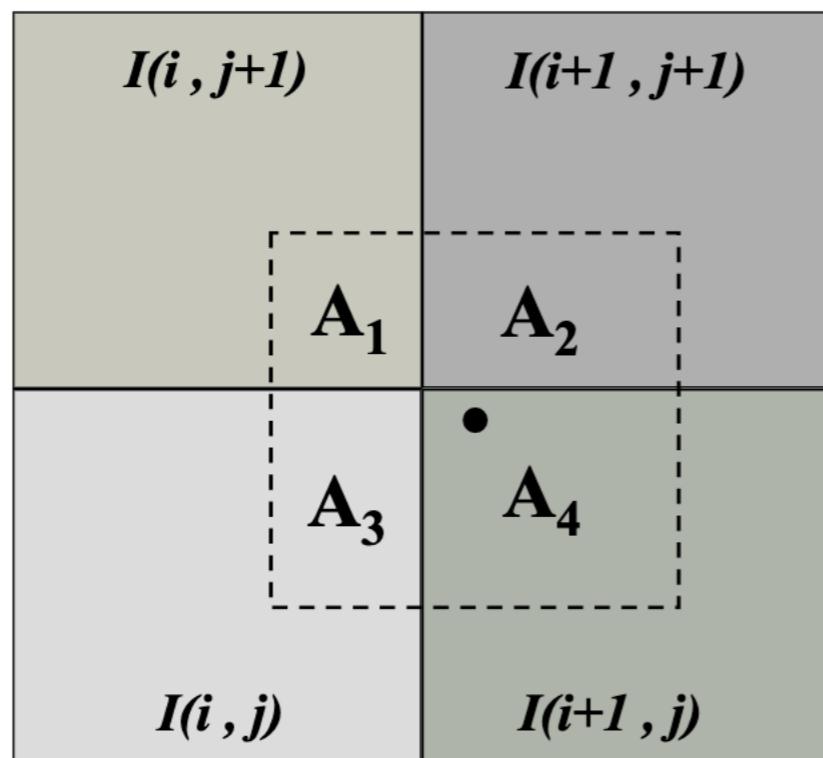
# Backward Warping



- For each pixel  $x$  in the destination image
- Determine where it comes from as  $H^{-1}(x)$
- Get color from that location

# Bilinear interpolation

What do we mean by “get color from that location”?  
Consider grey values. What is intensity at (x,y)?



**Bilinear Interpolation:**  
Weighted average

$$\begin{aligned}I(x,y) = & A_3 * I(i,j) \\& + A_4 * I(i+1,j) \\& + A_2 * I(i+1,j+1) \\& + A_1 * I(i,j+1)\end{aligned}$$

Bilinear interpolation; the output pixel value is a weighted average of pixels in the nearest 2-by-2 neighborhood

# Linear Interpolation (recall)

First, consider linear interpolation



**Intuition:** Given two pixel values, what should the value be at some intermediate point between them?



If close to  $(i, j)$ , should be intensity similar to  $I(i, j)$

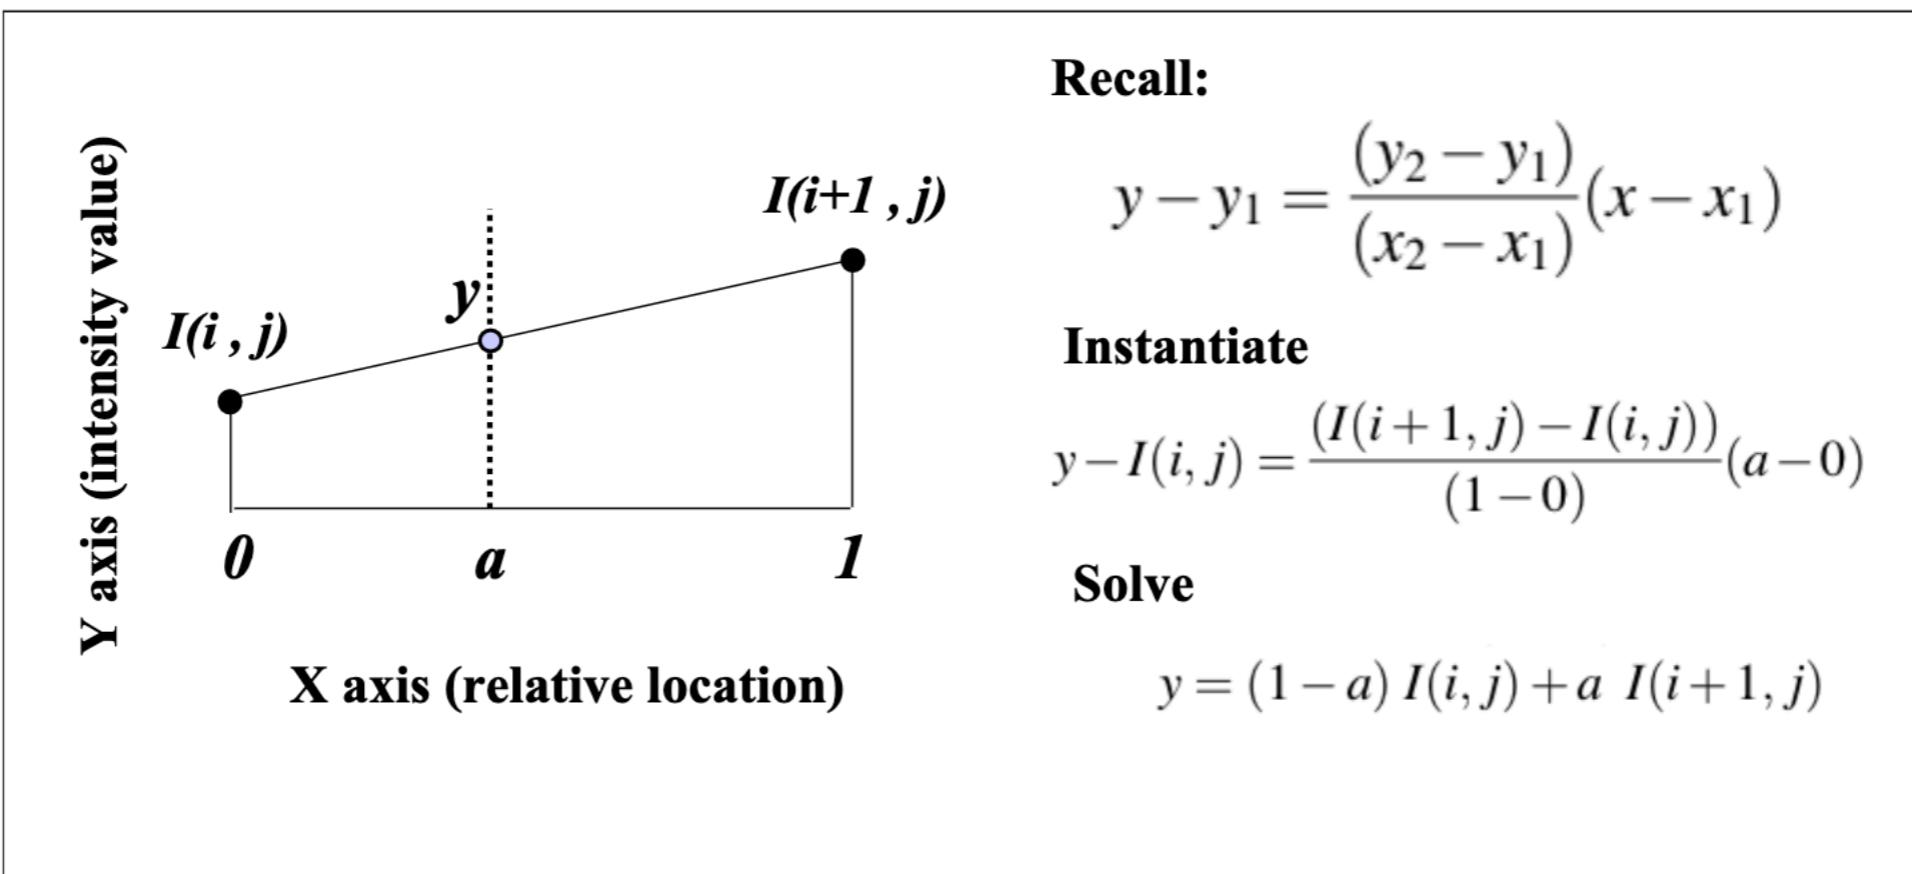
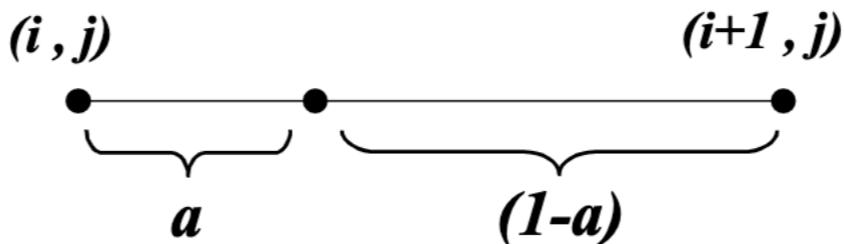


If equidistant from both, should be average of the two intensities

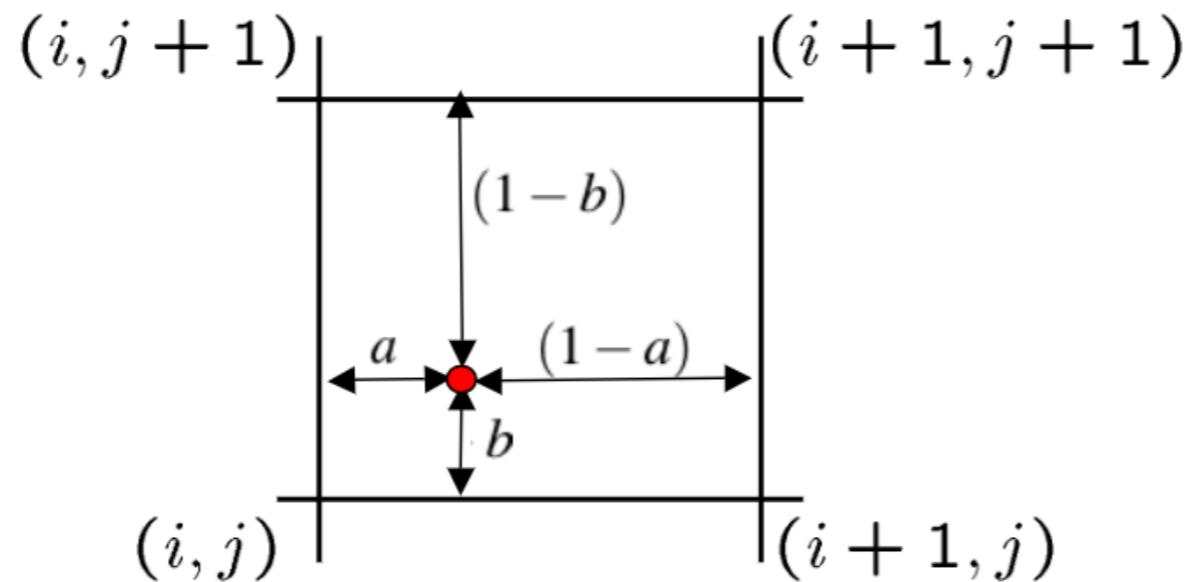


If close to  $(i+1, j)$ , should be intensity similar to  $I(i+1, j)$

# Linear Interpolation (recall)



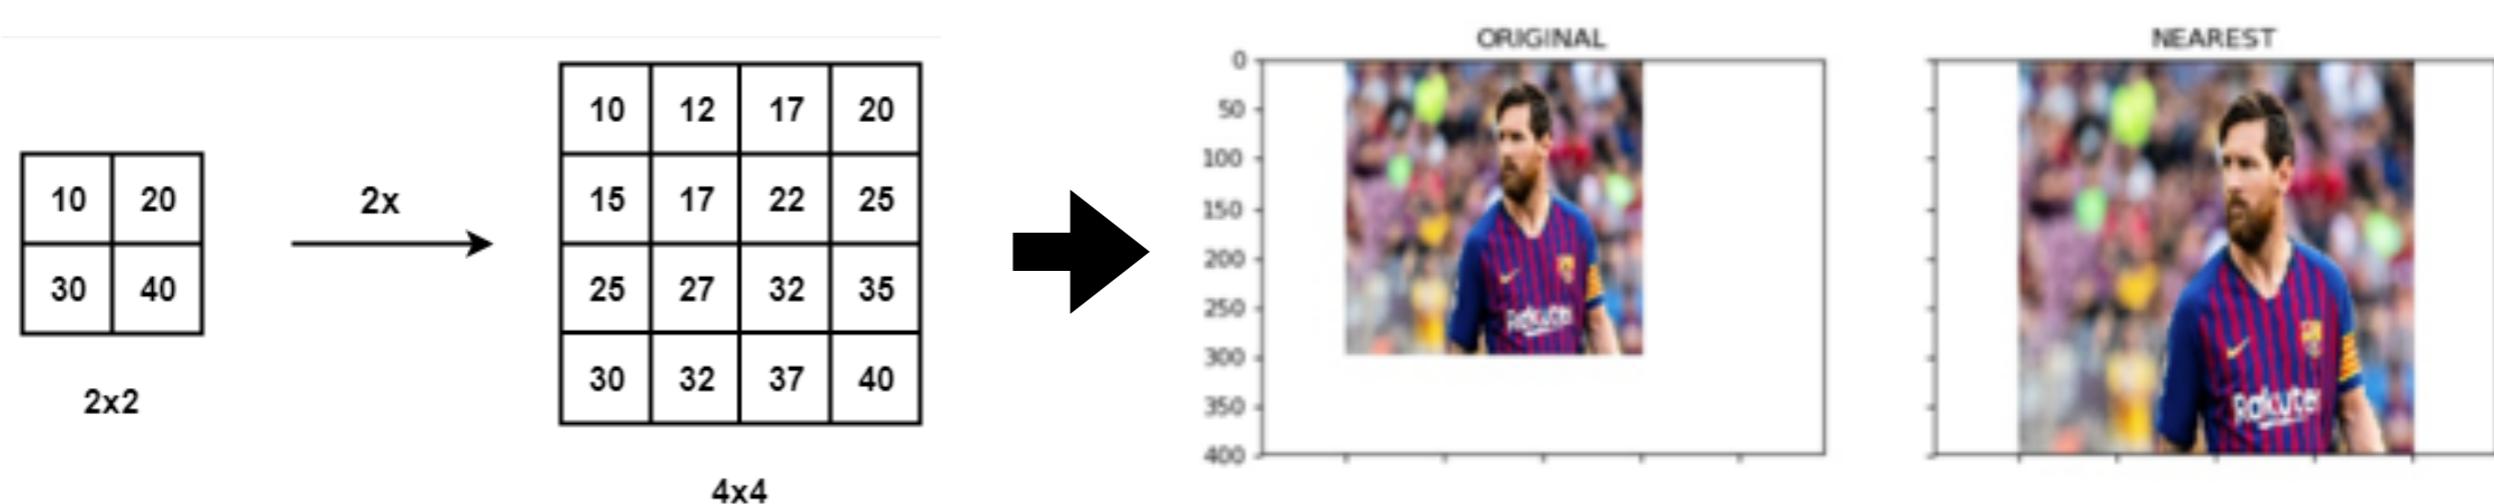
# Bilinear Interpolation



$$\begin{aligned}\mathbf{I} = & (1 - a)(1 - b) I(i, j) \\ & + a (1 - b) I(i + 1, j) \\ & + (1 - a) b I(i, j + 1) \\ & + a b I(i + 1, j + 1)\end{aligned}$$

# Image resizing

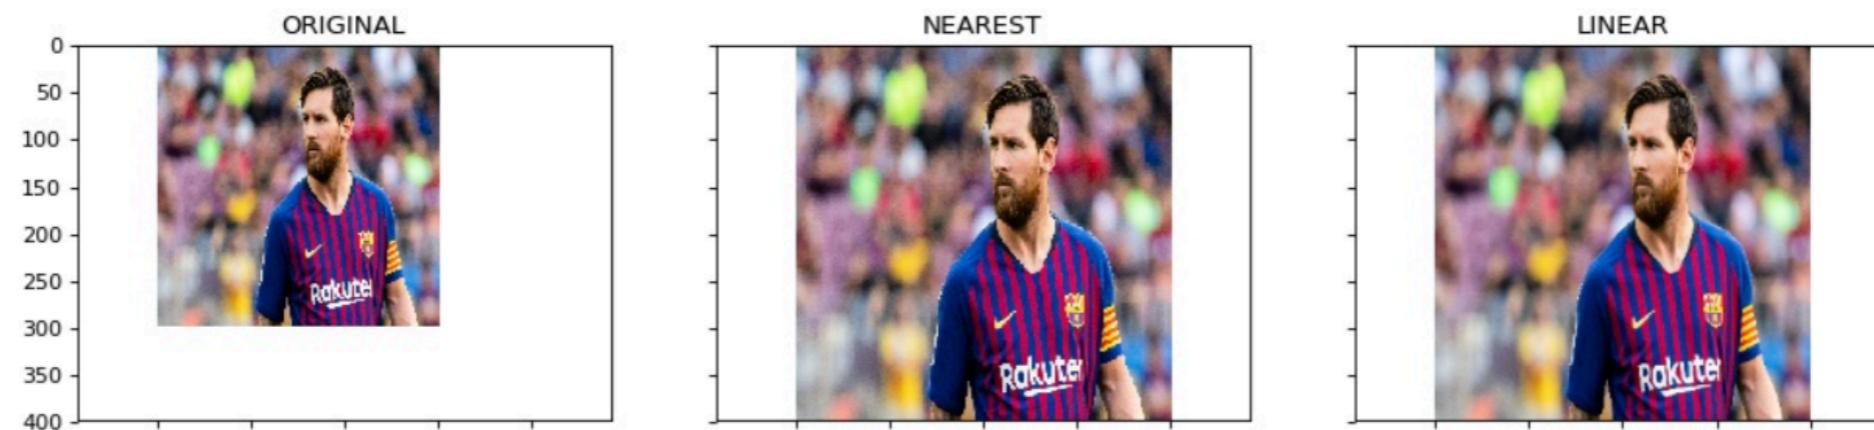
- Machine learning models work with a fixed sized input. The same idea applies to computer vision models as well. The images we use for training our model must be of the **same size**.
- Images can be easily scaled up and down



- Different interpolation and downsampling methods are supported by OpenCV. OpenCV's resize function uses bilinear interpolation by default.

# Resizing with OpenCV

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 #reading the image
6 image = cv2.imread('index.jpg')
7 #converting image to size (100,100,3)
8 smaller_image = cv2.resize(image,(100,100),interpolation='linear')
9 #plot the resized image
10 plt.imshow(smaller_image)
```



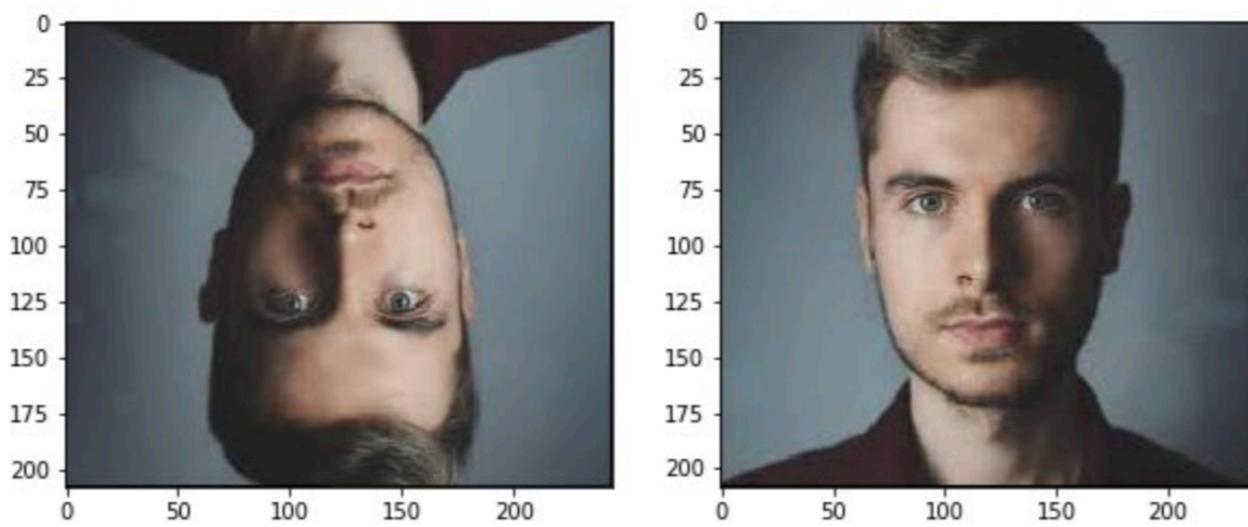
# Rotate

- Suppose we are building an image classification model for identifying the animal present in an image. So, both the images shown below should be classified as ‘dog’:



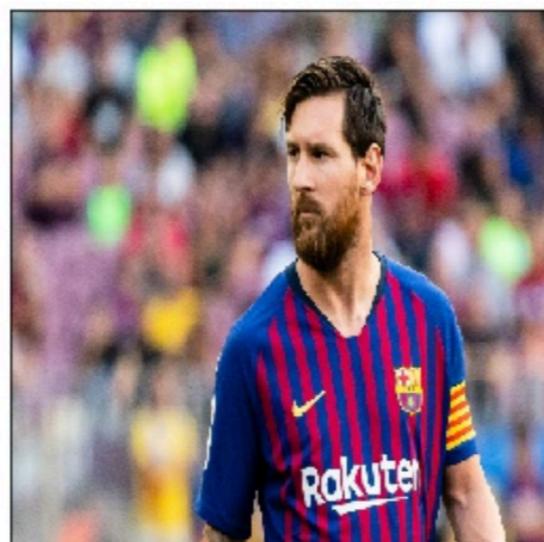
# Rotation with OpenCV

```
1 #importing the required libraries
2 import numpy as np
3 import cv2
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6 image = cv2.imread('index.png')
7 rows,cols = image.shape[:2]
8 #(col/2,rows/2) is the center of rotation for the image
9 # M is the coordinates of the center
10 M = cv2.getRotationMatrix2D((cols/2,rows/2),90,1)
11 dst = cv2.warpAffine(image,M,(cols,rows))
12 plt.imshow(dst)
```

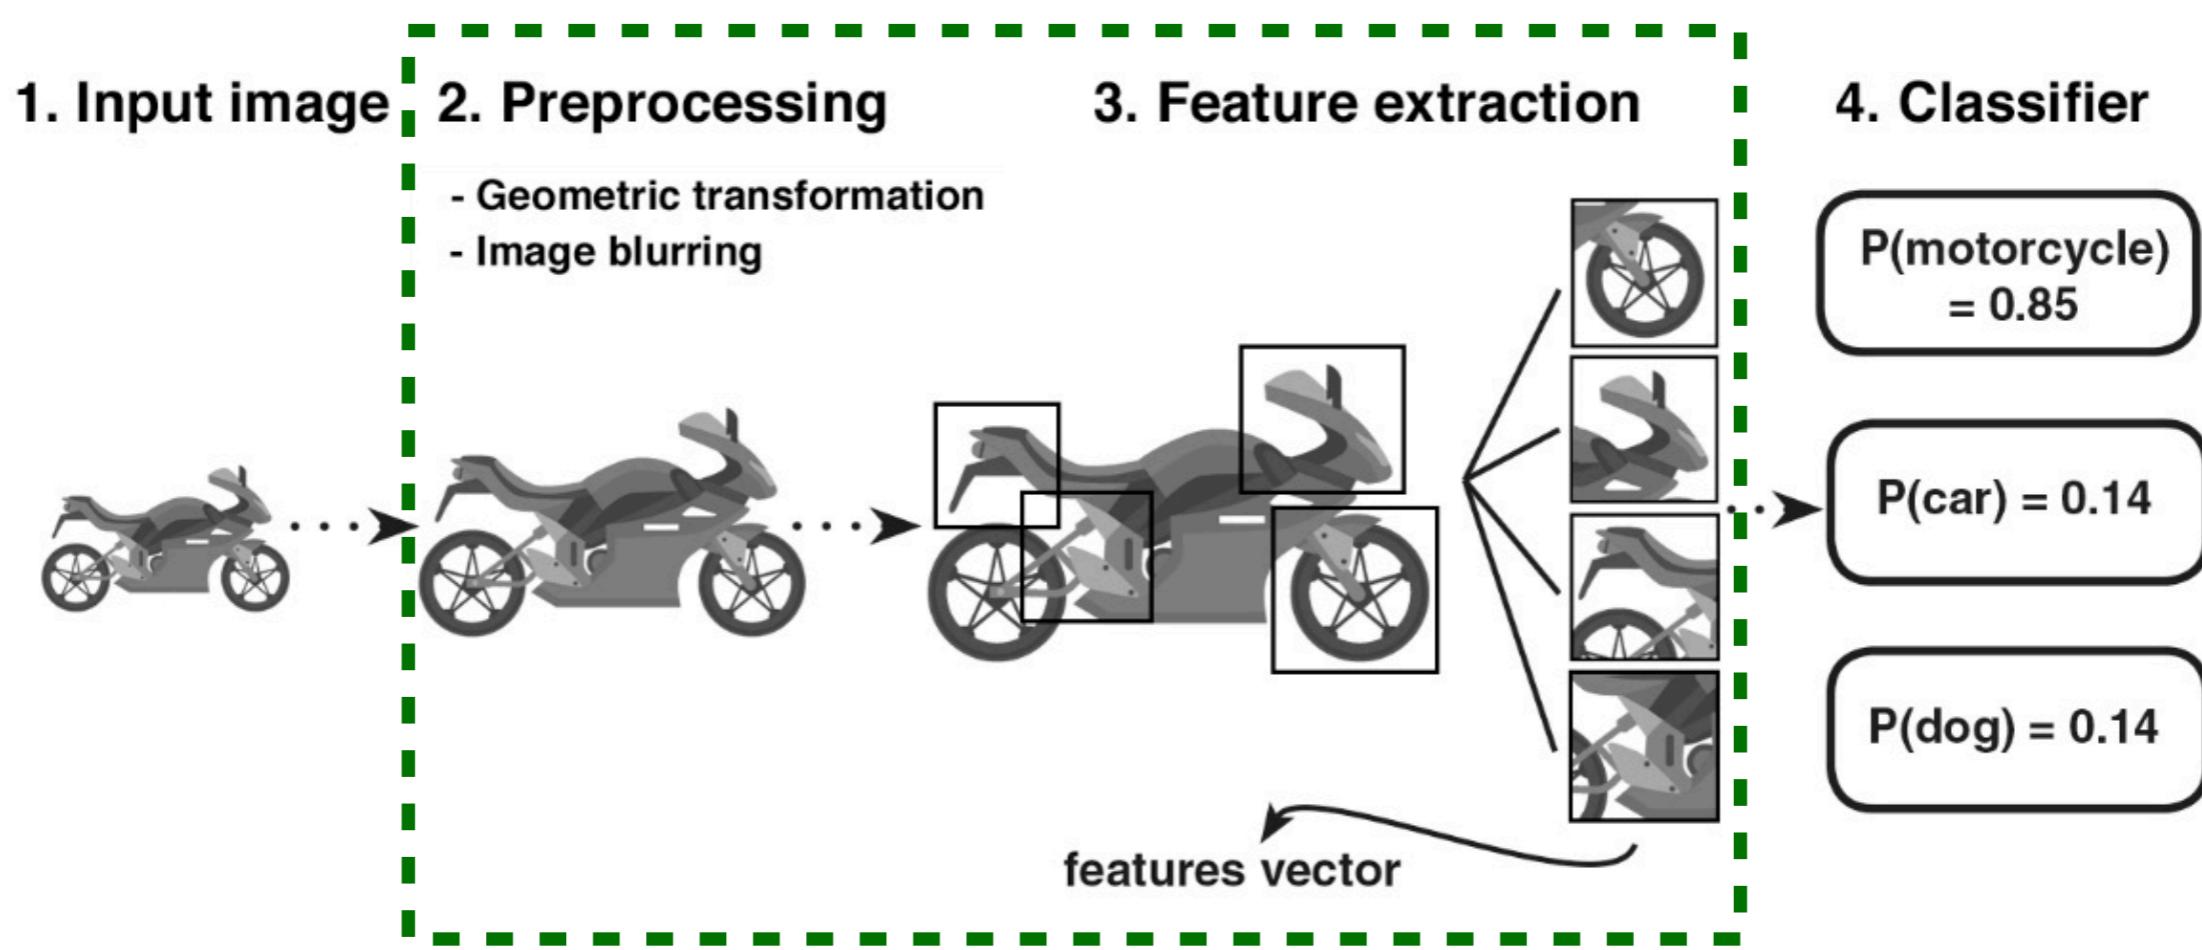


# Shifting

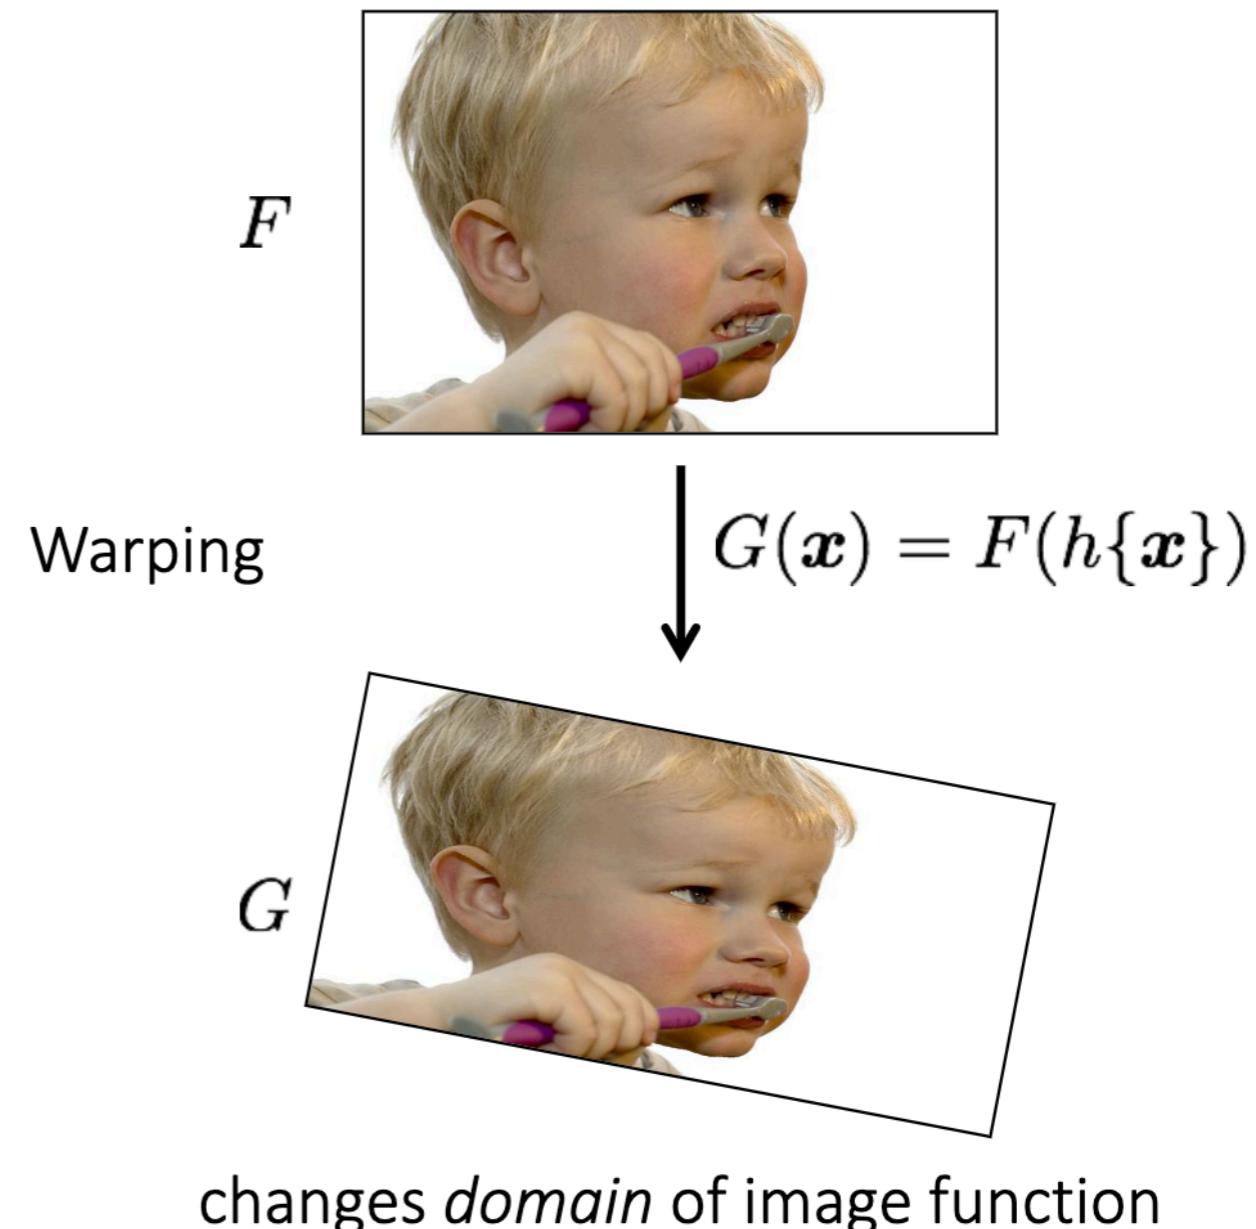
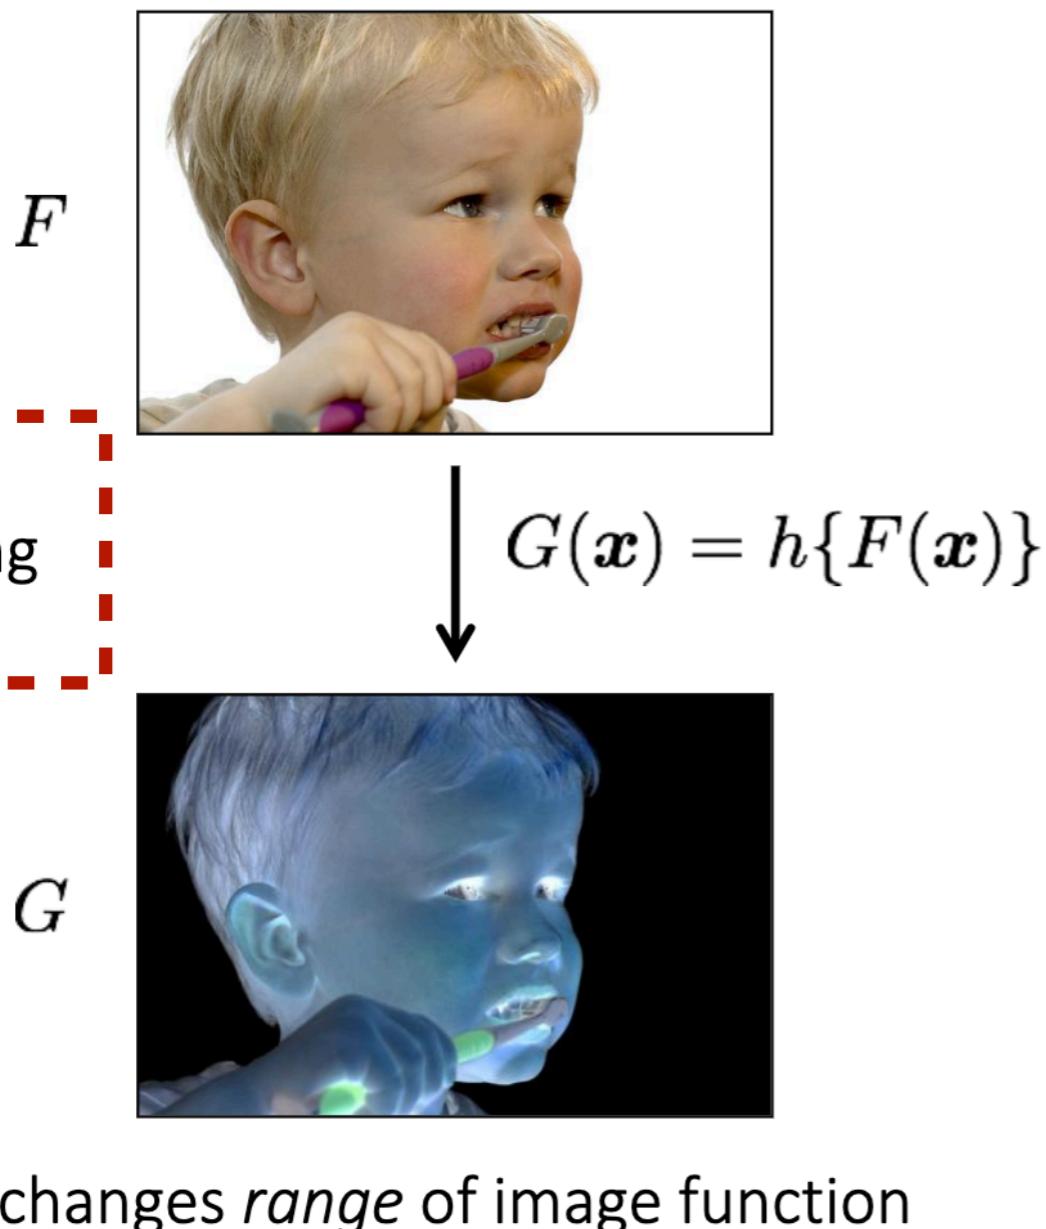
```
1 #importing the required libraries
2 import numpy as np
3 import cv2
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6 #reading the image
7 image = cv2.imread('index.png')
8 #shifting the image 100 pixels in both dimensions
9 M = np.float32([[1,0,-100],[0,1,-100]])
10 dst = cv2.warpAffine(image,M,(cols,rows))
11 plt.imshow(dst)
```



# CV Pipeline

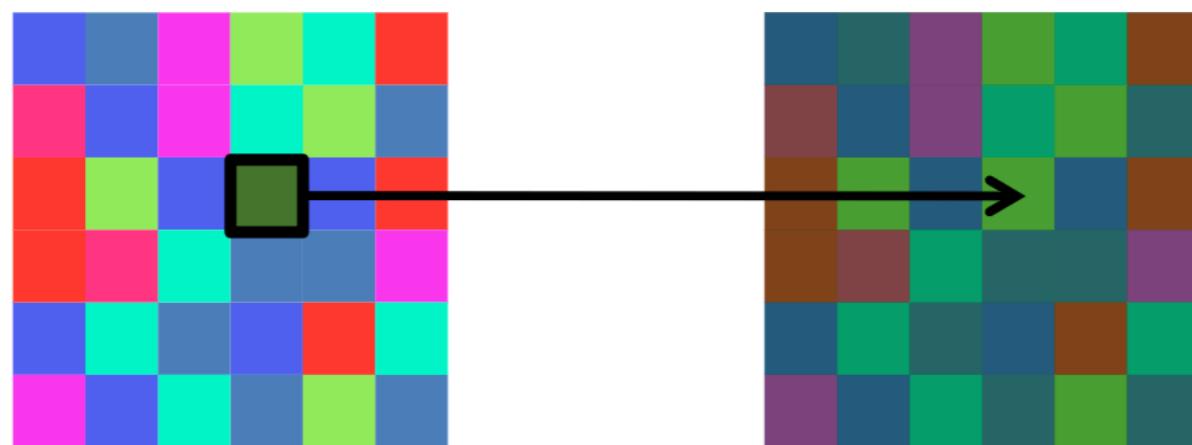


# Image transformations



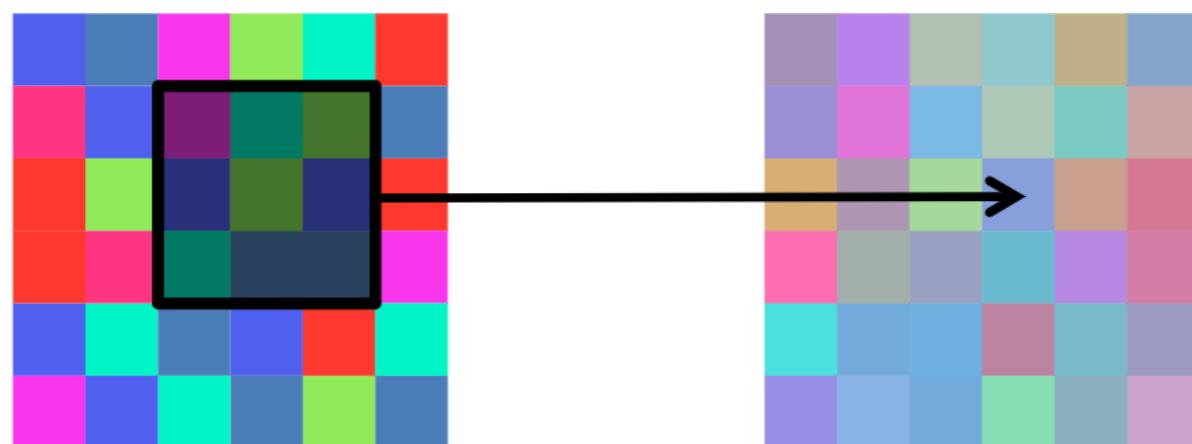
# Image filtering

Point Operation



1D  
point processing

Neighborhood Operation



2D  
“filtering”

# Point processing

original



darken



lower contrast



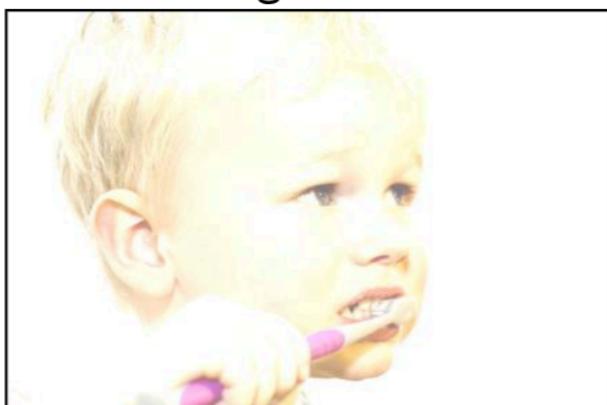
non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast



# Point processing (How to implement them?)

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear raise contrast



$$\left(\frac{x}{255}\right)^2 \times 255$$

# Filtering

- **Filtering:**
  - Form a new image whose pixels are a combination original pixel values

## Goals:

- Extract useful information from the images
  - Features (edges, corners, blobs...)
- Modify or enhance image properties:
  - super-resolution; in-painting; de-noising

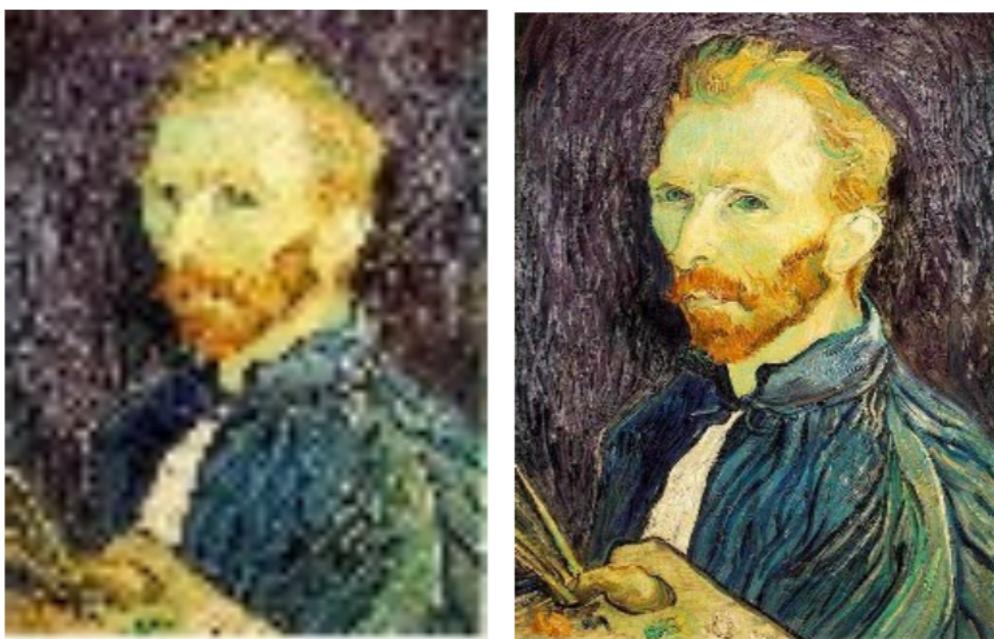
# Enhancing Examples

De-noising



Salt and pepper noise

Super-resolution



In-painting



Bertamio et al

# Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function



	7	

Modified image data

# 2D discrete-space systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

$$g = \mathcal{S}[f], \quad g[n, m] = \mathcal{S}\{f[n, m]\}$$

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$$

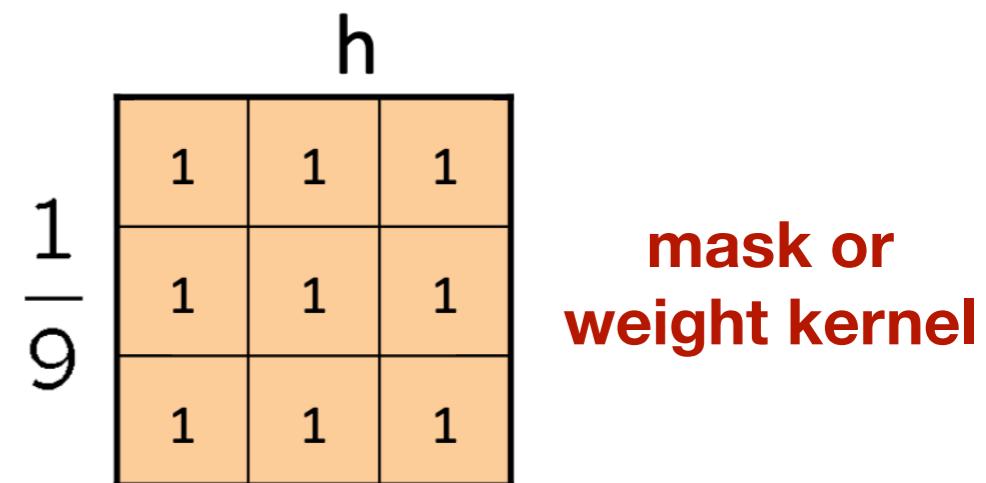
# Filter example: Moving average

- Also known as **Box filter**

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l]$$

$$= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$(f * h)[m, n] = \frac{1}{9} \sum_{k,l} f[k, l] h[m - k, n - l]$$



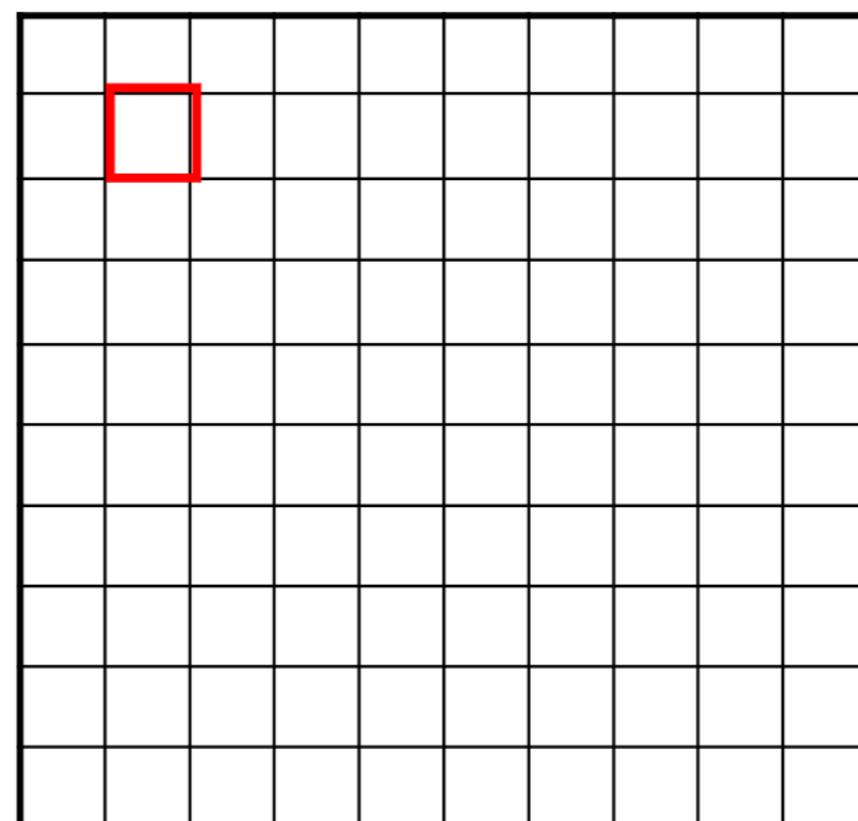
- 2D moving average over a 3x3 window of neighborhood

# Filter example: Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$



$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

# Filter example: Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

# Filter example: Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$


$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

# Filter example: Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

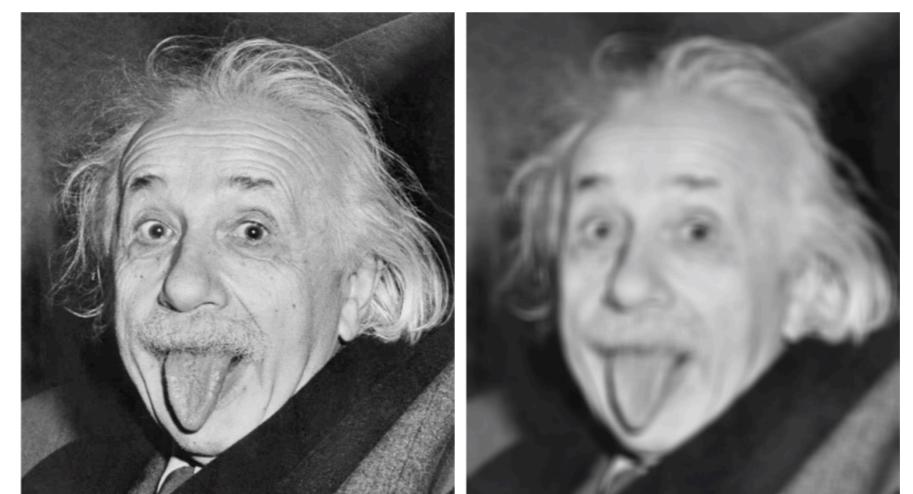
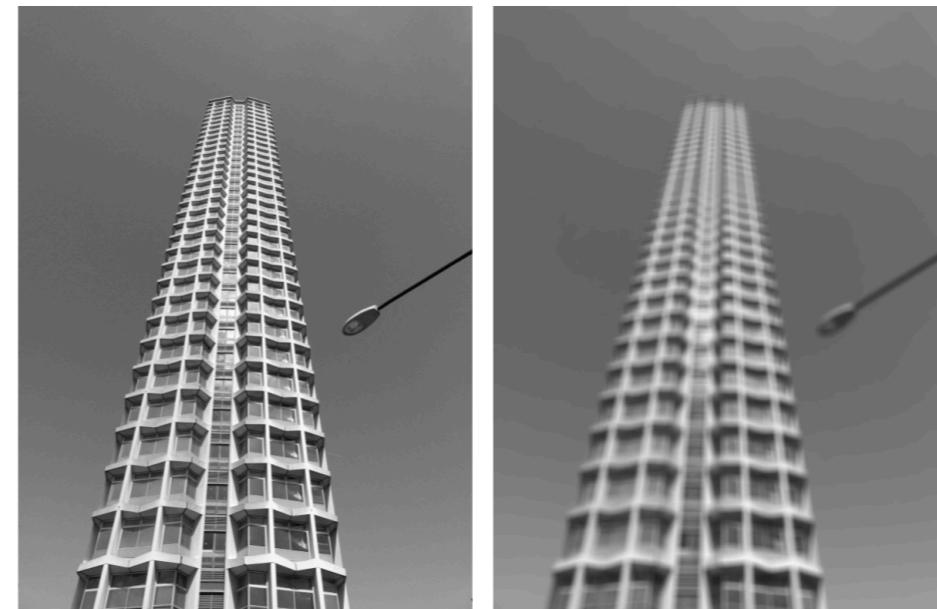
$G[x, y]$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Source: S. Seitz

# Filter example: Moving average



Achieve smoothing effect (remove sharp features)

# Filter example: Image Segmentation

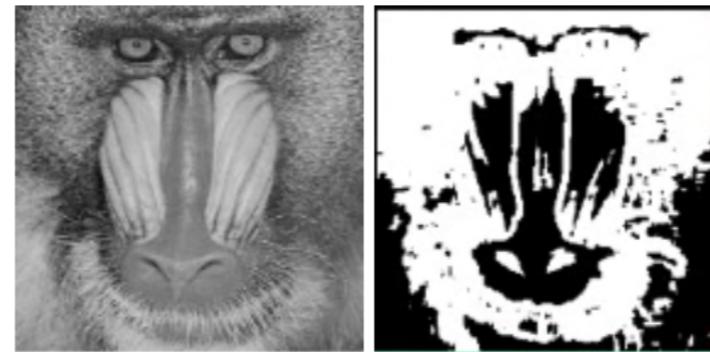
- Image segmentation based on a simple threshold:

$$g[n, m] = \begin{cases} 255, & f[n, m] > 100 \\ 0, & \text{otherwise.} \end{cases}$$

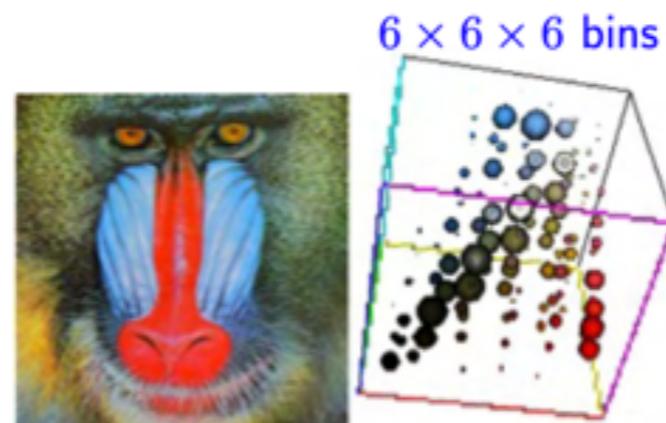


# Filter example: Image Segmentation

- Non-contextual: grouping pixels with similar global features.



- Contextual: grouping pixels with similar features and in close locations.



# Filtering properties

Commutative:  $f * g = g * f$

Associative:  $(f * g) * h = f * (g * h)$

Distributive:  $(f + g) * h = f * h + g * h$

- | Linear:  $(a f + b g) * h = a f * h + b g * h$
- | Shift Invariant:  $f(x+t) * h = (f * h)(x+t)$

Differentiation rule:

$$\frac{\partial}{\partial x} (f * g) = \frac{\partial f}{\partial x} * g$$

# Shift invariant

- Filter replaces each pixel by a linear combination of its neighbors (and possibly itself). The combination is determined by the filter's kernel.

If  $f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$  then

$$f[n - n_0, m - m_0] \xrightarrow{\mathcal{S}} g[n - n_0, m - m_0]$$

for every input image  $f[n, m]$  and shifts  $n_0, m_0$

# Shift invariant

Is the moving average system shift invariant?

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

# Shift invariant

Is the moving average system shift invariant?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$f[n - n_0, m - m_0]$$

$$\xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[(n - n_0) - k, (m - m_0) - l]$$

$$= g[n - n_0, m - m_0]$$

Yes!

# Linear filtering

Linear filtering means linear combination of neighboring pixel values.

- **S** is a linear system (function) iff it *S satisfies*

$$\mathcal{S}[\alpha f_1 + \beta f_2] = \alpha \mathcal{S}[f_1] + \beta \mathcal{S}[f_2]$$

superposition property

- Is the moving average system a linear system?
- Is thresholding a linear system?

# Convolution

- Any linear, shift invariant operator can be represented as **convolution!**

