

Deep Learning

IFT6758 - Data Science

Sources:

<http://www.cs.cmu.edu/~16385/>

<http://cs231n.stanford.edu/syllabus.html>

<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

<https://www.cs.ubc.ca/labs/lci/mlrg/slides/rnn.pdf>

Announcements

- Grades of Assignment 2 is published on Gradescope!
- Check Evaluation 7, the scores are on scoreboard!
- Grade of mid-term will be published on Gradescope by the end of this week!
- Homework 3 is on Gradescope and it is due on **November 28**.
- Homework 4 will be published on Gradescope on Monday.

Crash Course to Deep Learning

1950s Age of the Perceptron

1957 The Perceptron (Rosenblatt)
1969 Perceptrons (Minsky, Papert)



1980s Age of the Neural Network


1986 Back propagation (Hinton)

1990s Age of the Graphical Model
2000s Age of the Support Vector Machine



2010s Age of the Deep Network

deep learning = known algorithms + computing power + big data

Perceptron



Inspiration from Biology



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Neural nets/perceptrons are **loosely** inspired by biology.

But they certainly are **not** a model of how the brain works, or even how neurons work.

1: **function** PERCEPTRON ALGORITHM

2: $\mathbf{w}^{(0)} \leftarrow \mathbf{0}$

3: **for** $t = 1, \dots, T$ **do**

4: RECEIVE($\mathbf{x}^{(t)}$) $\mathbf{x} \in \{0, 1\}^N$ N-d binary vector

5: $\hat{y}_A^{(t)} = \text{sign} \left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle \right)$ perceptron is just one line of code!
 sign of zero is +1

6: RECEIVE($y^{(t)}$) $y \in \{1, -1\}$


7: $w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$

RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



RECEIVE($\mathbf{x}^{(t)}$)


$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

update w

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)


$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

update w

no match!

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1) (0,0) -1 (1,-1) 1



RECEIVE($\mathbf{x}^{(t)}$)


$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)

RECEIVE($\mathbf{x}^{(t)}$)
observation (-1,1)



RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$


RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)

$$\begin{aligned}\hat{y}_A^{(t)} &= \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right) \\ &= 1\end{aligned}$$

observation (-1,1)



RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)


$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1)

$$\begin{aligned}\hat{y}_A^{(t)} &= \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right) \\ &= 1\end{aligned}$$

observation (-1,1)

label +1



RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)


$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

update w

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$

(-1,1) (-1,1) +1 (-1,1) 0

match!




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)

$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)

$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$




RECEIVE($\mathbf{x}^{(t)}$)


$$\hat{y}_A^{(t)} = \text{sign}\left(\langle \mathbf{w}^{(t-1)}, \mathbf{x}^{(t)} \rangle\right)$$

RECEIVE(y^t)


$$w_n^{(t)} = w_n^{(t-1)} + y_t \cdot x_n^{(t)} \cdot \mathbf{1}[y^{(t)} \neq \hat{y}^{(t)}]$$



Perceptron




Perceptron



Programming the 'forward pass'

Activation function (sigmoid, logistic function)

```
float f(float a)
{
    return 1.0 / (1.0+ exp(-a));
}
```




Perceptron function (logistic regression)

```
float perceptron(vector<float> x, vector<float> w)
{
    float a = dot(x, w);
    return f(a);
}
```


Neural Network

Connect a bunch of perceptrons together ...
a collection of connected perceptrons



Neural Network


Connect a bunch of perceptrons together ...
a collection of connected perceptrons



How many perceptrons in this neural network?


Neural Network

Connect a bunch of perceptrons together ...
a collection of connected perceptrons




Neural Network

Connect a bunch of perceptrons together ...
a collection of connected perceptrons




Neural Network




...also called a **Multi-layer Perceptron** (MLP)

this layer is a
'fully connected layer'



all pairwise neurons between layers are connected


so is this



all pairwise neurons between layers are connected

How many neurons (perceptrons)?

How many weights (edges)?




How many learnable parameters total?

How many neurons (perceptrons)?

$$4 + 2 = 6$$

How many weights (edges)?




How many learnable parameters total?

How many neurons (perceptrons)?

$$4 + 2 = 6$$

How many weights (edges)?

$$(3 \times 4) + (4 \times 2) = 20$$




How many learnable parameters total?

How many neurons (perceptrons)?

$$4 + 2 = 6$$

How many weights (edges)?

$$(3 \times 4) + (4 \times 2) = 20$$




How many learnable parameters total?

$$20 + 4 + 2 = 26$$

bias terms


performance usually tops out at 2-3 layers,
deeper networks don't really improve performance...



...with the exception of **convolutional** networks for images

How to train perceptrons?


world's smallest perceptron!



$$y = wx$$

What does this look like?

world's smallest perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

Learning a Perceptron

Given a set of samples and a Perceptron


$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

Estimate the parameters of the Perceptron

$$w$$


Given training data:



What do you think the weight parameter is?

$$y = wx$$

Given training data:



What do you think the weight parameter is?

$$y = wx$$

not so obvious as the network gets more complicated so we use ...

An Incremental Learning Strategy (gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = w \cdot x$$

An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = w \cdot x$$

Modify weight w such that \hat{y} gets ‘**closer**’ to y

An Incremental Learning Strategy

(gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = w x$$

Modify weight w such that \hat{y} gets ‘**closer**’ to y



perceptron
parameter

perceptron
output

true
label
Université
de Montréal

An Incremental Learning Strategy

(gradient descent)


Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = wx$$

Modify weight w such that \hat{y} gets ‘closer’ to y



Before diving into gradient descent, we need to understand ...

Loss Function


defines what it means to be
close to the true solution

YOU get to chose the loss function!

(some are better than others depending on what you want to do)


Squared Error (L2)

(a popular loss function) ((why?))




L1 Loss

$$\ell(\hat{y}, y) = |\hat{y} - y|$$




L2 Loss

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$




Zero-One Loss

$$\ell(\hat{y}, y) = \mathbf{1}[\hat{y} = y]$$




Hinge Loss

$$\ell(\hat{y}, y) = \max(0, 1 - y \cdot \hat{y})$$



back to the...

world's smallest perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

function of **ONE** parameter!

Learning a Perceptron

Given a set of samples and a Perceptron

$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

*what is this
activation function?*

Estimate the parameter of the Perceptron

$$w$$

Learning a Perceptron

Given a set of samples and a Perceptron

$$\{x_i, y_i\}$$

$$y = f_{\text{PER}}(x; w)$$

*what is this
activation function?*  linear function! $f(x) = wx$

Estimate the parameter of the Perceptron

$$w$$

Learning Strategy (gradient descent)

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$\hat{y} = wx$$

Modify weight w such that \hat{y} gets ‘closer’ to y



perceptron
parameter

perceptron
output

true
label
Université
de Montréal 

Code to train your perceptron:

```
for n = 1 . . . N  
    w = w + (yn -  $\hat{y}$ )xi;
```


just one line of code!

Gradient descent

(partial) derivatives tell us how much one variable affects another

Gradient descent


Given a fixed-point or
a function,
move in the direction
opposite of the
gradient



Gradient descent

update rule:


$$w = w - \nabla w$$



Backpropagation

back to the...

World's Smallest Perceptron!



$$y = wx$$

(a.k.a. line equation, linear regression)

function of **ONE** parameter!

Training the world's smallest perceptron

for $n = 1 \dots N$

This is just gradient descent, that means...

$$w = w + \underline{(y_n - \hat{y})x_i};$$



this should be the gradient of the loss function

Now where does this come from?

$$\frac{d\mathcal{L}}{dw}$$

...is the rate at which **this** will change...

$$\mathcal{L} = \frac{1}{2}(y - \hat{y})^2$$

the loss function

... per unit change of **this**

$$y = wx$$

the weight parameter

Let's compute the derivative...

Compute the derivative

$$\begin{aligned}\frac{d\mathcal{L}}{dw} &= \frac{d}{dw} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{dwx}{dw} \\ &= -(y - \hat{y})x = \nabla w \quad \text{just shorthand}\end{aligned}$$

That means the weight update for **gradient descent** is:

$$\begin{aligned}w &= w - \nabla w \quad \text{move in direction of negative gradient} \\ &= w + (y - \hat{y})x\end{aligned}$$

Gradient Descent (world's smallest perceptron)

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = wx_i$$

b. Compute Loss

$$\mathcal{L}_i = \frac{1}{2}(y_i - \hat{y})^2$$

2. Update

a. Back Propagation

$$\frac{d\mathcal{L}_i}{dw} = -(y_i - \hat{y})x_i = \nabla w$$

b. Gradient update


$$w = w - \nabla w$$

Training the world's smallest perceptron

for $n = 1 \dots N$

$$w = w + (y_n - \hat{y})x_i;$$

world's (second) smallest perceptron!



function of **two** parameters!

Gradient Descent

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

b. Compute Loss

we just need to compute partial derivatives for this network

2. Update

a. Back Propagation

b. Gradient update



Derivative computation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial}{\partial w_1} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_1} \\ &= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_1} \\ &= -(y - \hat{y}) \frac{\partial w_1 x_1}{\partial w_1} \\ &= -(y - \hat{y}) x_1 = \nabla w_1\end{aligned}$$

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial}{\partial w_2} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_2} \\ &= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_2} \\ &= -(y - \hat{y}) \frac{\partial w_2 x_2}{\partial w_2} \\ &= -(y - \hat{y}) x_2 = \nabla w_2\end{aligned}$$

Why do we have partial derivatives now?

Derivative computation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial}{\partial w_1} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_1} \\ &= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_1} \\ &= -(y - \hat{y}) \frac{\partial w_1 x_1}{\partial w_1} \\ &= -(y - \hat{y}) x_1 = \nabla w_1\end{aligned}$$

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial}{\partial w_2} \left\{ \frac{1}{2}(y - \hat{y})^2 \right\} \\ &= -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_2} \\ &= -(y - \hat{y}) \frac{\partial \sum_i w_i x_i}{\partial w_2} \\ &= -(y - \hat{y}) \frac{\partial w_2 x_2}{\partial w_2} \\ &= -(y - \hat{y}) x_2 = \nabla w_2\end{aligned}$$

Gradient Update

$$\begin{aligned}w_1 &= w_1 - \eta \nabla w_1 \\ &= w_1 + \eta(y - \hat{y}) x_1\end{aligned}$$

$$\begin{aligned}w_2 &= w_2 - \eta \nabla w_2 \\ &= w_2 + \eta(y - \hat{y}) x_2\end{aligned}$$

Gradient Descent

For each sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i = \frac{1}{2}(y_i - \hat{y})$$

(side computation to track loss.
not needed for backprop)

2. Update

a. Back Propagation

$$\nabla w_{1i} = -(y_i - \hat{y})x_{1i}$$
$$\nabla w_{2i} = -(y_i - \hat{y})x_{2i}$$

b. Gradient update


$$w_{1i} = w_{1i} + \eta(y - \hat{y})x_{1i}$$
$$w_{2i} = w_{2i} + \eta(y - \hat{y})x_{2i}$$

(adjustable step size)





We haven't seen a lot of 'propagation' yet
because our perceptrons only had one layer...


Multi-layer perceptron




function of **FOUR** parameters and **FOUR** layers!








$$a_1 = w_1 \cdot x + b_1$$




$$a_1 = w_1 \cdot x + b_1$$




$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$



$$a_1 = w_1 \cdot x + b_1$$


$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$


$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$

$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$



$$a_1 = w_1 \cdot x + b_1$$

$$a_2 = w_2 \cdot f_1(w_1 \cdot x + b_1)$$

$$a_3 = w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1))$$

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$

Entire network can be written out as one long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$

We need to train the network:

What is known? What is unknown?

Entire network can be written out as a long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$



We need to train the network:

What is known? What is unknown?

Entire network can be written out as a long equation

$$y = f_3(w_3 \cdot f_2(w_2 \cdot f_1(w_1 \cdot x + b_1)))$$

activation function
sometimes has
parameters

unknown

We need to train the network:

What is known? What is unknown?

Learning an MLP

Given a set of samples and a MLP

$$\{x_i, y_i\}$$

$$y = f_{\text{MLP}}(x; \theta)$$

Estimate the parameters of the MLP

$$\theta = \{f, w, b\}$$

Gradient Descent

For each **random** sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

2. Update

a. Back Propagation

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives

b. Gradient update

$$\theta \leftarrow \theta - \eta \nabla \theta$$


vector of parameter update equations

So we need to compute the partial derivatives

$$\frac{\partial \mathcal{L}}{\partial \theta} = \left[\frac{\partial \mathcal{L}}{\partial w_3} \frac{\partial \mathcal{L}}{\partial w_2} \frac{\partial \mathcal{L}}{\partial w_1} \frac{\partial \mathcal{L}}{\partial b} \right]$$


Remember,

Partial derivative $\frac{\partial L}{\partial w_1}$ describes...




So, how do you compute it?


The Chain Rule



According to the chain rule...


$$\frac{\partial L}{\partial w_3} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$





$$\frac{\partial L}{\partial w_3} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$


Chain Rule!




$$\begin{aligned}\frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}\end{aligned}$$

Just the partial derivative of L2 loss






$$\begin{aligned}\frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}\end{aligned}$$



Let's use a Sigmoid function


$$\frac{ds(x)}{dx} = s(x)(1 - s(x))$$




$$\begin{aligned}
 \frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) f_3(1 - f_3) \frac{\partial a_3}{\partial w_3}
 \end{aligned}$$

Let's use a Sigmoid function


$$\frac{ds(x)}{dx} = s(x)(1 - s(x))$$



$$\begin{aligned}
 \frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) f_3(1 - f_3) \frac{\partial a_3}{\partial w_3} \\
 &= -\eta(y - \hat{y}) f_3(1 - f_3) f_2
 \end{aligned}$$




$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$



$$\frac{\partial L}{\partial w_2} = \frac{\frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3}}{\frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2}} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$


already computed.
re-use (propagate)!

The Chain Rule




a.k.a. backpropagation

The chain rule says...




$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

The chain rule says...



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

already computed.
re-use (propagate)!



$$\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3}$$

$$\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2}$$

$$\frac{\partial \mathcal{L}}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}$$

Gradient Descent

For each example sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i$$

2. Update

a. Back Propagation

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_3} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial w_3} \\ \frac{\partial \mathcal{L}}{\partial w_2} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial w_2} \\ \frac{\partial \mathcal{L}}{\partial w_1} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial b} &= \frac{\partial \mathcal{L}}{\partial f_3} \frac{\partial f_3}{\partial a_3} \frac{\partial a_3}{\partial f_2} \frac{\partial f_2}{\partial a_2} \frac{\partial a_2}{\partial f_1} \frac{\partial f_1}{\partial a_1} \frac{\partial a_1}{\partial b}\end{aligned}$$

b. Gradient update

$$w_3 = w_3 - \eta \nabla w_3$$

$$w_2 = w_2 - \eta \nabla w_2$$

$$w_1 = w_1 - \eta \nabla w_1$$

$$b = b - \eta \nabla b$$

Gradient Descent

For each example sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i$$

2. Update

a. Back Propagation

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives

b. Gradient update

$$\theta \leftarrow \theta + \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter update equations

Stochastic gradient descent

What we are truly minimizing:

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

The gradient is:

What we are truly minimizing:

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

The gradient is:

$$\sum_{i=1}^N \frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta}$$

What we use for gradient update is:

What we are truly minimizing:

$$\min_{\theta} \sum_{i=1}^N L(y_i, f_{MLP}(x_i))$$

The gradient is:

$$\sum_{i=1}^N \frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta}$$

What we use for gradient update is:

$$\frac{\partial L(y_i, f_{MLP}(x_i))}{\partial \theta} \quad \text{for some } i$$

Stochastic Gradient Descent

For each example sample

$$\{x_i, y_i\}$$

1. Predict

a. Forward pass

$$\hat{y} = f_{\text{MLP}}(x_i; \theta)$$

b. Compute Loss

$$\mathcal{L}_i$$

2. Update

a. Back Propagation

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter partial derivatives

b. Gradient update

$$\theta \leftarrow \theta + \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

vector of parameter update equations

How do we select which sample?

- Select randomly!

Do we need to use only one sample?

- You can use a *minibatch* of size $B < N$.

Why not do gradient descent with all samples?


- It's very expensive when N is large (big data).

Do I lose anything by using stochastic GD?

- Same convergence guarantees and complexity!
- Better generalization.

Convolution Neural Networks (ConvNet)

Convolution Neural Networks



Motivation



10 BREAKTHROUGH TECHNOLOGIES 2013

Introduction

The 10 Technologies

Past Years

Deep Learning

With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart.

Temporary Social Media

Messages that quickly self-destruct could enhance the privacy of online communications and make people freer to be spontaneous.

Prenatal DNA Sequencing

Reading the DNA of fetuses will be the next frontier of the genomic revolution. But do you really want to know about the genetic problems or musical aptitude of your unborn child?

Additive Manufacturing

Skeptical about 3-D printing? GE, the world's largest manufacturer, is on the verge of using the technology to make jet parts.

Baxter: The Blue-Collar Robot

Rodney Brooks's newest creation is easy to interact with, but the complex innovations behind the robot show just how hard it is to get along with people.

Memory Implants


Smart Watches

Ultra-Efficient Solar

Big Data from

Supergrids

Recap: Before Deep Learning



*Input
Pixels*


*Extract
Features*

*Concatenate into
a vector x*

*Linear
Classifier*

The last layer of (most) CNNs are linear classifiers

This piece is just a linear classifier




*Input
Pixels*


*Perform everything with a big neural
network, trained end-to-end*

Key: perform enough processing so that by the time you get to the end of the network, the classes are linearly separable


What shape should the activations have?



- The input is an image, which is 3D (RGB channel, height, width)




What shape should the activations have?



- The input is an image, which is 3D (RGB channel, height, width)
- We could flatten it to a 1D vector, but then we lose structure

What shape should the activations have?



- The input is an image, which is 3D (RGB channel, height, width)
- We could flatten it to a 1D vector, but then we lose structure
- What about keeping everything in 3D?

ConvNets

They're just neural networks with
3D activations and weight sharing

3D Activations

before:





Figure: Andrej Karpathy


3D Activations

before:



(1D vectors)

now:



(3D arrays)

Figure: Andrej Karpathy

3D Activations

All Neural Net
activations
arranged in 3
dimensions:





Figure: Andrej Karpathy

3D Activations

All Neural Net activations arranged in **3 dimensions**:



For example, a CIFAR-10 image is a $3 \times 32 \times 32$ volume
(3 depth — RGB channels, 32 height, 32 width)

Figure: Andrej Karpathy

3D Activations

1D Activations:





Figure: Andrej Karpathy

3D Activations

1D Activations:



3D Activations:





Figure: Andrej Karpathy


3D Activations



- The input is $3 \times 32 \times 32$
- This neuron depends on a $3 \times 5 \times 5$ chunk of the input
- The neuron also has a $3 \times 5 \times 5$ set of weights and a bias (scalar)

Figure: Andrej Karpathy

3D Activations




Example: consider the region of the input " x^r "

With output neuron h^r

Figure: Andrej Karpathy

3D Activations



Example: consider the region of the input “ x^r ”


With output neuron h^r

Then the output is:

$$h^r = \sum_{ijk} x^r_{ijk} W_{ijk} + b$$

Figure: Andrej Karpathy

3D Activations



Example: consider the region of the input “ x^r ”

With output neuron h^r

Then the output is:

$$h^r = \sum_{ijk} x^r_{ijk} W_{ijk} + b$$



Sum over 3 axes

Figure: Andrej Karpathy

3D Activations




Figure: Andrej Karpathy

3D Activations





Figure: Andrej Karpathy

3D Activations




With **2** output neurons

$$h^r_1 = \sum_{ijk} x^r_{ijk} W_{1ijk} + b_1$$

$$h^r_2 = \sum_{ijk} x^r_{ijk} W_{2ijk} + b_2$$

Figure: Andrej Karpathy

3D Activations



With **2** output neurons

$$h^r_1 = \sum_{ijk} x^r_{ijk} W_{1ijk} + b_1$$

$$h^r_2 = \sum_{ijk} x^r_{ijk} W_{2ijk} + b_2$$

Figure: Andrej Karpathy

3D Activations





Figure: Andrej Karpathy

3D Activations




We can keep adding more outputs

These form a column in the output volume:
[depth x 1 x 1]

Figure: Andrej Karpathy

3D Activations




We can keep adding more outputs

These form a column in the output volume:
[depth x 1 x 1]

Figure: Andrej Karpathy


3D Activations



Now repeat this
across the input

Figure: Andrej Karpathy

3D Activations




Now repeat this across the input

Weight sharing:
Each filter shares
the same weights
(but each depth
index has its own
set of weights)

Figure: Andrej Karpathy


3D Activations



With weight sharing,
this is called
convolution

Figure: Andrej Karpathy

3D Activations




With weight sharing,
this is called
convolution

Without weight sharing,
this is called a
locally connected layer

Figure: Andrey Karpathy

3D Activations




One set of weights gives
one slice in the output

To get a 3D output of depth D ,
use D different filters

In practice, ConvNets use
many filters (~ 64 to 1024)

3D Activations



One set of weights gives
one slice in the output

To get a 3D output of depth D ,
use D different filters

In practice, ConvNets use
many filters (~ 64 to 1024)

All together, the weights are **4** dimensional:
(output depth, input depth, kernel height, kernel width)

3D Activations

We can unravel the 3D cube and show each layer separately:

(Input)




Figure: Andrej Karpathy

3D Activations

We can unravel the 3D cube and show each layer separately:

(Input)




Figure: Andrej Karpathy

3D Activations

We can unravel the 3D cube and show each layer separately:

(Input)





Figure: Andrej Karpathy

3D Activations


We can unravel the 3D cube and show each layer separately:

(Input)



ConvNet


A **ConvNet** is a sequence of convolutional layers, interspersed with activation functions (and possibly other layer types)



ConvNet

Convolution Layer


32x32x3 image



ConvNet

Convolution Layer

32x32x3 image



5x5x3 filter




Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

ConvNet

Convolution Layer

32x32x3 image



Filters always extend the full depth of the input volume


5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”


ConvNet

Convolution Layer




ConvNet

Convolution Layer




ConvNet

Convolution Layer



ConvNet

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size $28 \times 28 \times 6$!

CNNs Notations



Stride


Padding



Pooling




Stride



Animations: https://github.com/vdumoulin/conv_arithmetic


Stride

During convolution, the weights “slide” along the input to generate each output




Stride

During convolution, the weights “slide” along the input to generate each output




Input




Output

Stride

During convolution, the weights “slide” along the input to generate each output




Input




Output

Stride

During convolution, the weights “slide” along the input to generate each output




Input




Output

Stride

During convolution, the weights “slide” along the input to generate each output




Input




Output

Stride

During convolution, the weights “slide” along the input to generate each output




Input



Output

Stride

During convolution, the weights “slide” along the input to generate each output



Input


Recall that at each position,
we are doing a **3D** sum:

$$h^r = \sum_{ijk} x^r_{ijk} W_{ijk} + b$$

(channel, row, column)


Stride

But we can also convolve with a **stride**, e.g. stride = 2




Stride

But we can also convolve with a **stride**, e.g. stride = 2




Input




Output

Stride

But we can also convolve with a **stride**, e.g. stride = 2




Input




Output

Stride

But we can also convolve with a **stride**, e.g. stride = 2



Input



Output

- Notice that with certain strides, we may not be able to cover all of the input
- The output is also half the size of the input

CNNs Notations



Stride

Padding




Pooling



PADDING

CONTENT

Padding



Padding

We can also pad the input with zeros.

Here, **pad = 1, stride = 2**

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input

0			

Output

Padding

We can also pad the input with zeros.

Here, **pad = 1, stride = 2**

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input

Output

Padding

We can also pad the input with zeros.

Here, **pad = 1, stride = 2**

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input

Output

Padding

We can also pad the input with zeros.


Here, **pad = 1, stride = 2**

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input

Output


How big is the output?



In general, the output has size:

$$w_{\text{out}} = \left\lfloor \frac{w_{\text{in}} + 2p - k}{s} \right\rfloor + 1$$

How big is the output?




Example: $k=3$, $s=1$, $p=1$


$$\begin{aligned}w_{out} &= \left\lfloor \frac{w_{in} + 2p - k}{s} \right\rfloor + 1 \\&= \left\lfloor \frac{w_{in} + 2 - 3}{1} \right\rfloor + 1 \\&= w_{in}\end{aligned}$$

VGGNet [Simonyan 2014]
uses filters of this shape

Other variations?



Transposed



Dilation

More info? Check this <https://arxiv.org/abs/1603.07285>

CNNs Notations



Stride


Padding



Pooling



Pooling



Convolved
feature

Pooled
feature

Pooling

For most ConvNets, **convolution** is often followed by **pooling**:

- Creates a smaller representation while retaining the most important information
- The “max” operation is the most common
- Why might “avg” be a poor choice?






Figure: Andrej Karpathy

Pooling

- makes the representations smaller and more manageable
- operates over each activation map independently:



Max Pooling




What's the backprop rule for max pooling?

- In the forward pass, store the index that took the max
- The backprop gradient is the input gradient at that index

Figure: Andrej Karpathy

Example: AlexNet [Krizhevsky 2012]



“max”: max pooling


“norm”: local response normalization

“full”: fully connected


Figure: [Karnowski 2015] (*with corrections*)

Example ConvNet


Deep neural networks learn hierarchical feature representations



Hierarchical Feature representation



Visual embedding (Img2Vec)



A ‘feature vector’ of an image is simply a list of numbers taken from the output of a neural network layer.


This vector is a dense representation of the input image, and can be used for a variety of tasks such as ranking, classification, or clustering.

How to train ConvNets?

How to train ConvNets?

Roughly speaking:

Gather
labeled data



Find a ConvNet
architecture



Minimize
the loss



How to train ConvNets?

- Split and preprocess your data
- Choose your network architecture
- Initialize the weights
- Find a learning rate and regularization strength
- Minimize the loss and monitor progress
- Fiddle with knobs

Mini-batch Gradient Decent

Loop:


1. Sample a batch of training data (~100 images)
2. Forwards pass: compute loss (avg. over batch)
3. Backwards pass: compute gradient
4. Update all parameters

Regularization

Regularization reduces overfitting:

$$L = L_{\text{data}} + L_{\text{reg}}$$

$$L_{\text{reg}} = \lambda \frac{1}{2} \|W\|_2^2$$




[Andrej Karpathy <http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>]

Regularization

Overfitting: modeling noise in the training set instead of the “true” underlying relationship

Underfitting: insufficiently modeling the relationship in the training set

General rule: models that are “bigger” or have more capacity are more likely to overfit



[Image: https://en.wikipedia.org/wiki/File:Overfitted_Data.png]

1) Data pre-processing

Preprocess the data so that learning is better conditioned:





Figure: Andrej Karpathy

1) Data pre-processing

For ConvNets, typically only the mean is subtracted.



An input image (256x256)



Minus sign




The mean input image

A per-channel mean also works (one value per R,G,B).

Figure: Alex Krizhevsky

1) Data pre-processing

Augment the data — extract random crops from the input, with slightly jittered offsets. Without this, typical ConvNets (e.g. [Krizhevsky 2012]) overfit the data.



E.g. 224x224 patches
extracted from 256x256 images

Randomly reflect horizontally

Perform the augmentation live
during training


Figure: Alex Krizhevsky

1) Data pre-processing

Here are few tricks used by the AlexNet team.



Mirror Image



Without data augmentation, the authors would not have been able to use such a large network because it would have suffered from substantial overfitting.



256


Random Crops

227



2) Choose your architecture

Toy example: one hidden layer of size 50



Slide: Andrej Karpathy

3) Initialize your weights

Set the weights to small random numbers:

```
W = np.random.randn(D, H) * 0.001
```

(matrix of small random numbers drawn from a Gaussian distribution)

Set the bias to zero (or small nonzero):

```
b = np.zeros(H)
```

Slide: Andrej Karpathy

4) Find a learning rate


Let's start with small regularization and find the learning rate that makes the loss decrease:

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                   model, two_layer_net,
                                   num_epochs=10, reg=0.000001, ←
                                   update='sgd', learning_rate_decay=1,
                                   sample_batches = True,
                                   learning_rate=1e-6, verbose=True)
```

new weight = weight - learning rate*gradient

4) Find a learning rate

Learning rate: $1e6$ — what could go wrong?



4) Find a learning rate

Normally, you don't have the budget for lots of cross-validation —> visualize as you go

Plot the loss

For very small learning rates, the loss decreases linearly and slowly

(*Why linearly?*)

Larger learning rates tend to look more exponential





Figure: Andrej Karpathy

4) Find a learning rate

Visualize the accuracy




Big gap: overfitting
(increase regularization)

No gap: underfitting
(increase model capacity,
make layers bigger
or decrease regularization)

Figure: Andrej Karpathy

4) Find a learning rate

Visualize the weights



Nice clean weights:
training is proceeding well




Figure: Alex Krizhevsky , Andrej Karpathy


What to fiddle?

- Network architecture
- Learning rate, decay schedule, update type
- Regularization (L2, L1, maxnorm, dropout, ...)
- Loss function (softmax, SVM, ...)
- Weight initialization

Neural network
parameters




Example: AlexNet [Krizhevsky 2012]



Dropout



Dropout is yet another approach to reduce overfitting!



When a neuron is dropped, it does not contribute to either forward or backward propagation. So every input goes through a different network architecture, as shown in the animation. As a result, the learnt weight parameters are more robust and do not get overfitted easily.

Dropout

Simple but powerful technique to reduce overfitting:




During testing, there is no dropout and the whole network is used.

[Srivasta et al, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”, JMLR 2014]

Dropout

Simple but powerful technique to reduce overfitting:



Note: Dropout can be interpreted as an approximation to taking the geometric mean of an ensemble of exponentially many models


[Srivasta et al, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”, JMLR 2014]

Dropout

Case study: [Krizhevsky 2012]

“Without dropout, our network exhibits substantial overfitting.”

Dropout here



But not here – why?


[Krizhevsky et al, “ImageNet Classification with Deep Convolutional Neural Networks”, NIPS 2012]

Transfer Learning

Transfer Learning with CNNs

Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014


1. Train on Imagenet



2. Small Dataset (C classes)



3. Bigger dataset



Transfer Learning

FC-1000
FC-4096
FC-4096

MaxPool

Conv-512

Conv-512

MaxPool

Conv-512

Conv-512

MaxPool

Conv-256

Conv-256

MaxPool

Conv-128

Conv-128

MaxPool


Conv-64

Conv-64

Image

More specific


More generic



	very similar dataset	very different dataset
very little data	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

Recurrent Neural Networks

Neural Network



Temporal dependencies

Analyzing temporal dependencies → Improved decisions

Frame 0

Stem: seen
Petals: hidden



$P(Iris): 0.1$
 $P(\neg Iris): 0.9$

Frame 1

Stem: seen
Petals: hidden



$P(Iris): 0.11$
 $P(\neg Iris): 0.89$

Frame 2

Stem: seen
Petals: partial



$P(Iris): 0.2$
 $P(\neg Iris): 0.8$

Frame 3

Stem: partial
Petals: partial



$P(Iris): 0.45$
 $P(\neg Iris): 0.55$

Frame 4


Stem: hidden
Petals: seen



$P(Iris): 0.9$
 $P(\neg Iris): 0.1$



Decision on
sequence of
observations





Memory is important → Reasoning relies on experience

Sequential Data

- Sometimes the sequence of data matters.
 - Text generation
 - Stock price prediction
- **For example: The clouds are in the ?**
 - **sky**


Sequential Data

- Sometimes the sequence of data matters.
 - Text generation
 - Stock price prediction
- **For example: The clouds are in the ?**
 - **sky**
- Simple solution: N-grams?
 - Hard to represent patterns with more than a few words
(possible patterns increases exponentially)

Sequential Data

- Sometimes the sequence of data matters.
 - Text generation
 - Stock price prediction
- **For example: The clouds are in the ?**
 - **sky**
- Simple solution: N-grams?
 - Hard to represent patterns with more than a few words
(possible patterns increases exponentially)
- Simple solution: Neural networks?
 - Fixed input/output size
 - Fixed number of steps

Time-delay neural network



Pro: Dependencies between features at different timestamps

Cons:


- **Limited** history of the input (< 10 timestamps)
- **Delay values** should be set explicitly
- **Not general**, can not solve complex tasks

Recurrent neural networks

- Recurrent neural networks (RNNs) are networks with loops, allowing information to persist [Rumelhart et al., 1986].

$$h_t = f_W(h_{t-1}, x_t)$$


new state old state input vector at some time step
some function with parameters W



- Have **memory** that keeps track of information observed so far
- Maps from the entire history of previous inputs to each output
- Handle sequential data


Neural Networks

one to one




Vanilla Neural Networks

Recurrent Neural Networks




Recurrent Neural Networks




Recurrent Neural Networks


one to one




one to many




many to one



many to many




many to many



e.g. **Machine Translation**
seq of words -> seq of words


Recurrent Neural Networks



Recurrent neural networks

$$\mathbf{h}_t = \theta \phi(\mathbf{h}_{t-1}) + \theta_x \mathbf{x}_t$$

$$\mathbf{y}_t = \theta_y \phi(\mathbf{h}_t)$$




- \mathbf{x}_t is the **input** at time t .
- \mathbf{h}_t is the **hidden state** (memory) at time t .
- \mathbf{y}_t is the **output** at time t .
- $\theta, \theta_x, \theta_y$ are distinct **weights**.
 - weights are the same at all time steps.

Recurrent neural networks

We can process a sequence of vectors x by applying a recurrence formula at every time step:


$$h_t = f_W(h_{t-1}, x_t)$$




Notice: the same function and the same set of parameters are used at every time step.

Recurrent neural networks


- RNNs can be thought of as multiple copies of the same network, each passing a message to a successor.




RNN: Computational Graph



RNN: Computational Graph




RNN: Computational Graph




RNN: Computational Graph

First words get transformed into machine-readable vectors. Then the RNN processes the sequence of vectors one by one.




Animations by Michael Nguyen)

RNN: Computational Graph




The hidden state acts as the neural networks internal memory. It holds information on previous data the network has seen before.




Animations by Michael Nguyen)

RNN: Computational Graph


Re-use the same weight matrix at every time-step



RNN: Computational Graph: Many to One




RNN: Computational Graph: One to Many



Sequence to Sequence: Many-to-one + one-to-many

Many to one: Encode input sequence in a single vector




One to many: Produce output sequence from single input vector


How to train RNNs?

Back-Propagation Through Time (BPTT)


- Using the generalized back-propagation algorithm one can obtain the so-called **Back-Propagation Through Time** algorithm.
- The recurrent model is represented as a multi-layer one (with an unbounded number of layers) and backpropagation is applied on the unrolled model.



BPTT




Truncated BPTT




Run forward and backward
through chunks of the
sequence instead of whole
sequence

Truncated BPTT




Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps

Truncated BPTT



How does gradient flow in RNN?


RNN Gradient Flow



$$\begin{aligned} h_t &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\ &= \tanh \left(\begin{pmatrix} W_{hh} & W_{hx} \end{pmatrix} \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \\ &= \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \end{aligned}$$

Why the activation function is Tanh?


- The tanh activation is used to help regulate the values flowing through the network. The tanh function squishes values to always be between -1 and 1.



Animations from Michael Nguyen


RNN Gradient Flow

Backpropagation from h_t
to h_{t-1} multiplies by W
(actually W_{hh}^T)



$$\begin{aligned} h_t &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\ &= \tanh \left(\begin{pmatrix} W_{hh} & W_{hx} \end{pmatrix} \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \\ &= \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \end{aligned}$$

RNN Gradient Flow




Computing gradient
of h_0 involves many
factors of W
(and repeated tanh)

RNN Gradient Flow

Bengio et al, "Learning long-term dependencies with gradient descent is difficult", IEEE Transactions on Neural Networks, 1994

Pascanu et al, "On the difficulty of training recurrent neural networks", ICML 2013




Computing gradient
of h_0 involves many
factors of W
(and repeated tanh)

Largest singular value > 1 :
Exploding gradients

Largest singular value < 1 :
Vanishing gradients

RNN Gradient Flow



Computing gradient of h_0 involves many factors of W (and repeated tanh)


Largest singular value > 1 :
Exploding gradients

Largest singular value < 1 :
Vanishing gradients

Gradient clipping: Scale gradient if its norm is too big

```
grad_norm = np.sum(grad * grad)
if grad_norm > threshold:
    grad *= (threshold / grad_norm)
```

RNN Gradient Flow



Computing gradient
of h_0 involves many
factors of W
(and repeated tanh)

Largest singular value > 1 :
Exploding gradients


Largest singular value < 1 :
Vanishing gradients

new weight = weight - learning rate * gradient




The Problem of Long-term Dependencies

- RNNs connect previous information to present task:
 - may be enough for predicting the next word for "the clouds are in the **sky**"



- may not be enough when more context is needed




Short-Term memory

- RNNs suffer from what is known as short-term memory!

I was born in France, but I have been working in South Africa
working for ... (another 200 words) ... Therefore my mother tongue
is:



RNN Gradient Flow




Computing gradient
of h_0 involves many
factors of W
(and repeated tanh)

Largest singular value > 1 :
Exploding gradients

Largest singular value < 1 :
Vanishing gradients

→ Change RNN architecture

new weight = weight - learning rate * gradient



Long Short-Term Memory Networks (LSTM)

Vanilla RNN

$$h_t = \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right)$$

LSTM

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$


$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

Hochreiter and Schmidhuber, “Long Short Term Memory”, Neural Computation
1997



Long Short-Term Memory Networks

- Long Short-Term Memory (LSTM) networks are RNNs capable of learning long-term dependencies [Hochreiter and Schmidhuber, 1997].



- A memory cell using logistic and linear units with multiplicative interactions:
 - Information gets into the cell whenever its **input gate** is on.
 - Information is thrown away from the cell whenever its **forget gate** is off.
 - Information can be read from the cell by turning on its **output gate**

Notation



Adapted from: C. Olah

LSTM overview

- We define the LSTM unit at each time step t to be a collection of vectors in \mathbb{R}^d :

- Memory cell c_t**

$$\tilde{c}_t = \text{Tanh}(W_c \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_c) \text{ vector of new candidate values}$$

$$c_t = f_t * c_{t-1} + i_t * \tilde{c}_t$$

- Forget gate f_t** in $[0, 1]$: scales old memory cell value (**reset**)

$$f_t = \sigma(W_f \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f)$$

- Input gate i_t** in $[0, 1]$: scales input to memory cell (**write**)

$$i_t = \sigma(W_i \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i)$$

- Output gate o_t** in $[0, 1]$: scales output from memory cell (**read**)


$$o_t = \sigma(W_o \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o)$$

- Output \mathbf{h}_t**

$$\mathbf{h}_t = o_t * \text{Tanh}(c_t)$$


Memory Cell

- Information can flow along the **memory cell unchanged**.
- Information can be **removed** or **written** to the **memory cell**, regulated by gates.




Gates

- **Gates** are a way to optionally let information through.
 - A **sigmoid layer** outputs number between 0 and 1, **deciding** how much of each component should be let through.
 - A pointwise multiplication operation applies the decision.



Sigmoid activation function

- Gates contains **sigmoid activations**. A sigmoid activation is similar to the tanh activation. Instead of squishing values between -1 and 1, it squishes values between 0 and 1.




- That is helpful to update or forget data because any number getting multiplied by 0 is 0, causing values to disappear or be “forgotten.” Any number multiplied by 1 is the same value therefore that value stay’s the same or is “kept.”

Animations from Michael Nguyen


Forget Gate

- A **sigmoid** layer, **forget gate**, **decides** which values of the **memory cell** to **reset**.



$$\mathbf{f}_t = \sigma(W_f \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f)$$


Forget Gate



Animations from Michael Nguyen


Input Gate

- A **sigmoid** layer, **input gate**, **decides** which values of the **memory cell** to **write** to.



$$\mathbf{i}_t = \sigma(W_i \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i)$$


Input Gate



Animations from Michael Nguyen

Vector of New Candidate Values


- A **Tanh** layer creates a **vector of new candidate values** \tilde{c}_t to **write** to the **memory cell**.



$$\tilde{c}_t = \text{Tanh}(W_c \cdot [h_{t-1}, x_t] + b_c)$$


Memory Cell Update

- The previous steps decided which values of the **memory cell** to **reset** and **overwrite**.
- Now the LSTM **applies the decisions** to the **memory cell**.



$$c_t = f_t * c_{t-1} + i_t * \tilde{c}_t$$


Memory Cell Update



Animations from Michael Nguyen

Output Gate


- A sigmoid layer, **output gate**, decides which values of the memory cell to **output**.



$$\mathbf{o}_t = \sigma(W_o \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o)$$


Output Update

- The **memory cell** goes through **Tanh** and is multiplied by the **output gate**.



$$h_t = \textcolor{red}{o}_t * \text{Tanh}(c_t)$$


Output Update



Animations from Michael Nguyen

How does gradient flow in LSTM?

Long Short-Term Memory Networks (LSTM)




$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

LSTM Gradient Flow

- f:** Forget gate, Whether to erase cell
- i:** Input gate, whether to write to cell
- g:** Gate gate (?), How much to write to cell
- o:** Output gate, How much to reveal cell



Backpropagation from c_t to c_{t-1} only elementwise multiplication by f , no matrix multiply by W


$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

LSTM Gradient Flow


Uninterrupted gradient flow!




The gradient behaves similarly to the forget gate, and if the forget gate decides that a certain piece of information should be remembered, it will be open and have values closer to 1 to allow for information flow.

RNN vs. LSTM

RNN



LSTM
(ignoring
forget gates)




Variants on LSTM

- Gate layers look at the memory cell [Gers and Schmidhuber, 2000].

$$\mathbf{f}_t = \sigma(W_f \cdot [\mathbf{c}_{t-1}, \mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f)$$


$$\mathbf{i}_t = \sigma(W_i \cdot [\mathbf{c}_{t-1}, \mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(W_o \cdot [\mathbf{c}_{t-1}, \mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o)$$



Variants on LSTM

- Use coupled **forget** and **input** gates. Instead of separately deciding what to **forget** and what to **add**, make those decisions together.




$$\mathbf{c}_t = \mathbf{f}_t * \mathbf{c}_{t-1} + (1 - \mathbf{f}_t) * \tilde{\mathbf{c}}_t$$

Variants on LSTM

- Gated Recurrent Unit (GRU) [Cho et al., 2014]:
 - Combine the **forget** and **input** gates into a single **update** gate.
 - **Merge the memory cell and the hidden state.**
 - ...

$$\begin{aligned} z_t &= \sigma(W_z \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t]) \\ r_t &= \sigma(W_r \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t]) \\ \tilde{\mathbf{h}}_t &= \text{Tanh}(W \cdot [r_t * \mathbf{h}_{t-1}, \mathbf{x}_t]) \\ \mathbf{h}_t &= (1 - z_t) * \mathbf{h}_{t-1} + (z_t) * \tilde{\mathbf{h}}_t \end{aligned}$$




Summary

- RNNs allow a lot of flexibility in architecture design
- Vanilla RNNs are simple but don't work very well
- Common to use LSTM or GRU: their additive interactions improve gradient flow
- Backward flow of gradients in RNN can explode or vanish.
- Exploding is controlled with gradient clipping. Vanishing is controlled with additive interactions (LSTM)
- Better/simpler architectures are a hot topic of current research
- Better understanding (both theoretical and empirical) is needed

Application: Image Captioning

Recurrent Neural Network



Convolutional Neural Network

Explain Images with Multimodal Recurrent Neural Networks, Mao et al.

Deep Visual-Semantic Alignments for Generating Image Descriptions, Karpathy and Fei-Fei

Show and Tell: A Neural Image Caption Generator, Vinyals et al.

Long-term Recurrent Convolutional Networks for Visual Recognition and Description, Donahue et al.

Learning a Recurrent Visual Representation for Image Caption Generation, Chen and Zitnick



Additional resources

- [1] Kyunghyun Cho et al. “Learning phrase representations using RNN encoder-decoder for statistical machine translation”. In: arXiv preprint arXiv:1406.1078 (2014).
- [2] Felix A Gers and Jürgen Schmidhuber. “Recurrent nets that time and count”. In: Neural Networks, 2000. IJCNN 2000. Vol. 3. IEEE. 2000, pp. 189–194.
- [3] Sepp Hochreiter and Jürgen Schmidhuber. “Long short-term memory”. In: Neural computation 9.8 (1997), pp. 1735–1780.
- [4] David E Rumelhart et al. “Sequential thought processes in PDP models”. In: V 2 (1986), pp. 3–57.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

<https://www.youtube.com/watch?v=56TYLaQN4N8&index=14&list=PLE6Wd9FR--EfW8dtjAuPoTuPcqmOV53Fu>

Additional resources

- Basic reading: No standard textbooks yet! Some good resources:
- <https://sites.google.com/site/deeplearningsummerschool/>
- <http://www.deeplearningbook.org/>
- <http://www.cs.toronto.edu/~hinton/absps/NatureDeepReview.pdf>